

# Reconstruction Metric - GPU Optimized

2026-05-19

## Table of contents

|      |                                                         |    |
|------|---------------------------------------------------------|----|
| 0.1  | 1. Setup e Imports                                      | 1  |
| 0.2  | 2. Cargar Datos                                         | 5  |
| 0.3  | 3. Cargar SAE y Extraer Features                        | 5  |
| 0.4  | 4. Split Train/Test y Filtrar BSPs de Piezas            | 7  |
| 0.5  | 5. Funciones GPU-Optimizadas                            | 8  |
| 0.6  | 6. Fase 1: Identificar Features de Alta Precisión (GPU) | 9  |
| 0.7  | 7. Fase 2: Reconstruir Tableros en Test Set (GPU)       | 12 |
| 0.8  | 8. Análisis de Resultados                               | 14 |
| 0.9  | 9. Guardar Resultados                                   | 16 |
| 0.10 | 10. Generar PDF con Quarto                              | 17 |

## 0.1 1. Setup e Imports

```
import numpy as np
import pandas as pd
import sys
import os
import subprocess
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt

from pathlib import Path
from tqdm import tqdm
from sae.tools.experiment_utils import get_experiment_dir, get_metrics_dir, get_report_name
from sae.tools.naming_utils import get_checkpoint_dir, get_model_name, get_extras_id

project_root = Path('...').resolve()
sys.path.insert(0, str(project_root))

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Device: {device}")
if torch.cuda.is_available():
    print(f"GPU: {torch.cuda.get_device_name(0)}")
    print(f"Memoria: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
```

Device: cuda  
GPU: NVIDIA GeForce RTX 5060  
Memoria: 8.55 GB

```
# Configuración para visualización en PDF
pd.set_option('display.max_colwidth', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 100)
```

```

np.set_printoptions(linewidth=100, edgeitems=3)

os.environ['COLUMNS'] = '100'

print(" Configuración para PDF lista")

```

Configuración para PDF lista

```

NAMING_META = {
    'variante': 'matryoshkabatchtopk',
    'tecnica': 'matryoshkabatchtopk',
    'capa': 6,
    'juegos': 50000,
    'base_path': 'C:\\Users\\Esposa\\Documents\\Repos\\sae-othello-gpt',
    'experiment_base_path': os.path.abspath('../../experiments'),
    'extras': {
        'expansion': 16,
        'k': 81,
        'matryoshka_widths': [2048, 4096, 8192],
        'epochs': 1000,
        'patience': 20,
        'lr': 1e-3
    }
}

print('\n Metadata de naming:')
for key, value in NAMING_META.items():
    print(f' {key}: {value}')

```

Metadata de naming:

```

variante: matryoshkabatchtopk
tecnica: matryoshkabatchtopk
capa: 6
juegos: 50000
base_path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt
experiment_base_path: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments
extras: {'expansion': 16, 'k': 81, 'matryoshka_widths': [2048, 4096, 8192], 'epochs': 1000, 'patience': 20, 'lr': 0.001}

```

```

class MatryoshkaBatchTopKSAE(nn.Module):
    """
    Sparse Autoencoder con activación MatryoshkaBatchTopK.
    Corregido según paper (ec. 3-5) e implementación de los autores (sae.py).

    Correcciones aplicadas:
    - P1: encoding único con W_enc completo; BatchTopK global; prefijos compartidos
    - P1: BatchTopK preserva valores originales (x * indicator), threshold EMA para eval
    - P1: auxiliary loss para dead features (sae.py:198-221)
    - P2: make_decoder_weights_and_grad_unit_norm proyecta gradiente (sae.py:51-57)
    - P2: preprocessing input_unit_norm configurable (sae.py:36-43)
    - P3: W_dec inicializado como transpuesta de W_enc (sae.py:93-94)
    """
    def __init__(self, d_in, d_sae, k, matryoshka_widths,
                  n_batches_to_dead=20, top_k_aux=512, aux_penalty=1/32,
                  input_unit_norm=False):
        super().__init__()
        self.d_in = d_in
        self.d_sae = d_sae
        self.k = k
        self.matryoshka_widths = sorted(matryoshka_widths)
        assert self.matryoshka_widths[-1] == d_sae, \

```

```

        f"El último ancho ({self.matryoshka_widths[-1]}) debe ser d_sae ({d_sae})"

self.n_batches_to_dead = n_batches_to_dead
self.top_k_aux         = top_k_aux
self.aux_penalty       = aux_penalty
self.input_unit_norm   = input_unit_norm

self.W_enc = nn.Parameter(torch.empty(d_in, d_sae))
self.W_dec = nn.Parameter(torch.empty(d_sae, d_in))
self.b_dec = nn.Parameter(torch.zeros(d_in))

# P3: W_dec = W_enc^T normalizado (sae.py:93-94) - inicialización tied
nn.init.kaiming_uniform_(self.W_enc)
self.W_dec.data[:] = self.W_enc.t().data
self.W_dec.data[:] = self.W_dec / self.W_dec.norm(dim=-1, keepdim=True)

# P1: threshold persistente calibrado con EMA durante training (sae.py:97)
self.register_buffer('threshold', torch.tensor(0.0))
# P1: contador de batches sin activación por feature - para L_aux (sae.py:30-32)
self.register_buffer('num_batches_not_active', torch.zeros(d_sae))

# -----
# P2: Preprocessing del input (sae.py:36-48)
# -----

def preprocess_input(self, x):
    if self.input_unit_norm:
        x_mean = x.mean(dim=-1, keepdim=True)
        x = x - x_mean
        x_std = x.std(dim=-1, keepdim=True)
        x = x / (x_std + 1e-5)
        return x, x_mean, x_std
    return x, None, None

def postprocess_output(self, x_reconstruct, x_mean, x_std):
    if self.input_unit_norm and x_mean is not None:
        x_reconstruct = x_reconstruct * x_std + x_mean
    return x_reconstruct

# -----
# P1: Encoding único + BatchTopK global (sae.py:100-119)
# -----

def compute_activations(self, x_cent):
    # P1: encoding ÚNICO con W_enc completo - un solo pase para todos los anchos
    # Las features 0:w son literalmente las mismas en todos los niveles (paper ec. 3)
    pre_acts = x_cent @ self.W_enc      # [batch, d_sae]
    acts     = F.relu(pre_acts)         # ReLU antes del topk (sae.py:102)

    if self.training:
        # P1: BatchTopK global sobre d_sae*batch_size valores (paper ec. 6)
        # Preserva valores originales con scatter - NO resta threshold (sae.py:104-114)
        topk_result = torch.topk(acts.flatten(), self.k * x_cent.shape[0])
        acts_topk = (
            torch.zeros_like(acts.flatten())
            .scatter(-1, topk_result.indices, topk_result.values)
            .reshape(acts.shape)
        )
        self._update_threshold(acts_topk)
    else:
        acts.mul_(acts > self.threshold)
    return acts, acts

```

```

        return acts, acts_topk

@torch.no_grad()
def _update_threshold(self, acts_topk, lr=0.01):
    # P1: EMA del mínimo valor positivo - calibra para eval (sae.py:224-228)
    positive_mask = acts_topk > 0
    if positive_mask.any():
        min_positive = acts_topk[positive_mask].min()
        self.threshold = (1 - lr) * self.threshold + lr * min_positive

@torch.no_grad()
def _update_inactive_features(self, acts_topk):
    # P1: necesario para L_aux - trackea qué features no se activan (sae.py:60-62)
    self.num_batches_not_active += (acts_topk.sum(0) == 0).float()
    self.num_batches_not_active[acts_topk.sum(0) > 0] = 0

# -----
# Forward
# -----

def forward(self, x):
    """
    Retorna:
        outputs      : lista de (recon_w, hidden_w) para cada ancho matryoshka
                       recon_w está en espacio normalizado (= espacio original si input_unit_norm=False)
        all_acts      : pre-topk [batch, d_sae] - para L_aux
        all_acts_topk: post-topk [batch, d_sae]
        x_proc        : x preprocesado (= x si input_unit_norm=False)
    """
    x_proc, x_mean, x_std = self.preprocess_input(x)
    x_cent = x_proc - self.b_dec
    all_acts, all_acts_topk = self.compute_activations(x_cent)

    # P1: reconstrucciones por PREFIJO compartido (paper ec. 4, sae.py:149-155)
    # all_acts_topk[:, 0:w] son las mismas features en todos los niveles anidados
    # recon_w se mantiene en espacio normalizado para que la loss compare espacios iguales
    outputs = []
    for w in self.matryoshka_widths:
        hidden_w = all_acts_topk[:, :w]
        recon_w = hidden_w @ self.W_dec[:, w, :] + self.b_dec
        outputs.append((recon_w, hidden_w))

    self._update_inactive_features(all_acts_topk)
    return outputs, all_acts, all_acts_topk, x_proc

# -----
# P1: Auxiliary loss para dead features (sae.py:198-221)
# -----

def get_auxiliary_loss(self, x_proc, x_reconstruct_full, all_acts):
    residual = x_proc.float() - x_reconstruct_full.float()
    aux_reconstruct = torch.zeros_like(residual)

    dead_features = self.num_batches_not_active >= self.n_batches_to_dead

    if dead_features.sum() > 0:
        acts_topk_aux = torch.topk(
            all_acts[:, dead_features],
            min(self.top_k_aux, int(dead_features.sum()))),
            dim=-1,
        )

```

```

        acts_aux = torch.zeros_like(all_acts[:, dead_features]).scatter(
            -1, acts_topk_aux.indices, acts_topk_aux.values
        )
        aux_reconstruct = aux_reconstruct + acts_aux @ self.W_dec[dead_features]

    if aux_reconstruct.abs().sum() > 0:
        return self.aux_penalty * (aux_reconstruct.float() - residual.float()).pow(2).mean()
    return torch.tensor(0.0, device=x_proc.device)

# -----
# P2: Constraint de norma unitaria con proyección del gradiente (sae.py:51-57)
# -----

@torch.no_grad()
def make_decoder_weights_and_grad_unit_norm(self):
    # P2: proyecta gradiente al espacio tangente de la esfera unitaria + normaliza
    # @torch.no_grad() evita trazar grafo interno - no impide leer .grad (sae.py:50)
    # Llamar ANTES de optimizer.step() (training.py:31)
    W_dec_normed = self.W_dec / self.W_dec.norm(dim=-1, keepdim=True)
    W_dec_grad_proj = (
        (self.W_dec.grad * W_dec_normed).sum(-1, keepdim=True) * W_dec_normed
    )
    self.W_dec.grad -= W_dec_grad_proj
    self.W_dec.data = W_dec_normed

```

## 0.2 2. Cargar Datos

```

# Cargar activaciones
activations_path = project_root / "activations" / "data" / "layer6_2000games.npy"
activations = np.load(activations_path)

print(f"Activaciones del modelo:")
print(f"  Shape: {activations.shape}")
print(f"  Memoria: {activations.nbytes / (1024**2):.2f} MB")

```

Activaciones del modelo:  
 Shape: (118000, 512)  
 Memoria: 230.47 MB

```

# Cargar ground truth de BSPs
bsp_gt_path = project_root / "metrics" / "02_data" / "bsp_ground_truth_2000games.npy"
bsp_names_path = project_root / "metrics" / "02_data" / "bsp_ground_truth_2000games.names.npy"

bsp_ground_truth = np.load(bsp_gt_path)
bsp_names = np.load(bsp_names_path, allow_pickle=True)

print(f"\nGround truth BSPs:")
print(f"  Shape: {bsp_ground_truth.shape}")
print(f"  Total BSPs: {len(bsp_names)}")

```

Ground truth BSPs:  
 Shape: (118000, 326)  
 Total BSPs: 326

## 0.3 3. Cargar SAE y Extraer Features

```

# Configuración del SAE
input_dim = 512
hidden_dim = 8192

# Construir ruta del modelo usando naming_utils
checkpoint_dir = get_checkpoint_dir(
    NAMING_META['base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos']
)
extras_id = get_extras_id(checkpoint_dir, NAMING_META['extras']) if NAMING_META.get('extras') else None

model_filename = get_model_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    'final',
    extras_id=extras_id
)
model_path = checkpoint_dir / model_filename

# Cargar modelo Matryoska SAE
sae = MatryoshkaBatchTopKSAE(d_in=input_dim, d_sae=hidden_dim, k=NAMING_META['extras']['k'], matryoshka_widths=NAMING_META['extras']['widths'])

checkpoint = torch.load(model_path, map_location=device, weights_only=False)
sae.load_state_dict(checkpoint['model_state_dict'])
sae.eval()

print(f"SAE cargado:")
print(f"  Path: {model_path}")
print(f"  Input: {input_dim}, Hidden: {hidden_dim}, k: {NAMING_META['extras']['k']}")
print(f"  Expansion: {hidden_dim/input_dim}x")

```

SAE cargado:

```

Path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_matryoshkabatchtopk\sae_matryoshkabatchtopk.pth
Input: 512, Hidden: 8192, k: 81
Expansion: 16.0x

```

```

# Extraer features del SAE en GPU
print("Extrayendo features del SAE...")
activations_tensor = torch.from_numpy(activations).float().to(device)

with torch.no_grad():
    outputs, all_acts, all_acts_topk, x_proc = sae(activations_tensor)
    sae_features = all_acts_topk

# Liberar tensores intermedios que ya no se necesitan
del outputs, all_acts, all_acts_topk, x_proc, activations_tensor
torch.cuda.empty_cache()

print(f"\nFeatures SAE:")
print(f"  Shape: {sae_features.shape}")
print(f"  Device: {sae_features.device}")
print(f"  Sparsity: {(sae_features == 0).float().mean():.2%}")
#print(f"  Activaciones promedio: {(sae_features > 0).sum(dim=1).float().mean():.1f}")

```

Extrayendo features del SAE...

```

Features SAE:
  Shape: torch.Size([118000, 8192])
  Device: cuda:0
  Sparsity: 99.01%

```

#### 0.4 4. Split Train/Test y Filtrar BSPs de Piezas

```

# Split 50/50:
n_moves = 59
n_games = 2000 # número de partidas para prueba de métricas
train_games = n_games // 2
split_idx = train_games * n_moves

sae_features_train = sae_features[:split_idx]
sae_features_test = sae_features[split_idx:]

print(f"Total partidas: {n_games}")
print(f"Train: {train_games} partidas ({split_idx} activaciones)")
print(f"Test: {n_games - train_games} partidas ({len(activations) - split_idx} activaciones)")
print(f"Train set: {sae_features_train.shape}")
print(f"Test set: {sae_features_test.shape}")

```

```

Total partidas: 2000
Train: 1000 partidas (59000 activaciones)
Test: 1000 partidas (59000 activaciones)
Train set: torch.Size([59000, 8192])
Test set: torch.Size([59000, 8192])

```

```

# =====
# CONFIGURACIÓN DE FILTRADO DE BSPs
# =====
USE_PLAYER = False # BSPs perspectiva del jugador (1 y 2)
USE_ABSOLUTE = True # BSPs absolutas negro/blanco (B y W)
USE_STRATEGIC = False # BSPs especiales (BSP_)

bsp_pieces_indices = []
bsp_pieces_names = []

for i, name in enumerate(bsp_names):
    include = False
    if USE_PLAYER and len(name) == 6 and name.startswith('BSP') and (name.endswith('1') or name.endswith('2')):
        include = True
    if USE_ABSOLUTE and (name.endswith('B') or name.endswith('W')):
        include = True
    if USE_STRATEGIC and name.startswith('BSP_'):
        include = True
    if include:
        bsp_pieces_indices.append(i)
        bsp_pieces_names.append(name)

bsp_pieces_indices = np.array(bsp_pieces_indices)

# Para reconstruction necesitas split train/test
bsp_gt_train_pieces = torch.from_numpy(bsp_ground_truth[:split_idx, bsp_pieces_indices]).bool().to(device)
bsp_gt_test_pieces = torch.from_numpy(bsp_ground_truth[split_idx:, bsp_pieces_indices]).bool().to(device)

# Convertir a tensor en GPU
# Identificador del filtrado (para naming de archivos de salida)
parts = []
if USE_PLAYER: parts.append("player")

```

```

if USE_ABSOLUTE: parts.append("absolute")
if USE_STRATEGIC: parts.append("strategic")
filter_suffix = "_".join(parts) if parts else "none"

print(f"BSPs seleccionadas para Reconstruction:")
print(f" Player: {USE_PLAYER} | Absolute: {USE_ABSOLUTE} | Strategic: {USE_STRATEGIC}")
print(f" Total: {len(bsp_pieces_indices)}")
print(f" Train shape: {bsp_gt_train_pieces.shape}")
print(f" Test shape: {bsp_gt_test_pieces.shape}")
print(f" Device: {bsp_gt_train_pieces.device}")
print(f" Primeras 5: {'', '.join(bsp_pieces_names[:5])}")
print(f" Últimas 5: {'', '.join(bsp_pieces_names[-5:])}")
print(f" Filter suffix: {filter_suffix}")

```

```

BSPs seleccionadas para Reconstruction:
Player: False | Absolute: True | Strategic: False
Total: 128
Train shape: torch.Size([59000, 128])
Test shape: torch.Size([59000, 128])
Device: cuda:0
Primeras 5: BSPA1B, BSPA1W, BSPA2B, BSPA2W, BSPA3B
Últimas 5: BSPH6W, BSPH7B, BSPH7W, BSPH8B, BSPH8W
Filter suffix: absolute

```

## 0.5 5. Funciones GPU-Optimizadas

```

def fast_precision_score_gpu(y_true, y_pred, eps=1e-8):
    """
    Precisión vectorizada en GPU.

    Args:
        y_true: Tensor (n_positions,) booleano
        y_pred: Tensor (n_positions, n_candidates) booleano

    Returns:
        precision_scores: Tensor (n_candidates,)
    """
    y_true_expanded = y_true.unsqueeze(1)

    tp = (y_true_expanded & y_pred).sum(dim=0).float()
    fp = (~y_true_expanded & y_pred).sum(dim=0).float()

    precision = tp / (tp + fp + eps)

    return precision

def fast_f1_score_gpu(y_true, y_pred, eps=1e-8):
    """
    F1 score vectorizado en GPU.

    Args:
        y_true: Tensor (n_positions,) booleano
        y_pred: Tensor (n_positions,) booleano

    Returns:
        f1: Float
    """
    tp = (y_true & y_pred).sum().float()
    fp = (~y_true & y_pred).sum().float()

```



```

fn = (y_true & ~y_pred).sum().float()

precision = tp / (tp + fp + eps)
recall = tp / (tp + fn + eps)

f1 = 2 * (precision * recall) / (precision + recall + eps)

return f1.item()

print(" Funciones GPU definidas")

```

Funciones GPU definidas

## 0.6 6. Fase 1: Identificar Features de Alta Precisión (GPU)

Para cada BSP, encontrar features con precisión 0.95 en el train set.

```

def identify_high_precision_features_gpu(
    sae_features_train,
    bsp_gt_train,
    f_max,
    precision_threshold=0.95,
    significance_threshold=1,
    thresholds=[0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9],
    feature_batch_size=512,
    verbose=True
):
    """
    Para cada (threshold t, feature f, BSP b): determina si f es clasificador
    de alta precisión para b al threshold t.

    Criterios (igual que implementación oficial):
        - precision(f, b, t) >= precision_threshold (default 0.95)
        - on_count(f, t) >= significance_threshold (default 10)

    Args:
        sae_features_train: Tensor (P, n_features) GPU
        bsp_gt_train: Tensor (P, n_bsps) GPU, bool
        f_max: Tensor (n_features,) GPU - calculado FUERA para
            garantizar identidad exacta con Fase 2
        precision_threshold: float, default 0.95
        significance_threshold: int, min. activaciones requeridas (default 10)
        thresholds: lista de fracciones t [0, 1)
        feature_batch_size: int, features por mini-batch interno

    Returns:
        hp_mask_TFB: BoolTensor (T, n_active, n_bsps)
        active_feature_indices: LongTensor (n_active,)
    """
    n_positions, n_features = sae_features_train.shape
    n_bsps = bsp_gt_train.shape[1]
    device = sae_features_train.device

    thresholds_T = torch.tensor(thresholds, device=device)
    n_thresholds = len(thresholds_T)

    # Features vivas (f_max > 0)
    active_feature_indices = (f_max > 0).nonzero(as_tuple=True)[0]
    n_active = len(active_feature_indices)

    if verbose:

```

```

print("=" * 60)
print("FASE 1: IDENTIFICAR FEATURES DE ALTA PRECISIÓN")
print("=" * 60)
print(f"Precision threshold:      {precision_threshold}")
print(f"Significance threshold:    >= {significance_threshold} activaciones")
print(f"Features activas:           {n_active}/{n_features} ({n_active/n_features*100:.1f}%)")
print(f"BSPs:                     {n_bsps}")
print(f"Thresholds:                  {thresholds}")
print("=" * 60)

# Thresholds absolutos: thresh_TF[t, f] = t * f_max[f]
active_f_max = f_max[active_feature_indices] # (n_active,)
thresholds_TF = thresholds_T[:, None] * active_f_max # (T, n_active)

# Acumuladores sobre todo el train set
on_count_TF = torch.zeros(n_thresholds, n_active, device=device)
tp_count_TFB = torch.zeros(n_thresholds, n_active, n_bsps, device=device)

bsp_float_PB = bsp_gt_train.float() # (P, B)

n_iters = (n_active + feature_batch_size - 1) // feature_batch_size

for fi in tqdm(range(n_iters), desc="Fase 1"):
    f_start = fi * feature_batch_size
    f_end = min(f_start + feature_batch_size, n_active)

    global_idx = active_feature_indices[f_start:f_end]
    acts_PF = sae_features_train[:, global_idx] # (P, bs)
    thresh_TF = thresholds_TF[:, f_start:f_end] # (T, bs)

    # active_TFP[t, f, p] = True si acts[p,f] > thresh[t,f]
    active_TFP = acts_PF.T.unsqueeze(0) > thresh_TF.unsqueeze(2) # (T, bs, P)

    # Conteo de activaciones por (threshold, feature)
    on_count_TF[:, f_start:f_end] += active_TFP.sum(dim=2).float()

    # True positives: f activa Y BSP verdadera
    # tp[t, f, b] = Σ_p active[t,f,p] * bsp[p,b]
    tp_count_TFB[:, f_start:f_end, :] += torch.einsum(
        'tfp,pb->tfb', active_TFP.float(), bsp_float_PB
    )

# Precisión: tp / on_count
eps = 1e-8
precision_TFB = tp_count_TFB / (on_count_TF.unsqueeze(2) + eps) # (T, F, B)

# Máscara: alta precisión AND significancia suficiente
sig_mask_TF1 = (on_count_TF >= significance_threshold).unsqueeze(2) # (T, F, 1)
hp_mask_TFB = (precision_TFB >= precision_threshold) & sig_mask_TF1 # (T, F, B)

if verbose:
    hp_per_t = hp_mask_TFB.sum(dim=(1, 2)) # (T,)
    hp_per_b = hp_mask_TFB.sum(dim=1) # (T, B)
    best_t = hp_per_t.argmax().item()
    print(f"\nPares (feature, BSP) de alta precisión por threshold:")
    for i, t in enumerate(thresholds):
        marker = " ← más pares" if i == best_t else ""
        print(f" t={t:.1f}: {hp_per_t[i].item():5d} pares{marker}")
    print(f"\nFeatures HP por BSP en t={thresholds[best_t]:.1f}:")
    print(f" Media: {hp_per_b[best_t].float().mean():.1f}")
    print(f" Máx: {hp_per_b[best_t].max().item()}")
    print(f" BSPs sin features: {(hp_per_b[best_t] == 0).sum().item()}")

```

```

    return hp_mask_TFB, active_feature_indices

print(" Fase 1 rediseñada: umbral de significancia (10) + estructura (T, F, B)")

```

Fase 1 rediseñada: umbral de significancia (10) + estructura (T, F, B)

```

THRESHOLDS = [0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]

# f_max calculado UNA SOLA VEZ aquí - se pasa a Fase 1 y Fase 2
f_max = sae_features_train.max(dim=0)[0] # (n_features,)
print(f"f_max calculado: {(f_max > 0).sum().item()} features activas de {f_max.shape[0]}")

start_time = time.time()

hp_mask_TFB, active_feature_indices = identify_high_precision_features_gpu(
    sae_features_train,
    bsp_gt_train_pieces,
    f_max,
    precision_threshold=0.95,
    significance_threshold=1,
    thresholds=THRESHOLDS,
    feature_batch_size=512,
    verbose=True
)

phase1_time = time.time() - start_time

print(f"\nTiempo Fase 1: {phase1_time:.2f} seg")
print(f"hp_mask_TFB shape: {hp_mask_TFB.shape} (T, n_active, n_bsps)")

```

```

f_max calculado: 8191 features activas de 8192
=====
FASE 1: IDENTIFICAR FEATURES DE ALTA PRECISIÓN
=====
Precision threshold:      0.95
Significance threshold:   >= 1 activaciones
Features activas:        8191/8192 (100.0%)
BSPs:                    128
Thresholds:               [0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
=====

```

Fase 1: 100%| | 16/16 [00:00<00:00, 1596.08it/s]

Pares (feature, BSP) de alta precisión por threshold:

```

t=0.0: 1313 pares
t=0.1: 1335 pares
t=0.2: 2029 pares
t=0.3: 4233 pares
t=0.4: 6566 pares
t=0.5: 9438 pares
t=0.6: 13898 pares
t=0.7: 24005 pares
t=0.8: 50511 pares
t=0.9: 111362 pares ← más pares

```

Features HP por BSP en t=0.9:

```

Media:      870.0
Máx:       3195
BSPs sin features: 0

```

Tiempo Fase 1: 0.89 seg

hp\_mask\_TFB shape: torch.Size([10, 8191, 128]) (T, n\_active, n\_bsps)

## 0.7 7. Fase 2: Reconstruir Tableros en Test Set (GPU)

```
def reconstruct_boards_gpu(
    sae_features_test,
    bsp_gt_test,
    hp_mask_TFB,
    active_feature_indices,
    f_max,
    thresholds=[0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9],
    verbose=True
):
    """
    Fase 2: barrido simultáneo de thresholds T → toma el mejor F1.

    Para cada threshold t:
    1. Binarizar features de test: active[p, f] = acts[p, f] > t * f_max[f]
    2. Predicción por (posición, BSP):
        predict[p, b] = OR_f ( active[p,f] AND hp_mask[t, f, b] )
        implementado como: (active_PF @ hp_FB) > 0
    3. F1 global micro-averaged (acumula TP/FP/FN sobre todas las posiciones y BSPs)

    Retorna el max F1 sobre todos los thresholds (Eq. 5 del paper: max_t).

    Args:
        sae_features_test:      Tensor (P, n_features) GPU
        bsp_gt_test:            Tensor (P, n_bsps) GPU, bool
        hp_mask_TFB:            BoolTensor (T, n_active, n_bsps) de Fase 1
        active_feature_indices: LongTensor (n_active,) de Fase 1
        f_max:                  Tensor (n_features,) GPU - MISMO objeto que Fase 1
        thresholds:             lista de fracciones (debe coincidir con Fase 1)

    Returns:
        best_f1:                float, max F1 sobre thresholds
        best_threshold:         float, threshold óptimo
        f1_per_threshold:       np.array (T,), F1 por threshold
    """
    device = sae_features_test.device
    n_thresholds = len(thresholds)
    thresholds_T = torch.tensor(thresholds, device=device)

    # Activaciones del test solo para features activas
    test_active_PF = sae_features_test[:, active_feature_indices] # (P, n_active)
    active_f_max = f_max[active_feature_indices] # (n_active,)

    gt_PB = bsp_gt_test.bool() # (P, B)
    eps = 1e-8

    f1_scores = torch.zeros(n_thresholds, device=device)
    precision_T = torch.zeros(n_thresholds, device=device)
    recall_T = torch.zeros(n_thresholds, device=device)

    if verbose:
        print("=" * 60)
        print("FASE 2: RECONSTRUCCIÓN - BARRIDO DE THRESHOLDS")
        print("=" * 60)
        print(f"{'t':>5} {'F1':>8} {'P':>8} {'R':>8} {'TP':>8} {'FP':>8} {'FN':>8}")
        print("-" * 60)
```

```

for t_idx in range(n_thresholds):
    t = thresholds[t_idx]

    # Threshold absoluto por feature: t * f_max[f]
    abs_thresh_F = thresholds_T[t_idx] * active_f_max # (n_active,)

    # Binarizar features en test
    active_PF = test_active_PF > abs_thresh_F.unsqueeze(0) # (P, n_active)

    # predict[p, b] = any( active[p,f] AND hp[t,f,b] )
    # = (active_PF.float() @ hp_mask[t].float()) > 0
    hp_FB = hp_mask_TFB[t_idx].float() # (n_active, B)
    predictions_PB = (active_PF.float() @ hp_FB) > 0 # (P, B)

    # Métricas globales micro-averaged
    tp = (predictions_PB & gt_PB).sum().float()
    fp = (predictions_PB & ~gt_PB).sum().float()
    fn = (~predictions_PB & gt_PB).sum().float()

    p = tp / (tp + fp + eps)
    r = tp / (tp + fn + eps)
    f1 = 2 * p * r / (p + r + eps)

    f1_scores[t_idx] = f1
    precision_T[t_idx] = p
    recall_T[t_idx] = r

    if verbose:
        print(f"t:>5.1f} {f1.item():>8.4f} {p.item():>8.4f} {r.item():>8.4f}"
              f" {tp.item():>8.0f} {fp.item():>8.0f} {fn.item():>8.0f}")

    best_idx = f1_scores.argmax().item()
    best_f1 = f1_scores[best_idx].item()
    best_threshold = thresholds[best_idx]

    if verbose:
        print("=" * 60)
        print(f"Mejor threshold: t={best_threshold:.1f}")
        print(f"Reconstruction Score (max_t F1): {best_f1:.4f}")
        print("=" * 60)

    return best_f1, best_threshold, f1_scores.cpu().numpy()

print(" Fase 2 rediseñada: barrido T simultáneo + F1 global micro-averaged + max_t")

```

Fase 2 rediseñada: barrido T simultáneo + F1 global micro-averaged + max\_t

```

start_time = time.time()

reconstruction_score, best_threshold, f1_per_threshold = reconstruct_boards_gpu(
    sae_features_test,
    bsp_gt_test_pieces,
    hp_mask_TFB,
    active_feature_indices,
    f_max, # mismo tensor calculado antes de Fase 1
    thresholds=THRESHOLDS,
    verbose=True
)

phase2_time = time.time() - start_time
total_time = phase1_time + phase2_time

```

```
print(f"\nTiempo Fase 2: {phase2_time:.2f} seg")
print(f"Tiempo total: {total_time:.2f} seg ({total_time/60:.2f} min)")
```

#### FASE 2: RECONSTRUCCIÓN - BARRIDO DE THRESHOLDS

| t   | F1     | P      | R      | TP     | FP    | FN      |
|-----|--------|--------|--------|--------|-------|---------|
| 0.0 | 0.0175 | 0.8974 | 0.0088 | 17704  | 2025  | 1988296 |
| 0.1 | 0.0193 | 0.8973 | 0.0098 | 19575  | 2241  | 1986425 |
| 0.2 | 0.0289 | 0.9145 | 0.0147 | 29447  | 2753  | 1976553 |
| 0.3 | 0.0410 | 0.9198 | 0.0210 | 42042  | 3667  | 1963958 |
| 0.4 | 0.0472 | 0.9214 | 0.0242 | 48541  | 4139  | 1957459 |
| 0.5 | 0.0508 | 0.9007 | 0.0261 | 52440  | 5783  | 1953560 |
| 0.6 | 0.0543 | 0.8475 | 0.0281 | 56281  | 10127 | 1949719 |
| 0.7 | 0.0603 | 0.7500 | 0.0314 | 62985  | 21000 | 1943015 |
| 0.8 | 0.0722 | 0.6197 | 0.0383 | 76928  | 47201 | 1929072 |
| 0.9 | 0.0922 | 0.5189 | 0.0506 | 101465 | 94070 | 1904535 |

Mejor threshold: t=0.9

Reconstruction Score (max\_t F1): 0.0922

Tiempo Fase 2: 11.11 seg

Tiempo total: 12.01 seg (0.20 min)

## 0.8 8. Análisis de Resultados

```
print("=" * 60)
print("BARRIDO DE THRESHOLDS - F1 POR THRESHOLD")
print("=" * 60)
for i, (t, f1) in enumerate(zip(THRESHOLDS, f1_per_threshold)):
    marker = " ← MEJOR" if i == f1_per_threshold.argmax() else ""
    print(f"t={t:.1f}: F1={f1:.4f}{marker}")
print("=" * 60)

plt.figure(figsize=(10, 6))
plt.plot(THRESHOLDS, f1_per_threshold, marker='o', linewidth=2, markersize=8)
plt.axhline(y=reconstruction_score, color='r', linestyle='--',
            label=f'Best F1={reconstruction_score:.4f} (t={best_threshold:.1f})')
plt.axhline(y=0.95, color='g', linestyle=':', label='Target Othello SAE (paper)=0.95')
plt.xlabel('Threshold t', fontsize=12)
plt.ylabel('F1 Score (micro-averaged global)', fontsize=12)
plt.title('Board Reconstruction: F1 por Threshold', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

print(f"\n Threshold óptimo: t={best_threshold:.1f}")
print(f" Reconstruction Score: {reconstruction_score:.4f}")
```


#### BARRIDO DE THRESHOLDS - F1 POR THRESHOLD

```
t=0.0: F1=0.0175
t=0.1: F1=0.0193
t=0.2: F1=0.0289
t=0.3: F1=0.0410
```

```

t=0.4: F1=0.0472
t=0.5: F1=0.0508
t=0.6: F1=0.0543
t=0.7: F1=0.0603
t=0.8: F1=0.0722
t=0.9: F1=0.0922 ← MEJOR
=====

```



```

Threshold óptimo: t=0.9
Reconstruction Score: 0.0922

```

```

# Estadísticas sobre features de alta precisión por BSP
hp_per_bsp = hp_mask_TFB.sum(dim=1) # (T, n_bsps)
best_t_idx = torch.tensor(f1_per_threshold).argmax().item()
n_features_per_bsp = hp_per_bsp[best_t_idx].cpu().numpy()

print("="*60)
print("ESTADÍSTICAS - FEATURES POR BSP")
print("="*60)
print(f"Total BSPs: {len(n_features_per_bsp)}")
print(f"Features promedio por BSP: {np.mean(n_features_per_bsp):.1f}")
print(f"Features mediana por BSP: {np.median(n_features_per_bsp):.1f}")
print(f"BSPs con 0 features: {np.sum(n_features_per_bsp == 0)}")
print(f"BSPs con 1+ features: {np.sum(n_features_per_bsp > 0)}")
print(f"BSPs con 5+ features: {np.sum(n_features_per_bsp >= 5)}")
print(f"BSPs con 10+ features: {np.sum(n_features_per_bsp >= 10)}")
print(f"Max features en una BSP: {np.max(n_features_per_bsp)}")
print("="*60)

```

```

=====
ESTADÍSTICAS - FEATURES POR BSP
=====

```

```

Total BSPs: 128
Features promedio por BSP: 870.0
Features mediana por BSP: 703.0
BSPs con 0 features: 0
BSPs con 1+ features: 128
BSPs con 5+ features: 128
BSPs con 10+ features: 128
Max features en una BSP: 3195
=====

```

## 0.9 9. Guardar Resultados

```

results = {
    'reconstruction_score': reconstruction_score,
    'best_threshold': best_threshold,
    'f1_per_threshold': f1_per_threshold,
    'thresholds': np.array(THRESHOLDS),
    'phase1_time_seconds': phase1_time,
    'phase2_time_seconds': phase2_time,
    'total_time_seconds': total_time,
    'bsp_names': bsp_pieces_names,
}

metrics_dir = get_metrics_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

output_path = metrics_dir / f"reconstruction_results_{filter_suffix}.npz"
np.savez(output_path, **results)

print(f" Resultados guardados en: {output_path}")
print(f"\nContenido:")
print(f" reconstruction_score: {reconstruction_score:.4f}")
print(f" best_threshold: {best_threshold:.1f}")
print(f" f1_per_threshold: {f1_per_threshold}")

# Comparación con el paper
print()
print("=" * 60)
print("COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)")
print("=" * 60)
print(f"{'Modelo':<25} {'Reconstruction':>14}")
print("-" * 40)
print(f"{'SAE trained (paper)':<25} {'0.95':>14} ← Objetivo Othello")
print(f"{'Nuestro SAE':<25} {reconstruction_score:>14.4f} ← Resultado")
print("=" * 60)

if reconstruction_score >= 0.90:
    print("\n Excelente: Muy cercano al objetivo del paper")
elif reconstruction_score >= 0.50:
    print("\n~ Moderado: Por encima de random, revisar SAE")
else:
    print("\n Bajo: Revisar entrenamiento del SAE o calidad de las activaciones")

```

Resultados guardados en: c:\Users\Esposa\Documents\Repos\othello\_world\sae\experiments\layer\_06\sae\_matryoshkabatchto



Contenido:

```
reconstruction_score: 0.0922
best_threshold:      0.9
f1_per_threshold:    [0.01747914 0.01930649 0.0288951  0.04098242 0.0471574  0.05080846 0.05431459 0.06027316 0.0722
0.09217659]
```

=====

COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)

=====

| Modelo              | Reconstruction          |
|---------------------|-------------------------|
| SAE trained (paper) | 0.95 ← Objetivo Othello |
| Nuestro SAE         | 0.0922 ← Resultado      |

=====

Bajo: Revisar entrenamiento del SAE o calidad de las activaciones

## 0.10 10. Generar PDF con Quarto

```
output_dir = get_metrics_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

pdf_name = get_report_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    f'reconstruction_{filter_suffix}',
    extras_id=extras_id
)

notebook_name = 'reconstruction_sae_matryoska.ipynb'
output_path = output_dir / pdf_name

print(f' Generando PDF con Quarto...')
print(f' Archivo de salida: {output_path}')
```

```
result = subprocess.run(
    f'quarto render {notebook_name} --to pdf --output-dir "{output_dir}" --output {pdf_name}',
    shell=True,
    capture_output=True,
    text=True
)

if result.returncode == 0:
    print(f' PDF generado exitosamente: {output_path}')
else:
    print(f' Error al generar PDF:')
    print(result.stderr)
```

Generando PDF con Quarto...

Archivo de salida: c:\Users\Esposa\Documents\Repos\othello\_world\sae\experiments\layer\_06\sae\_matryoshkabatchtopk\sae

PDF generado exitosamente: c:\Users\Esposa\Documents\Repos\othello\_world\sae\experiments\layer\_06\sae\_matryoshkabatch