

1. Setup e Imports

```
import json as _json
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import os
import sys
from pathlib import Path

# Agregar path del proyecto para imports
sys.path.append(os.path.abspath('../..'))
from sae.models.sae import SparseAutoencoder
from sae.tools.naming_utils import save_checkpoint, get_model_name, get_extras_id, get_checkpoint_dir, get_model_dir
from sae.tools.experiment_utils import get_experiment_dir, get_plots_dir, get_metrics_dir, get_logs_dir

# Configuración de visualización para Quarto/PDF
pd.set_option('display.width', 100)
pd.set_option('display.max_columns', None)
np.set_printoptions(linewidth=100, precision=4)
os.environ['COLUMNS'] = '100'

# Configuración de reproducibilidad
torch.manual_seed(42)
np.random.seed(42)

# Device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
# Print config con formato profesional
config_df = pd.DataFrame({
    'Configuración': ['Device', 'Seed', 'Output Width'],
    'Valor': [str(device), '42', '100']
})

print('\n' + '='*50)
print(' CONFIGURACIÓN INICIAL')
print('='*50)
print(config_df.to_string(index=False))
print('='*50)
```

```
=====
CONFIGURACIÓN INICIAL
=====
Configuración Valor
      Device  cuda
      Seed    42
Output Width  100
=====
```

2 Metadata para Naming de Checkpoints

```
# Metadata para naming de checkpoints (separada de hiperparámetros)
NAMING_META = {
    'variante': 'standard',          # Arquitectura del SAE
    'tecnica': 'l1-decay',           # L1 con penalty annealing
    'capa': 6,                       # Layer del modelo OthelloGPT
    'juegos': 500,                   # Número de partidas
    'base_path': 'C:\\Users\\Esposa\\Documents\\Repos\\sae-othello-gpt', # Path base para cl
    'experiment_base_path': os.path.abspath('../../experiments') ,      # sae/experiments,
    'extras': {
        'expansion': 32,
        'sparsity_coef': 5e-3,
        'epochs': 100,
        'patience': 10,
    }
}
```

```
print('\n Metadata de naming:')
for key, value in NAMING_META.items():
    print(f' {key}: {value}')
```

```
Metadata de naming:
variante: standard
tecnic: l1-decay
capa: 6
juegos: 500
base_path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt
experiment_base_path: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments
extras: {'expansion': 32, 'sparsity_coef': 0.005, 'epochs': 100, 'patience': 10}
```

```
extras_id = get_extras_id(
    get_checkpoint_dir(
        NAMING_META['base_path'],
        NAMING_META['variante'],
        NAMING_META['tecnic'],
        NAMING_META['capa'],
        NAMING_META['juegos']
    ),
    NAMING_META['extras']
) if NAMING_META.get('extras') else None

print(f'extras_id: {extras_id}')
```

Nuevo extras registrado: ex1 → {'expansion': 32, 'sparsity_coef': 0.005, 'epochs': 100, 'patience': 10}

extras_id: ex1

3. Configuración del SAE

Las activaciones ya fueron extraídas previamente con `sae/activations/run_extraction.py`. Aquí especificamos qué archivo cargar y los hiperparámetros del SAE.

```
# Configuración del SAE
CONFIG = {
    # Arquitectura
    'input_dim': 512,
```

```

'hidden_dim': 16384,
'tied_weights': False,

# Entrenamiento
'batch_size': 128,
'num_epochs': 100,
'learning_rate': 1e-3,
'sparsity_coef': 5e-3,
'sparsity_decay': 0.95,

# Early Stopping
'early_stopping': True,
'patience': 10,

# Datos
'activations_file': Path(f"../../activations/data/layer{NAMING_META['capa']}-{NAMING_I
'save_dir': Path('./saved_model'),
'plots_dir': get_plots_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)
}

# Crear directorio local del modelo (plots_dir ya fue creado por get_plots_dir)
CONFIG['save_dir'].mkdir(parents=True, exist_ok=True)

# Mostrar configuración con formato profesional
config_data = {
    'Categoría': [],
    'Parámetro': [],
    'Valor': []
}

# Arquitectura
for key in ['input_dim', 'hidden_dim', 'tied_weights']:
    config_data['Categoría'].append('Arquitectura')
    config_data['Parámetro'].append(key)

```

```

        config_data['Valor'].append(str(CONFIG[key]))

# Entrenamiento
for key in ['batch_size', 'num_epochs', 'learning_rate', 'sparsity_coef', 'sparsity_decay']:
    config_data['Categoría'].append('Entrenamiento')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

# Early Stopping
for key in ['early_stopping', 'patience']:
    config_data['Categoría'].append('Early Stopping')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

config_table = pd.DataFrame(config_data)

print('\n' + '='*70)
print(' CONFIGURACIÓN DEL SAE')
print('='*70)
print(config_table.to_string(index=False))
print('='*70)

print(f'\nArchivo de activaciones: {CONFIG["activations_file"]}')
print(f'Expansión: {CONFIG["hidden_dim"]/CONFIG["input_dim"]:.1f}x')
print(f'Plots dir: {CONFIG["plots_dir"]}')
print('='*70)

```

```

=====
CONFIGURACIÓN DEL SAE
=====

```

Categoría	Parámetro	Valor
Arquitectura	input_dim	512
Arquitectura	hidden_dim	16384
Arquitectura	tied_weights	False
Entrenamiento	batch_size	128
Entrenamiento	num_epochs	100
Entrenamiento	learning_rate	0.001
Entrenamiento	sparsity_coef	0.005
Entrenamiento	sparsity_decay	0.95
Early Stopping	early_stopping	True
Early Stopping	patience	10

```

=====
Archivo de activaciones: ..\..\..\activations\data\layer6_500games.npy
Expansión: 32.0x
Plots dir: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\decay_standard\500games\ex1\plots
=====

```

4. Dataset de Activaciones (Train/Validation Split 80-20%)

```

class ActivationsDataset(Dataset):
    """Dataset para cargar activaciones extraídas"""

    def __init__(self, activations_path):
        self.activations = np.load(activations_path)

    def __len__(self):
        return len(self.activations)

    def __getitem__(self, idx):
        return torch.tensor(self.activations[idx], dtype=torch.float32)

# Cargar dataset completo
full_dataset = ActivationsDataset(CONFIG['activations_file'])

# Información del dataset con formato profesional
dataset_info = pd.DataFrame({
    'Propiedad': ['Total muestras', 'Shape', 'Dtype', 'Tamaño en memoria'],
    'Valor': [
        f"{len(full_dataset):,}",
        str(full_dataset.activations.shape),
        str(full_dataset.activations.dtype),
        f"{full_dataset.activations.nbytes / 1024**2:.2f} MB"
    ]
})

print('\n' + '='*50)
print(' DATASET DE ACTIVACIONES')
print('='*50)
print(dataset_info.to_string(index=False))
print('='*50)

```

```

# Split train/validation (80-20)
train_size = int(0.8 * len(full_dataset))
val_size = len(full_dataset) - train_size
train_dataset, val_dataset = torch.utils.data.random_split(
    full_dataset,
    [train_size, val_size],
    generator=torch.Generator().manual_seed(42)
)

# Resumen del split con formato profesional
split_info = pd.DataFrame({
    'Conjunto': ['Train', 'Validation', 'Total'],
    'Muestras': [
        f"{len(train_dataset):,}",
        f"{len(val_dataset):,}",
        f"{len(full_dataset):,}"
    ],
    'Porcentaje': [
        f"{len(train_dataset)/len(full_dataset):.1%}",
        f"{len(val_dataset)/len(full_dataset):.1%}",
        "100.0%"
    ]
})
print('\n' + '='*50)
print(' SPLIT DE DATOS (80-20) ')
print('='*50)
print(split_info.to_string(index=False))
print('='*50)

# Crear DataLoaders
train_loader = DataLoader(
    train_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=True,
    num_workers=0,
    pin_memory=(device.type == 'cuda')
)

val_loader = DataLoader(
    val_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=False, # No shuffle para validación

```

```

    num_workers=0,
    pin_memory=(device.type == 'cuda')
)

# Resumen de DataLoaders
loader_info = pd.DataFrame({
    'DataLoader': ['Train', 'Validation'],
    'Batches': [len(train_loader), len(val_loader)],
    'Batch Size': [CONFIG['batch_size'], CONFIG['batch_size']],
    'Shuffle': ['Sí', 'No']
})

print('\n' + '='*50)
print(' DATALOADERS')
print('='*50)
print(loader_info.to_string(index=False))
print('='*50)

```

```

=====
DATASET DE ACTIVACIONES
=====

      Propiedad      Valor
Total muestras      29,500
      Shape (29500, 512)
      Dtype      float32
Tamaño en memoria      57.62 MB
=====

=====
SPLIT DE DATOS (80-20)
=====

Conjunto Muestras Porcentaje
Train      23,600      80.0%
Validation      5,900      20.0%
Total      29,500      100.0%
=====

=====
DATALOADERS
=====
DataLoader  Batches  Batch Size  Shuffle

```


Train	185	128	Sí
Validation	47	128	No
=====			

5. Modelo: Sparse Autoencoder

Modelo importado desde `sae.models.sae.SparseAutoencoder`

Arquitectura:

- **Encoder:** Linear $\rightarrow$ ReLU (sparsity natural)- **Decoder:** Linear (reconstrucción)

```
# Crear modelo desde el módulo
model = SparseAutoencoder(
    input_dim=CONFIG['input_dim'],
    hidden_dim=CONFIG['hidden_dim'],
    tied_weights=CONFIG['tied_weights']
).to(device)

# Calcular parámetros
num_params = sum(p.numel() for p in model.parameters())
encoder_params = sum(p.numel() for p in model.encoder.parameters())
if CONFIG['tied_weights']:
    decoder_params = CONFIG['input_dim'] # Solo bias
else:
    decoder_params = sum(p.numel() for p in model.decoder.parameters())

# Resumen del modelo con formato profesional
model_info = pd.DataFrame({
    'Componente': ['Input', 'Encoder', 'Hidden (Latent)', 'Decoder', 'Output', 'Total'],
    'Dimensión': [
        CONFIG['input_dim'],
        f"{CONFIG['input_dim']}  $\rightarrow$  {CONFIG['hidden_dim']}",
        CONFIG['hidden_dim'],
        f"{CONFIG['hidden_dim']}  $\rightarrow$  {CONFIG['input_dim']}",
        CONFIG['input_dim'],
        '-',
    ],
    'Parámetros': [
        '-',
        f"{encoder_params:,}",
        '-',
        f"{decoder_params:,}",
    ]
})
```

```

        '-',
        f"{num_params: ,}"
    ]
})

print('\n' + '='*70)
print(' ARQUITECTURA DEL MODELO')
print('='*70)
print(model_info.to_string(index=False))
print('='*70)
print(f'\n Factor de expansión: {CONFIG["hidden_dim"]/CONFIG["input_dim"]:.1f}x')
print(f' Tied weights: {"Sí" if CONFIG["tied_weights"] else "No"}')
print(f' Device: {device}')
print('='*70)

```

```

=====
ARQUITECTURA DEL MODELO
=====

```

	Componente	Dimensión	Parámetros
	Input	512	-
	Encoder	512 → 16384	8,404,992
Hidden	(Latent)	16384	-
	Decoder	16384 → 512	8,389,120
	Output	512	-
	Total	-	16,794,112

```

=====

Factor de expansión: 32.0x
Tied weights: No
Device: cuda
=====

```

6. Función de Pérdida

$\text{Loss} = \text{MSE}(\text{reconstrucción}) + \lambda \times \text{L1}(\text{activaciones})$

- **MSE**: Mide qué tan bien reconstruimos las activaciones originales
- **L1**: Penaliza activaciones grandes → fuerza sparsity

```

def loss_function(x, reconstruction, hidden, sparsity_coef):
    """
    Pérdida compuesta: MSE + L1 penalty

    Args:
        x: Activaciones originales
        reconstruction: Activaciones reconstruidas
        hidden: Activaciones latentes (sparse)
        sparsity_coef: Peso de la penalización L1
    """
    # Error de reconstrucción
    mse_loss = F.mse_loss(reconstruction, x)

    # Sumar sobre las features (dim=1), promediar sobre el batch
    sparsity_loss = hidden.abs().sum(dim=1).mean()

    # Pérdida total
    total_loss = mse_loss + sparsity_coef * sparsity_loss

    return total_loss, mse_loss, sparsity_loss

# Optimizador
optimizer = torch.optim.Adam(model.parameters(), lr=CONFIG['learning_rate'])

# Resumen del optimizador
optim_info = pd.DataFrame({
    'Parámetro': ['Optimizador', 'Learning Rate', 'Sparsity Coef (inicial)', 'Sparsity Decay'],
    'Valor': [
        'Adam',
        f"{CONFIG['learning_rate']:.0e}",
        f"{CONFIG['sparsity_coef']:.0e}",
        f"{CONFIG['sparsity_decay']:.2f}"
    ]
})

print('\n' + '='*50)
print(' OPTIMIZACIÓN')
print('='*50)
print(optim_info.to_string(index=False))
print('='*50)

```

```
=====
OPTIMIZACIÓN
=====
                Parámetro Valor
            Optimizador  Adam
            Learning Rate 1e-03
Sparsity Coef (inicial) 5e-03
            Sparsity Decay 0.95
=====
```

7. Loop de Entrenamiento

```
# Historial de entrenamiento
history = {
    'train_total_loss': [],
    'train_mse_loss': [],
    'train_sparsity_loss': [],
    'train_l0_sparsity': [],
    'train_fraction_active': [],
    'val_total_loss': [],
    'val_mse_loss': [],
    'val_sparsity_loss': [],
    'val_l0_sparsity': [],
    'val_fraction_active': [],
    'sparsity_coef_history': [] # Guardar evolución del coeficiente L1
}

# Early stopping variables
best_val_mse = float('inf')
best_epoch = 0
epochs_no_improve = 0

print(f'\n Iniciando entrenamiento: {CONFIG["num_epochs"]} épocas máximo')
print(f' L1 decay: {CONFIG["sparsity_coef"]:.2e} × {CONFIG["sparsity_decay"]:.2f}^época')
print(f' Early stopping: patience={CONFIG["patience"]} épocas')
print('=' * 70)

for epoch in range(CONFIG['num_epochs']):
    # ===== L1 DECAY =====
    # Reducir sparsity_coef exponencialmente
    current_sparsity_coef = CONFIG['sparsity_coef'] * (CONFIG['sparsity_decay'] ** epoch)
```

```

history['sparsity_coef_history'].append(current_sparsity_coef)

# ===== ENTRENAMIENTO =====
model.train()

epoch_total_loss = 0
epoch_mse_loss = 0
epoch_sparsity_loss = 0
epoch_l0 = 0
epoch_fraction_active = 0

# Barra de progreso
pbar = tqdm(train_loader, desc=f'Época {epoch+1}/{CONFIG["num_epochs"]} [Train] L1={curr

for batch in pbar:
    # Mover al device
    x = batch.to(device)

    # Forward pass
    reconstruction, hidden = model(x)

    # Calcular pérdida con sparsity_coef actualizado
    total_loss, mse_loss, sparsity_loss = loss_function(
        x, reconstruction, hidden, current_sparsity_coef
    )

    # Backward pass
    optimizer.zero_grad()
    total_loss.backward()

    optimizer.step()

    # Métricas
    with torch.no_grad():
        l0 = (hidden > 0).float().sum(dim=1).mean().item()
        fraction_active = (hidden > 0).float().mean().item()

    # Acumular
    epoch_total_loss += total_loss.item()
    epoch_mse_loss += mse_loss.item()
    epoch_sparsity_loss += sparsity_loss.item()
    epoch_l0 += l0

```

```

epoch_fraction_active += fraction_active

# Actualizar barra
pbar.set_postfix({
    'loss': f'{total_loss.item():.4f}',
    'mse': f'{mse_loss.item():.4f}',
    'l0': f'{l0:.1f}'
})

# Promedios de entrenamiento
num_train_batches = len(train_loader)
avg_train_total_loss = epoch_total_loss / num_train_batches
avg_train_mse_loss = epoch_mse_loss / num_train_batches
avg_train_sparsity_loss = epoch_sparsity_loss / num_train_batches
avg_train_l0 = epoch_l0 / num_train_batches
avg_train_fraction_active = epoch_fraction_active / num_train_batches

# ===== VALIDACIÓN =====
model.eval()

val_total_loss = 0
val_mse_loss = 0
val_sparsity_loss = 0
val_l0 = 0
val_fraction_active = 0

with torch.no_grad():
    for batch in val_loader:
        x = batch.to(device)

        # Forward pass
        reconstruction, hidden = model(x)

        # Calcular pérdida con sparsity_coef actualizado
        total_loss, mse_loss, sparsity_loss = loss_function(
            x, reconstruction, hidden, current_sparsity_coef
        )

        # Métricas
        l0 = (hidden > 0).float().sum(dim=1).mean().item()
        fraction_active = (hidden > 0).float().mean().item()

```

```

        # Acumular
        val_total_loss += total_loss.item()
        val_mse_loss += mse_loss.item()
        val_sparsity_loss += sparsity_loss.item()
        val_l0 += l0
        val_fraction_active += fraction_active

# Promedios de validación
num_val_batches = len(val_loader)
avg_val_total_loss = val_total_loss / num_val_batches
avg_val_mse_loss = val_mse_loss / num_val_batches
avg_val_sparsity_loss = val_sparsity_loss / num_val_batches
avg_val_l0 = val_l0 / num_val_batches
avg_val_fraction_active = val_fraction_active / num_val_batches

# Guardar en historial
history['train_total_loss'].append(avg_train_total_loss)
history['train_mse_loss'].append(avg_train_mse_loss)
history['train_sparsity_loss'].append(avg_train_sparsity_loss)
history['train_l0_sparsity'].append(avg_train_l0)
history['train_fraction_active'].append(avg_train_fraction_active)

history['val_total_loss'].append(avg_val_total_loss)
history['val_mse_loss'].append(avg_val_mse_loss)
history['val_sparsity_loss'].append(avg_val_sparsity_loss)
history['val_l0_sparsity'].append(avg_val_l0)
history['val_fraction_active'].append(avg_val_fraction_active)

# ===== EARLY STOPPING =====
if CONFIG['early_stopping']:
    if avg_val_mse_loss < best_val_mse:
        best_val_mse = avg_val_mse_loss
        best_epoch = epoch + 1
        epochs_no_improve = 0

    # Guardar mejor modelo (local + organizado)
    best_checkpoint = {
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'config': CONFIG,
        'history': history,
        'epoch': best_epoch,

```

```

        'val_mse': best_val_mse
    }

    # Guardado local
    best_model_name = get_model_name(
        variante=NAMING_META['variante'],
        tecnica=NAMING_META['tecnica'],
        capa=NAMING_META['capa'],
        juegos=NAMING_META['juegos'],
        estado='best'
    )
    torch.save(best_checkpoint, CONFIG['save_dir'] / best_model_name)

    # Guardado organizado
    save_checkpoint(
        model,
        base_path=NAMING_META['base_path'],
        variante=NAMING_META['variante'],
        tecnica=NAMING_META['tecnica'],
        capa=NAMING_META['capa'],
        juegos=NAMING_META['juegos'],
        estado='best',
        extras=NAMING_META['extras'],
        config=CONFIG,
        history=history,
        optimizer=optimizer
    )

    print(f' Mejor modelo guardado en época {best_epoch} (Val MSE: {best_val_mse:.4f})')
else:
    epochs_no_improve += 1
    if epochs_no_improve >= CONFIG['patience']:
        print(f'\n Early stopping activado en época {epoch+1}')
        print(f' Mejor modelo: época {best_epoch} con Val MSE = {best_val_mse:.4f}')
        break

# Log cada 5 épocas
if (epoch + 1) % 5 == 0 or epoch == 0:
    print(f'\n Época {epoch+1}/{CONFIG["num_epochs"]} | L1 coef: {current_sparsity_coef:.4f}')
    print(f' Train - Total: {avg_train_total_loss:.4f} | MSE: {avg_train_mse_loss:.4f}')
    print(f' Val - Total: {avg_val_total_loss:.4f} | MSE: {avg_val_mse_loss:.4f} | L0')
    if CONFIG['early_stopping']:

```



```

        print(f'    Best Val MSE: {best_val_mse:.4f} @ época {best_epoch} | Sin mejora: {best_val_mse:.4f}')

print('\n' + '='*70)
print(' RESUMEN DEL ENTRENAMIENTO')
print('='*70)

if CONFIG['early_stopping']:
    print(f'Mejor modelo en época {best_epoch} con Val MSE = {best_val_mse:.4f}')
    print(f'Early stopping tras {len(history["train_total_loss"])} épocas')
else:
    print(f'Entrenamiento completo: {CONFIG["num_epochs"]} épocas')

# Tabla de métricas finales
final_metrics = pd.DataFrame({
    'Métrica': ['Total Loss', 'MSE Loss', 'L0 Sparsity', 'Fraction Active'],
    'Train': [
        f"{history['train_total_loss'][-1]:.4f}",
        f"{history['train_mse_loss'][-1]:.4f}",
        f"{history['train_l0_sparsity'][-1]:.1f}",
        f"{history['train_fraction_active'][-1]:.2%}"
    ],
    'Validation': [
        f"{history['val_total_loss'][-1]:.4f}",
        f"{history['val_mse_loss'][-1]:.4f}",
        f"{history['val_l0_sparsity'][-1]:.1f}",
        f"{history['val_fraction_active'][-1]:.2%}"
    ]
})

print('\nMétricas Finales (Última Época):')
print(final_metrics.to_string(index=False))
print('='*70)

```

Iniciando entrenamiento: 100 épocas máximo
 L1 decay: $5.00e-03 \times 0.95^{\text{época}}$
 Early stopping: patience=10 épocas

=====

Época 1/100 [Train] L1=5.00e-03: 100% | 185/185 [00:01<00:00, 107.97it/s, loss=0.6476,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 1 (Val MSE: 0.5683)

Época 1/100 | L1 coef: 5.00e-03
Train - Total: 0.7895 | MSE: 0.6415 | L0: 169.8
Val - Total: 0.6103 | MSE: 0.5683 | L0: 11.6
Best Val MSE: 0.5683 @ época 1 | Sin mejora: 0/10

Época 2/100 [Train] L1=4.75e-03: 100% | 185/185 [00:01<00:00, 125.48it/s, loss=0.5347,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 2 (Val MSE: 0.5078)

Época 3/100 [Train] L1=4.51e-03: 100% | 185/185 [00:01<00:00, 122.80it/s, loss=0.5325,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 3 (Val MSE: 0.4596)

Época 4/100 [Train] L1=4.29e-03: 100% | 185/185 [00:01<00:00, 123.64it/s, loss=0.5033,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 4 (Val MSE: 0.4117)

Época 5/100 [Train] L1=4.07e-03: 100% | 185/185 [00:01<00:00, 123.94it/s, loss=0.4557,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 5 (Val MSE: 0.3774)

Época 5/100 | L1 coef: 4.07e-03
Train - Total: 0.4594 | MSE: 0.3908 | L0: 37.2
Val - Total: 0.4499 | MSE: 0.3774 | L0: 45.8
Best Val MSE: 0.3774 @ época 5 | Sin mejora: 0/10

Época 6/100 [Train] L1=3.87e-03: 100% | 185/185 [00:01<00:00, 123.73it/s, loss=0.4339,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 6 (Val MSE: 0.3542)

Época 7/100 [Train] L1=3.68e-03: 100% | 185/185 [00:01<00:00, 124.08it/s, loss=0.3979,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 7 (Val MSE: 0.3297)

Época 8/100 [Train] L1=3.49e-03: 100% | 185/185 [00:01<00:00, 124.67it/s, loss=0.3719,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 8 (Val MSE: 0.3105)

Época 9/100 [Train] L1=3.32e-03: 100% | 185/185 [00:01<00:00, 122.60it/s, loss=0.3663,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 9 (Val MSE: 0.2939)

Época 10/100 [Train] L1=3.15e-03: 100% | 185/185 [00:01<00:00, 126.63it/s, loss=0.3611,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 10 (Val MSE: 0.2806)

Época 10/100 | L1 coef: 3.15e-03

Train - Total: 0.3578 | MSE: 0.2833 | L0: 80.8

Val - Total: 0.3562 | MSE: 0.2806 | L0: 88.1

Best Val MSE: 0.2806 @ época 10 | Sin mejora: 0/10

Época 11/100 [Train] L1=2.99e-03: 100% | 185/185 [00:01<00:00, 123.67it/s, loss=0.3384,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 11 (Val MSE: 0.2703)

Época 12/100 [Train] L1=2.84e-03: 100%| | 185/185 [00:01<00:00, 125.36it/s, loss=0.3118,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 12 (Val MSE: 0.2578)

Época 13/100 [Train] L1=2.70e-03: 100%| | 185/185 [00:01<00:00, 125.85it/s, loss=0.3300,

Época 14/100 [Train] L1=2.57e-03: 100%| | 185/185 [00:01<00:00, 122.18it/s, loss=0.3169,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 14 (Val MSE: 0.2479)

Época 15/100 [Train] L1=2.44e-03: 100%| | 185/185 [00:01<00:00, 127.01it/s, loss=0.3210,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 15 (Val MSE: 0.2432)

Época 15/100 | L1 coef: 2.44e-03

Train - Total: 0.3047 | MSE: 0.2412 | L0: 108.8

Val - Total: 0.3071 | MSE: 0.2432 | L0: 114.3

Best Val MSE: 0.2432 @ época 15 | Sin mejora: 0/10

Época 16/100 [Train] L1=2.32e-03: 100%| | 185/185 [00:01<00:00, 122.74it/s, loss=0.2883,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 16 (Val MSE: 0.2415)

Época 17/100 [Train] L1=2.20e-03: 100%| | 185/185 [00:01<00:00, 126.77it/s, loss=0.2973,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 17 (Val MSE: 0.2381)

Época 18/100 [Train] L1=2.09e-03: 100%| | 185/185 [00:01<00:00, 126.28it/s, loss=0.2932,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 18 (Val MSE: 0.2346)

Época 19/100 [Train] L1=1.99e-03: 100%| | 185/185 [00:01<00:00, 123.89it/s, loss=0.2909,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 19 (Val MSE: 0.2321)

Época 20/100 [Train] L1=1.89e-03: 100%| | 185/185 [00:01<00:00, 121.56it/s, loss=0.2755,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 20 (Val MSE: 0.2310)

Época 20/100 | L1 coef: 1.89e-03
Train - Total: 0.2776 | MSE: 0.2273 | L0: 121.5
Val - Total: 0.2817 | MSE: 0.2310 | L0: 126.2
Best Val MSE: 0.2310 @ época 20 | Sin mejora: 0/10

Época 21/100 [Train] L1=1.79e-03: 100%| | 185/185 [00:01<00:00, 124.57it/s, loss=0.2796,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 21 (Val MSE: 0.2308)

Época 22/100 [Train] L1=1.70e-03: 100%| | 185/185 [00:01<00:00, 125.82it/s, loss=0.2637,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 22 (Val MSE: 0.2272)

Época 23/100 [Train] L1=1.62e-03: 100%| | 185/185 [00:01<00:00, 125.54it/s, loss=0.2584,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 23 (Val MSE: 0.2260)

Época 24/100 [Train] L1=1.54e-03: 100%| | 185/185 [00:01<00:00, 125.56it/s, loss=0.2787,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 24 (Val MSE: 0.2236)

Época 25/100 [Train] L1=1.46e-03: 100%| | 185/185 [00:01<00:00, 125.58it/s, loss=0.2623,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 25 (Val MSE: 0.2224)

Época 25/100 | L1 coef: 1.46e-03

Train - Total: 0.2594 | MSE: 0.2194 | L0: 129.3

Val - Total: 0.2628 | MSE: 0.2224 | L0: 134.0

Best Val MSE: 0.2224 @ época 25 | Sin mejora: 0/10

Época 26/100 [Train] L1=1.39e-03: 100%| | 185/185 [00:01<00:00, 126.17it/s, loss=0.2642,

Época 27/100 [Train] L1=1.32e-03: 100%| | 185/185 [00:01<00:00, 122.63it/s, loss=0.2741,

Época 28/100 [Train] L1=1.25e-03: 100%| | 185/185 [00:01<00:00, 125.24it/s, loss=0.2449,

Época 29/100 [Train] L1=1.19e-03: 100%| | 185/185 [00:01<00:00, 124.34it/s, loss=0.2499,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 29 (Val MSE: 0.2194)

Época 30/100 [Train] L1=1.13e-03: 100%| | 185/185 [00:01<00:00, 127.27it/s, loss=0.2519,

Época 30/100 | L1 coef: 1.13e-03

Train - Total: 0.2473 | MSE: 0.2153 | L0: 134.6

Val - Total: 0.2516 | MSE: 0.2196 | L0: 138.3

Best Val MSE: 0.2194 @ época 29 | Sin mejora: 1/10

Época 31/100 [Train] L1=1.07e-03: 100%| | 185/185 [00:01<00:00, 121.45it/s, loss=0.2307,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 31 (Val MSE: 0.2175)

Época 32/100 [Train] L1=1.02e-03: 100%| | 185/185 [00:01<00:00, 125.81it/s, loss=0.2462,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 32 (Val MSE: 0.2174)

Época 33/100 [Train] L1=9.69e-04: 100%| | 185/185 [00:01<00:00, 126.63it/s, loss=0.2452,

Época 34/100 [Train] L1=9.20e-04: 100%| | 185/185 [00:01<00:00, 125.26it/s, loss=0.2394,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 34 (Val MSE: 0.2165)

Época 35/100 [Train] L1=8.74e-04: 100%| | 185/185 [00:01<00:00, 124.79it/s, loss=0.2390,

Época 35/100 | L1 coef: 8.74e-04

Train - Total: 0.2384 | MSE: 0.2128 | L0: 138.2

Val - Total: 0.2429 | MSE: 0.2171 | L0: 142.0

Best Val MSE: 0.2165 @ época 34 | Sin mejora: 1/10

Época 36/100 [Train] L1=8.30e-04: 100%| | 185/185 [00:01<00:00, 122.48it/s, loss=0.2368,

Época 37/100 [Train] L1=7.89e-04: 100%| | 185/185 [00:01<00:00, 127.56it/s, loss=0.2389,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 37 (Val MSE: 0.2139)

Época 38/100 [Train] L1=7.49e-04: 100%| | 185/185 [00:01<00:00, 124.81it/s, loss=0.2397,

Época 39/100 [Train] L1=7.12e-04: 100%| | 185/185 [00:01<00:00, 125.73it/s, loss=0.2317,

Época 40/100 [Train] L1=6.76e-04: 100%| | 185/185 [00:01<00:00, 126.87it/s, loss=0.2319,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 40 (Val MSE: 0.2137)

Época 40/100 | L1 coef: 6.76e-04

Train - Total: 0.2313 | MSE: 0.2107 | L0: 141.4

Val - Total: 0.2344 | MSE: 0.2137 | L0: 145.5

Best Val MSE: 0.2137 @ época 40 | Sin mejora: 0/10

Época 41/100 [Train] L1=6.43e-04: 100%| | 185/185 [00:01<00:00, 122.11it/s, loss=0.2294,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 41 (Val MSE: 0.2135)

Época 42/100 [Train] L1=6.10e-04: 100%| | 185/185 [00:01<00:00, 126.97it/s, loss=0.2265,

Época 43/100 [Train] L1=5.80e-04: 100%| | 185/185 [00:01<00:00, 125.07it/s, loss=0.2262,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 43 (Val MSE: 0.2124)

Época 44/100 [Train] L1=5.51e-04: 100%| | 185/185 [00:01<00:00, 123.29it/s, loss=0.2054,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 44 (Val MSE: 0.2120)

Época 45/100 [Train] L1=5.23e-04: 100%| | 185/185 [00:01<00:00, 122.26it/s, loss=0.2262,

Época 45/100 | L1 coef: 5.23e-04

Train - Total: 0.2259 | MSE: 0.2092 | L0: 143.8

Val - Total: 0.2360 | MSE: 0.2195 | L0: 145.6

Best Val MSE: 0.2120 @ época 44 | Sin mejora: 1/10

Época 46/100 [Train] L1=4.97e-04: 100%| | 185/185 [00:01<00:00, 125.35it/s, loss=0.2130,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 46 (Val MSE: 0.2108)

Época 47/100 [Train] L1=4.72e-04: 100%| | 185/185 [00:01<00:00, 125.27it/s, loss=0.2182,

Época 48/100 [Train] L1=4.49e-04: 100%| | 185/185 [00:01<00:00, 123.39it/s, loss=0.2171,

Época 49/100 [Train] L1=4.26e-04: 100%| | 185/185 [00:01<00:00, 124.61it/s, loss=0.2311,

Época 50/100 [Train] L1=4.05e-04: 100%| | 185/185 [00:01<00:00, 126.63it/s, loss=0.2142,

Época 50/100 | L1 coef: 4.05e-04
Train - Total: 0.2217 | MSE: 0.2083 | L0: 145.4
Val - Total: 0.2273 | MSE: 0.2137 | L0: 149.5
Best Val MSE: 0.2108 @ época 46 | Sin mejora: 4/10

Época 51/100 [Train] L1=3.85e-04: 100%| | 185/185 [00:01<00:00, 125.28it/s, loss=0.2146,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 51 (Val MSE: 0.2092)

Época 52/100 [Train] L1=3.65e-04: 100%| | 185/185 [00:01<00:00, 122.70it/s, loss=0.2146,
Época 53/100 [Train] L1=3.47e-04: 100%| | 185/185 [00:01<00:00, 126.74it/s, loss=0.2039,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt
Mejor modelo guardado en época 53 (Val MSE: 0.2088)

Época 54/100 [Train] L1=3.30e-04: 100%| | 185/185 [00:01<00:00, 97.00it/s, loss=0.2195,
Época 55/100 [Train] L1=3.13e-04: 100%| | 185/185 [00:01<00:00, 121.58it/s, loss=0.2056,

Época 55/100 | L1 coef: 3.13e-04
Train - Total: 0.2188 | MSE: 0.2079 | L0: 146.8
Val - Total: 0.2229 | MSE: 0.2119 | L0: 150.3
Best Val MSE: 0.2088 @ época 53 | Sin mejora: 2/10

Época 56/100 [Train] L1=2.98e-04: 100%| | 185/185 [00:01<00:00, 126.55it/s, loss=0.2197,
Época 57/100 [Train] L1=2.83e-04: 100%| | 185/185 [00:01<00:00, 124.43it/s, loss=0.2053,
Época 58/100 [Train] L1=2.69e-04: 100%| | 185/185 [00:01<00:00, 115.73it/s, loss=0.2261,
Época 59/100 [Train] L1=2.55e-04: 100%| | 185/185 [00:01<00:00, 116.79it/s, loss=0.2016,
Época 60/100 [Train] L1=2.42e-04: 100%| | 185/185 [00:01<00:00, 99.66it/s, loss=0.2178,

Época 60/100 | L1 coef: 2.42e-04
Train - Total: 0.2135 | MSE: 0.2046 | L0: 148.6
Val - Total: 0.2192 | MSE: 0.2103 | L0: 151.6
Best Val MSE: 0.2088 @ época 53 | Sin mejora: 7/10

Época 61/100 [Train] L1=2.30e-04: 100%| | 185/185 [00:01<00:00, 123.17it/s, loss=0.2018,
Época 62/100 [Train] L1=2.19e-04: 100%| | 185/185 [00:01<00:00, 129.79it/s, loss=0.2294,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 62 (Val MSE: 0.2085)

Época 63/100 [Train] L1=2.08e-04: 100%| | 185/185 [00:01<00:00, 126.32it/s, loss=0.2177,
Época 64/100 [Train] L1=1.97e-04: 100%| | 185/185 [00:01<00:00, 126.39it/s, loss=0.2270,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 64 (Val MSE: 0.2079)

Época 65/100 [Train] L1=1.88e-04: 100%| | 185/185 [00:01<00:00, 129.57it/s, loss=0.2246,

Época 65/100 | L1 coef: 1.88e-04

Train - Total: 0.2112 | MSE: 0.2041 | L0: 149.5

Val - Total: 0.2159 | MSE: 0.2087 | L0: 152.9

Best Val MSE: 0.2079 @ época 64 | Sin mejora: 1/10

Época 66/100 [Train] L1=1.78e-04: 100%| | 185/185 [00:01<00:00, 126.66it/s, loss=0.2334,
Época 67/100 [Train] L1=1.69e-04: 100%| | 185/185 [00:01<00:00, 128.94it/s, loss=0.2135,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 67 (Val MSE: 0.2070)

Época 68/100 [Train] L1=1.61e-04: 100%| | 185/185 [00:01<00:00, 128.75it/s, loss=0.2180,
Época 69/100 [Train] L1=1.53e-04: 100%| | 185/185 [00:01<00:00, 126.60it/s, loss=0.2077,
Época 70/100 [Train] L1=1.45e-04: 100%| | 185/185 [00:01<00:00, 126.91it/s, loss=0.2130,

Época 70/100 | L1 coef: 1.45e-04

Train - Total: 0.2102 | MSE: 0.2045 | L0: 150.1

Val - Total: 0.2155 | MSE: 0.2098 | L0: 153.3

Best Val MSE: 0.2070 @ época 67 | Sin mejora: 3/10

Época 71/100 [Train] L1=1.38e-04: 100%| | 185/185 [00:01<00:00, 129.91it/s, loss=0.2188,
Época 72/100 [Train] L1=1.31e-04: 100%| | 185/185 [00:01<00:00, 128.57it/s, loss=0.2164,
Época 73/100 [Train] L1=1.24e-04: 100%| | 185/185 [00:01<00:00, 128.70it/s, loss=0.2082,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 73 (Val MSE: 0.2065)

Época 74/100 [Train] L1=1.18e-04: 100%| | 185/185 [00:01<00:00, 127.10it/s, loss=0.2042,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_best.pt

Mejor modelo guardado en época 74 (Val MSE: 0.2060)

Época 75/100 [Train] L1=1.12e-04: 100%| | 185/185 [00:01<00:00, 129.31it/s, loss=0.2250,

Época 75/100 | L1 coef: 1.12e-04

Train - Total: 0.2072 | MSE: 0.2026 | L0: 151.0

Val - Total: 0.2114 | MSE: 0.2067 | L0: 155.4

Best Val MSE: 0.2060 @ época 74 | Sin mejora: 1/10

Época 76/100 [Train] L1=1.07e-04: 100%| | 185/185 [00:01<00:00, 125.99it/s, loss=0.2067,
Época 77/100 [Train] L1=1.01e-04: 100%| | 185/185 [00:01<00:00, 127.20it/s, loss=0.2110,
Época 78/100 [Train] L1=9.63e-05: 100%| | 185/185 [00:01<00:00, 123.37it/s, loss=0.2130,
Época 79/100 [Train] L1=9.15e-05: 100%| | 185/185 [00:01<00:00, 124.17it/s, loss=0.2144,
Época 80/100 [Train] L1=8.69e-05: 100%| | 185/185 [00:01<00:00, 123.81it/s, loss=0.2077,

Época 80/100 | L1 coef: 8.69e-05

Train - Total: 0.2069 | MSE: 0.2032 | L0: 151.3

Val - Total: 0.2113 | MSE: 0.2076 | L0: 154.9

Best Val MSE: 0.2060 @ época 74 | Sin mejora: 6/10

Época 81/100 [Train] L1=8.26e-05: 100%| | 185/185 [00:01<00:00, 123.68it/s, loss=0.2210,
Época 82/100 [Train] L1=7.84e-05: 100%| | 185/185 [00:01<00:00, 124.33it/s, loss=0.2129,
Época 83/100 [Train] L1=7.45e-05: 100%| | 185/185 [00:01<00:00, 124.64it/s, loss=0.2174,
Época 84/100 [Train] L1=7.08e-05: 100%| | 185/185 [00:01<00:00, 125.80it/s, loss=0.2187,

Early stopping activado en época 84
Mejor modelo: época 74 con Val MSE = 0.2060

```
=====
RESUMEN DEL ENTRENAMIENTO
=====
Mejor modelo en época 74 con Val MSE = 0.2060
Early stopping tras 84 épocas

Métricas Finales (Última Época):
      Métrica  Train Validation
Total Loss 0.2084      0.2163
MSE Loss 0.2053      0.2132
LO Sparsity 151.3      154.6
Fraction Active 0.92%      0.94%
=====
```

8. Guardar Modelo

```
# Guardar checkpoint final (local + organizado)
checkpoint = {
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'config': CONFIG,
    'history': history
}

# Guardado local
final_model_name = get_model_name(
    variante=NAMING_META['variante'],
    tecnica=NAMING_META['tecnica'],
    capa=NAMING_META['capa'],
    juegos=NAMING_META['juegos'],
    estado='final'
)
checkpoint_path = CONFIG['save_dir'] / final_model_name
torch.save(checkpoint, checkpoint_path)

# Guardado organizado
save_checkpoint(
```

```

    model,
    base_path=NAMING_META['base_path'],
    variante=NAMING_META['variante'],
    tecnica=NAMING_META['tecnica'],
    capa=NAMING_META['capa'],
    juegos=NAMING_META['juegos'],
    estado='final',
    extras=NAMING_META['extras'],
    config=CONFIG,
    history=history,
    optimizer=optimizer
)

# Resumen de archivos guardados
saved_files = pd.DataFrame({
    'Tipo': ['Local', 'Organizado'],
    'Ubicación': [
        str(checkpoint_path),
        f"layer_{NAMING_META['capa']:02d}/models/..."
    ]
})

print('\n' + '='*70)
print(' CHECKPOINTS GUARDADOS')
print('='*70)
print(saved_files.to_string(index=False))
print('='*70)

```

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex1_final.pt

```

=====
CHECKPOINTS GUARDADOS
=====

```

Tipo	Ubicación
Local	saved_model\sae_standard_l1-decay_l6_500g_final.pt
Organizado	layer_06/models/...

```

=====

```

9. Métricas básicas del modelo

```
# Guardar metricas de entrenamiento
metrics_dir = get_metrics_dir(
    NAMING_META["experiment_base_path"],
    NAMING_META["variante"],
    NAMING_META["tecnica"],
    NAMING_META["capa"],
    NAMING_META["juegos"],
    extras_id=extras_id
)

training_metrics = {
    # Metricas de calidad
    "val_mse_final": history["val_mse_loss"][-1],
    "val_mse_best": best_val_mse,
    "val_l0_final": history["val_l0_sparsity"][-1],
    "val_fraction_active_final": history["val_fraction_active"][-1],
    # Contexto del entrenamiento
    "best_epoch": best_epoch,
    "total_epochs": len(history["train_total_loss"]),
    "train_mse_final": history["train_mse_loss"][-1],
    "train_l0_final": history["train_l0_sparsity"][-1],
    # Hiperparametros
    "hidden_dim": CONFIG["hidden_dim"],
    "expansion_factor": CONFIG["hidden_dim"] / CONFIG["input_dim"],
    "sparsity_coef": CONFIG["sparsity_coef"],
    "sparsity_decay": CONFIG["sparsity_decay"],
    "learning_rate": CONFIG["learning_rate"],
    "num_games": NAMING_META["juegos"],
}

metrics_path = metrics_dir / "training_metrics.json"
with open(metrics_path, "w") as f:
    _json.dump(training_metrics, f, indent=2)

# Resumen
metrics_df = pd.DataFrame({
    "Metrica": ["Val MSE (final)", "Val MSE (mejor)", "Val L0 (final)", "Fraccion Activa"],
    "Valor": [
        f"{training_metrics['val_mse_final']:.6f}",
        f"{training_metrics['val_mse_best']:.6f}",

```

```

        f"{training_metrics['val_l0_final']:.1f}",
        f"{training_metrics['val_fraction_active_final']:.2%}"
    ]
})
print(" " + "="*50)
print(" METRICAS DE ENTRENAMIENTO GUARDADAS")
print("="*50)
print(metrics_df.to_string(index=False))
print(f"Archivo: {metrics_path}")
print("="*50)

```

```

=====
METRICAS DE ENTRENAMIENTO GUARDADAS
=====

```

	Metrica	Valor
Val MSE (final)	0.213199	
Val MSE (mejor)	0.205962	
Val L0 (final)	154.6	
Fraccion Activa	0.94%	
Archivo:	c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\decay_standard\500games\ex1\metrics\training_metrics.json	

```

=====

```

10. Visualización de Resultados

```

# Gráficas de pérdidas (Train vs Validation) + L1 Decay
fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# Total loss
axes[0, 0].plot(history['train_total_loss'], label='Train', alpha=0.8)
axes[0, 0].plot(history['val_total_loss'], label='Validation', alpha=0.8)
axes[0, 0].set_title('Total Loss')
axes[0, 0].set_xlabel('Época')
axes[0, 0].set_ylabel('Loss')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)

# MSE loss
axes[0, 1].plot(history['train_mse_loss'], label='Train', color='orange', alpha=0.8)
axes[0, 1].plot(history['val_mse_loss'], label='Validation', color='red', alpha=0.8)

```

```

if CONFIG['early_stopping'] and best_epoch > 0:
    axes[0, 1].axvline(best_epoch - 1, color='green', linestyle='--', alpha=0.7, label=f'Best')
axes[0, 1].set_title('MSE Loss (Reconstrucción)')
axes[0, 1].set_xlabel('Época')
axes[0, 1].set_ylabel('MSE')
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)

# L0 sparsity
axes[0, 2].plot(history['train_l0_sparsity'], label='Train', color='green', alpha=0.8)
axes[0, 2].plot(history['val_l0_sparsity'], label='Validation', color='darkgreen', alpha=0.8)
axes[0, 2].set_title('L0 Sparsity (Features Activas)')
axes[0, 2].set_xlabel('Época')
axes[0, 2].set_ylabel('Número de features')
axes[0, 2].legend()
axes[0, 2].grid(True, alpha=0.3)

# Fraction active
axes[1, 0].plot(history['train_fraction_active'], label='Train', color='purple', alpha=0.8)
axes[1, 0].plot(history['val_fraction_active'], label='Validation', color='magenta', alpha=0.8)
axes[1, 0].set_title('Fracción de Features Activas')
axes[1, 0].set_xlabel('Época')
axes[1, 0].set_ylabel('Fracción')
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)

# Sparsity Coefficient (L1 Decay)
axes[1, 1].plot(history['sparsity_coef_history'], color='brown', linewidth=2)
axes[1, 1].set_title('L1 Coefficient (Decay)')
axes[1, 1].set_xlabel('Época')
axes[1, 1].set_ylabel('Sparsity Coef')
axes[1, 1].set_yscale('log')
axes[1, 1].grid(True, alpha=0.3)

# Train vs Val MSE Gap
mse_gap = [v - t for v, t in zip(history['val_mse_loss'], history['train_mse_loss'])]
axes[1, 2].plot(mse_gap, color='crimson', linewidth=2)
axes[1, 2].axhline(0, color='black', linestyle='--', alpha=0.5)
if CONFIG['early_stopping'] and best_epoch > 0:
    axes[1, 2].axvline(best_epoch - 1, color='green', linestyle='--', alpha=0.7, label=f'Best')
axes[1, 2].set_title('Overfitting Gap (Val MSE - Train MSE)')
axes[1, 2].set_xlabel('Época')

```



```


axes[1, 2].set_ylabel('Gap MSE')
axes[1, 2].legend()
axes[1, 2].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(CONFIG['plots_dir'] / 'training_curves.png', dpi=300, bbox_inches='tight')
plt.show()

print('\n' + '='*50)
print('  Gráficas guardadas exitosamente')

print('='*50)

```



=====
 Gráficas guardadas exitosamente
 =====

11. Resumen Final

```

# Resumen final completo
print('\n' + '='*70)
print(' RESUMEN FINAL DEL ENTRENAMIENTO')
print('='*70)

# Arquitectura
arch_summary = pd.DataFrame({
    'Componente': ['Input', 'Hidden', 'Parámetros'],
    'Valor': [
        f"{CONFIG['input_dim']}",
        f"{CONFIG['hidden_dim']} ({CONFIG['hidden_dim']/CONFIG['input_dim']:.1f}x)",
        f"{num_params:}"
    ]
})
print('\nArquitectura:')
print(arch_summary.to_string(index=False))

# Datos
data_summary = pd.DataFrame({
    'Conjunto': ['Train', 'Validation', 'Total'],
    'Muestras': [
        f"{len(train_dataset):}",
        f"{len(val_dataset):}",
        f"{len(full_dataset):}"
    ]
})
print('\nDatos:')
print(data_summary.to_string(index=False))

# Métricas finales
final_summary = pd.DataFrame({
    'Métrica': ['Train Loss', 'Val Loss', 'Train MSE', 'Val MSE', 'Train L0', 'Val L0'],
    'Valor': [
        f"{history['train_total_loss'][-1]:.4f}",
        f"{history['val_total_loss'][-1]:.4f}",
        f"{history['train_mse_loss'][-1]:.4f}",
        f"{history['val_mse_loss'][-1]:.4f}",
        f"{history['train_l0_sparsity'][-1]:.1f}",
        f"{history['val_l0_sparsity'][-1]:.1f}"
    ]
})
print('\nMétricas Finales:')

```

```

print(final_summary.to_string(index=False))

# Archivos
files_summary = pd.DataFrame({
    'Tipo': ['Modelo', 'Gráficas'],
    'Ubicación': [
        str(checkpoint_path),
        str(CONFIG['plots_dir'] / 'training_curves.png')
    ]
})
print('\nArchivos Guardados:')
print(files_summary.to_string(index=False))

print('\n' + '='*70)
print('  ENTRENAMIENTO COMPLETADO EXITOSAMENTE')
print('='*70)

```

```

=====
RESUMEN FINAL DEL ENTRENAMIENTO
=====

```

Arquitectura:

Componente	Valor
Input	512
Hidden	16384 (32.0x)
Parámetros	16,794,112

Datos:

Conjunto	Muestras
Train	23,600
Validation	5,900
Total	29,500

Métricas Finales:

Métrica	Valor
Train Loss	0.2084
Val Loss	0.2163
Train MSE	0.2053
Val MSE	0.2132
Train L0	151.3
Val L0	154.6

Archivos Guardados:

Tipo

Modelo

decay_l6_500g_final.pt

Gráficas c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard

decay_standard\500games\ex1\plots\training_curves.png

=====

ENTRENAMIENTO COMPLETADO EXITOSAMENTE

=====

12. Generar PDF con Quarto

Una vez ejecutado todo el notebook, ejecutar el siguiente comando en la terminal para generar el PDF con Quarto:

```
import subprocess

# Obtener directorio del experimento y nombre del PDF
exp_dir = get_experiment_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

pdf_name = get_report_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    'training',
    extras_id=extras_id
)

notebook_name = 'train_sae_standard_l1-decay.ipynb'
output_path = exp_dir / pdf_name

print(f' Generando PDF con Quarto...')
```

```

print(f' Archivo de salida: {output_path}')

# Quarto no acepta rutas en --output, solo nombre de archivo
# Se usa --output-dir para el directorio y --output solo para el nombre
result = subprocess.run(
    f'quarto render {notebook_name} --to pdf --output-dir "{exp_dir}" --output {pdf_name}',
    shell=True,
    capture_output=True,
    text=True
)

if result.returncode == 0:
    print(f' PDF generado exitosamente: {output_path}')
else:
    print(f' Error al generar PDF:')
    print(result.stderr)

```

```

Generando PDF con Quarto...
Archivo de salida: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\500games\ex1\sae_standard_l1-decay_16_500g_ex1_training.pdf
PDF generado exitosamente: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\500games\ex1\sae_standard_l1-decay_16_500g_ex1_training.pdf

```