

1. Setup e Imports

```
import json as _json
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import os
import sys
from pathlib import Path

# Agregar path del proyecto para imports
sys.path.append(os.path.abspath('../..'))
from sae.models.sae import SparseAutoencoder
from sae.tools.naming_utils import save_checkpoint, get_model_name, get_extras_id, get_checkpoint_dir, get_model_dir
from sae.tools.experiment_utils import get_experiment_dir, get_plots_dir, get_metrics_dir, get_logs_dir

# Configuración de visualización para Quarto/PDF
pd.set_option('display.width', 100)
pd.set_option('display.max_columns', None)
np.set_printoptions(linewidth=100, precision=4)
os.environ['COLUMNS'] = '100'

# Configuración de reproducibilidad
torch.manual_seed(42)
np.random.seed(42)

# Device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
# Print config con formato profesional
config_df = pd.DataFrame({
    'Configuración': ['Device', 'Seed', 'Output Width'],
    'Valor': [str(device), '42', '100']
})

print('\n' + '='*50)
print(' CONFIGURACIÓN INICIAL')
print('='*50)
print(config_df.to_string(index=False))
print('='*50)
```

```
=====
CONFIGURACIÓN INICIAL
=====
Configuración Valor
      Device  cuda
      Seed    42
Output Width  100
=====
```

2 Metadata para Naming de Checkpoints

```
# Metadata para naming de checkpoints (separada de hiperparámetros)
NAMING_META = {
    'variante': 'standard',          # Arquitectura del SAE
    'tecnica': 'l1-decay',           # L1 con penalty annealing
    'capa': 6,                        # Layer del modelo OthelloGPT
    'juegos': 500,                    # Número de partidas
    'base_path': 'C:\\Users\\Esposa\\Documents\\Repos\\sae-othello-gpt', # Path base para cl
    'experiment_base_path': os.path.abspath('../../experiments') ,      # sae/experiments,
    'extras': {
        'expansion': 32,
        'sparsity_coef': 5e-4,
        'epochs': 200,
        'patience': 15,
    }
}
```

```
print('\n Metadata de naming:')
for key, value in NAMING_META.items():
    print(f' {key}: {value}')
```

```
Metadata de naming:
variante: standard
tecnica: l1-decay
capa: 6
juegos: 500
base_path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt
experiment_base_path: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments
extras: {'expansion': 32, 'sparsity_coef': 0.0005, 'epochs': 200, 'patience': 15}
```

```
extras_id = get_extras_id(
    get_checkpoint_dir(
        NAMING_META['base_path'],
        NAMING_META['variante'],
        NAMING_META['tecnica'],
        NAMING_META['capa'],
        NAMING_META['juegos']
    ),
    NAMING_META['extras']
) if NAMING_META.get('extras') else None

print(f'extras_id: {extras_id}')
```

Nuevo extras registrado: ex3 → {'expansion': 32, 'sparsity_coef': 0.0005, 'epochs': 200, 'patience': 15}

extras_id: ex3

3. Configuración del SAE

Las activaciones ya fueron extraídas previamente con `sae/activations/run_extraction.py`. Aquí especificamos qué archivo cargar y los hiperparámetros del SAE.

```
# Configuración del SAE
CONFIG = {
    # Arquitectura
    'input_dim': 512,
```

```

'hidden_dim': 16384,
'tied_weights': False,

# Entrenamiento
'batch_size': 128,
'num_epochs': 200,
'learning_rate': 1e-3,
'sparsity_coef': 5e-4,
'sparsity_decay': 0.95,

# Early Stopping
'early_stopping': True,
'patience': 15,

# Datos
'activations_file': Path(f"../../activations/data/layer{NAMING_META['capa']}-{NAMING_I
'save_dir': Path('./saved_model'),
'plots_dir': get_plots_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)
}

# Crear directorio local del modelo (plots_dir ya fue creado por get_plots_dir)
CONFIG['save_dir'].mkdir(parents=True, exist_ok=True)

# Mostrar configuración con formato profesional
config_data = {
    'Categoría': [],
    'Parámetro': [],
    'Valor': []
}

# Arquitectura
for key in ['input_dim', 'hidden_dim', 'tied_weights']:
    config_data['Categoría'].append('Arquitectura')
    config_data['Parámetro'].append(key)

```

```

        config_data['Valor'].append(str(CONFIG[key]))

# Entrenamiento
for key in ['batch_size', 'num_epochs', 'learning_rate', 'sparsity_coef', 'sparsity_decay']:
    config_data['Categoría'].append('Entrenamiento')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

# Early Stopping
for key in ['early_stopping', 'patience']:
    config_data['Categoría'].append('Early Stopping')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

config_table = pd.DataFrame(config_data)

print('\n' + '='*70)
print(' CONFIGURACIÓN DEL SAE')
print('='*70)
print(config_table.to_string(index=False))
print('='*70)

print(f'\nArchivo de activaciones: {CONFIG["activations_file"]}')
print(f'Expansión: {CONFIG["hidden_dim"]/CONFIG["input_dim"]:.1f}x')
print(f'Plots dir: {CONFIG["plots_dir"]}')
print('='*70)

```

```

=====
CONFIGURACIÓN DEL SAE
=====

```

Categoría	Parámetro	Valor
Arquitectura	input_dim	512
Arquitectura	hidden_dim	16384
Arquitectura	tied_weights	False
Entrenamiento	batch_size	128
Entrenamiento	num_epochs	200
Entrenamiento	learning_rate	0.001
Entrenamiento	sparsity_coef	0.0005
Entrenamiento	sparsity_decay	0.95
Early Stopping	early_stopping	True
Early Stopping	patience	15

```

=====
Archivo de activaciones: ..\..\..\activations\data\layer6_500games.npy
Expansión: 32.0x
Plots dir: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\decay_standard\500games\ex3\plots
=====

```

4. Dataset de Activaciones (Train/Validation Split 80-20%)

```

class ActivationsDataset(Dataset):
    """Dataset para cargar activaciones extraídas"""

    def __init__(self, activations_path):
        self.activations = np.load(activations_path)

    def __len__(self):
        return len(self.activations)

    def __getitem__(self, idx):
        return torch.tensor(self.activations[idx], dtype=torch.float32)

# Cargar dataset completo
full_dataset = ActivationsDataset(CONFIG['activations_file'])

# Información del dataset con formato profesional
dataset_info = pd.DataFrame({
    'Propiedad': ['Total muestras', 'Shape', 'Dtype', 'Tamaño en memoria'],
    'Valor': [
        f"{len(full_dataset):,}",
        str(full_dataset.activations.shape),
        str(full_dataset.activations.dtype),
        f"{full_dataset.activations.nbytes / 1024**2:.2f} MB"
    ]
})

print('\n' + '='*50)
print(' DATASET DE ACTIVACIONES')
print('='*50)
print(dataset_info.to_string(index=False))
print('='*50)

```

```

# Split train/validation (80-20)
train_size = int(0.8 * len(full_dataset))
val_size = len(full_dataset) - train_size
train_dataset, val_dataset = torch.utils.data.random_split(
    full_dataset,
    [train_size, val_size],
    generator=torch.Generator().manual_seed(42)
)

# Resumen del split con formato profesional
split_info = pd.DataFrame({
    'Conjunto': ['Train', 'Validation', 'Total'],
    'Muestras': [
        f"{len(train_dataset):,}",
        f"{len(val_dataset):,}",
        f"{len(full_dataset):,}"
    ],
    'Porcentaje': [
        f"{len(train_dataset)/len(full_dataset):.1%}",
        f"{len(val_dataset)/len(full_dataset):.1%}",
        "100.0%"
    ]
})
print('\n' + '='*50)
print(' SPLIT DE DATOS (80-20) ')
print('='*50)
print(split_info.to_string(index=False))
print('='*50)

# Crear DataLoaders
train_loader = DataLoader(
    train_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=True,
    num_workers=0,
    pin_memory=(device.type == 'cuda')
)

val_loader = DataLoader(
    val_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=False, # No shuffle para validación

```

```

    num_workers=0,
    pin_memory=(device.type == 'cuda')
)

# Resumen de DataLoaders
loader_info = pd.DataFrame({
    'DataLoader': ['Train', 'Validation'],
    'Batches': [len(train_loader), len(val_loader)],
    'Batch Size': [CONFIG['batch_size'], CONFIG['batch_size']],
    'Shuffle': ['Sí', 'No']
})

print('\n' + '='*50)
print(' DATALOADERS')
print('='*50)
print(loader_info.to_string(index=False))
print('='*50)

```

```

=====
DATASET DE ACTIVACIONES
=====

```

Propiedad	Valor
Total muestras	29,500
Shape	(29500, 512)
Dtype	float32
Tamaño en memoria	57.62 MB

```

=====

```

```

=====
SPLIT DE DATOS (80-20)
=====

```

Conjunto	Muestras	Porcentaje
Train	23,600	80.0%
Validation	5,900	20.0%
Total	29,500	100.0%

```

=====

```

```

=====
DATALOADERS
=====

```

DataLoader	Batches	Batch Size	Shuffle
------------	---------	------------	---------

Train	185	128	Sí
Validation	47	128	No
=====			

5. Modelo: Sparse Autoencoder

Modelo importado desde `sae.models.sae.SparseAutoencoder`

Arquitectura:

- **Encoder:** Linear $\rightarrow$ ReLU (sparsity natural)- **Decoder:** Linear (reconstrucción)

```
# Crear modelo desde el módulo
model = SparseAutoencoder(
    input_dim=CONFIG['input_dim'],
    hidden_dim=CONFIG['hidden_dim'],
    tied_weights=CONFIG['tied_weights']
).to(device)

# Calcular parámetros
num_params = sum(p.numel() for p in model.parameters())
encoder_params = sum(p.numel() for p in model.encoder.parameters())
if CONFIG['tied_weights']:
    decoder_params = CONFIG['input_dim'] # Solo bias
else:
    decoder_params = sum(p.numel() for p in model.decoder.parameters())

# Resumen del modelo con formato profesional
model_info = pd.DataFrame({
    'Componente': ['Input', 'Encoder', 'Hidden (Latent)', 'Decoder', 'Output', 'Total'],
    'Dimensión': [
        CONFIG['input_dim'],
        f"{CONFIG['input_dim']}  $\rightarrow$  {CONFIG['hidden_dim']}",
        CONFIG['hidden_dim'],
        f"{CONFIG['hidden_dim']}  $\rightarrow$  {CONFIG['input_dim']}",
        CONFIG['input_dim'],
        '-',
    ],
    'Parámetros': [
        '-',
        f"{encoder_params:,}",
        '-',
        f"{decoder_params:,}",
    ]
})
```

```

        '-',
        f"{num_params: ,}"
    ]
})

print('\n' + '='*70)
print(' ARQUITECTURA DEL MODELO')
print('='*70)
print(model_info.to_string(index=False))
print('='*70)
print(f'\n Factor de expansión: {CONFIG["hidden_dim"]/CONFIG["input_dim"]:.1f}x')
print(f' Tied weights: {"Sí" if CONFIG["tied_weights"] else "No"}')
print(f' Device: {device}')
print('='*70)

```

```

=====
ARQUITECTURA DEL MODELO
=====

```

	Componente	Dimensión	Parámetros
	Input	512	-
	Encoder	512 → 16384	8,404,992
Hidden	(Latent)	16384	-
	Decoder	16384 → 512	8,389,120
	Output	512	-
	Total	-	16,794,112

```

=====

Factor de expansión: 32.0x
Tied weights: No
Device: cuda
=====

```

6. Función de Pérdida

$\text{Loss} = \text{MSE}(\text{reconstrucción}) + \lambda \times \text{L1}(\text{activaciones})$

- **MSE**: Mide qué tan bien reconstruimos las activaciones originales
- **L1**: Penaliza activaciones grandes → fuerza sparsity

```

def loss_function(x, reconstruction, hidden, sparsity_coef):
    """
    Pérdida compuesta: MSE + L1 penalty

    Args:
        x: Activaciones originales
        reconstruction: Activaciones reconstruidas
        hidden: Activaciones latentes (sparse)
        sparsity_coef: Peso de la penalización L1
    """
    # Error de reconstrucción
    mse_loss = F.mse_loss(reconstruction, x)

    # Sumar sobre las features (dim=1), promediar sobre el batch
    sparsity_loss = hidden.abs().sum(dim=1).mean()

    # Pérdida total
    total_loss = mse_loss + sparsity_coef * sparsity_loss

    return total_loss, mse_loss, sparsity_loss

# Optimizador
optimizer = torch.optim.Adam(model.parameters(), lr=CONFIG['learning_rate'])

# Resumen del optimizador
optim_info = pd.DataFrame({
    'Parámetro': ['Optimizador', 'Learning Rate', 'Sparsity Coef (inicial)', 'Sparsity Decay'],
    'Valor': [
        'Adam',
        f"{CONFIG['learning_rate']:.0e}",
        f"{CONFIG['sparsity_coef']:.0e}",
        f"{CONFIG['sparsity_decay']:.2f}"
    ]
})

print('\n' + '='*50)
print(' OPTIMIZACIÓN')
print('='*50)
print(optim_info.to_string(index=False))
print('='*50)

```

```
=====
OPTIMIZACIÓN
=====
                Parámetro Valor
                Optimizador Adam
                Learning Rate 1e-03
Sparsity Coef (inicial) 5e-04
                Sparsity Decay 0.95
=====
```

7. Loop de Entrenamiento

```
# Historial de entrenamiento
history = {
    'train_total_loss': [],
    'train_mse_loss': [],
    'train_sparsity_loss': [],
    'train_l0_sparsity': [],
    'train_fraction_active': [],
    'val_total_loss': [],
    'val_mse_loss': [],
    'val_sparsity_loss': [],
    'val_l0_sparsity': [],
    'val_fraction_active': [],
    'sparsity_coef_history': [] # Guardar evolución del coeficiente L1
}

# Early stopping variables
best_val_mse = float('inf')
best_epoch = 0
epochs_no_improve = 0

print(f'\n Iniciando entrenamiento: {CONFIG["num_epochs"]} épocas máximo')
print(f' L1 decay: {CONFIG["sparsity_coef"]:.2e} × {CONFIG["sparsity_decay"]:.2f}^época')
print(f' Early stopping: patience={CONFIG["patience"]} épocas')
print('=' * 70)

for epoch in range(CONFIG['num_epochs']):
    # ===== L1 DECAY =====
    # Reducir sparsity_coef exponencialmente
    current_sparsity_coef = CONFIG['sparsity_coef'] * (CONFIG['sparsity_decay'] ** epoch)
```

```

history['sparsity_coef_history'].append(current_sparsity_coef)

# ===== ENTRENAMIENTO =====
model.train()

epoch_total_loss = 0
epoch_mse_loss = 0
epoch_sparsity_loss = 0
epoch_l0 = 0
epoch_fraction_active = 0

# Barra de progreso
pbar = tqdm(train_loader, desc=f'Época {epoch+1}/{CONFIG["num_epochs"]} [Train] L1={curr

for batch in pbar:
    # Mover al device
    x = batch.to(device)

    # Forward pass
    reconstruction, hidden = model(x)

    # Calcular pérdida con sparsity_coef actualizado
    total_loss, mse_loss, sparsity_loss = loss_function(
        x, reconstruction, hidden, current_sparsity_coef
    )

    # Backward pass
    optimizer.zero_grad()
    total_loss.backward()

    optimizer.step()

    # Métricas
    with torch.no_grad():
        l0 = (hidden > 0).float().sum(dim=1).mean().item()
        fraction_active = (hidden > 0).float().mean().item()

    # Acumular
    epoch_total_loss += total_loss.item()
    epoch_mse_loss += mse_loss.item()
    epoch_sparsity_loss += sparsity_loss.item()
    epoch_l0 += l0

```

```

epoch_fraction_active += fraction_active

# Actualizar barra
pbar.set_postfix({
    'loss': f'{total_loss.item():.4f}',
    'mse': f'{mse_loss.item():.4f}',
    'l0': f'{l0:.1f}'
})

# Promedios de entrenamiento
num_train_batches = len(train_loader)
avg_train_total_loss = epoch_total_loss / num_train_batches
avg_train_mse_loss = epoch_mse_loss / num_train_batches
avg_train_sparsity_loss = epoch_sparsity_loss / num_train_batches
avg_train_l0 = epoch_l0 / num_train_batches
avg_train_fraction_active = epoch_fraction_active / num_train_batches

# ===== VALIDACIÓN =====
model.eval()

val_total_loss = 0
val_mse_loss = 0
val_sparsity_loss = 0
val_l0 = 0
val_fraction_active = 0

with torch.no_grad():
    for batch in val_loader:
        x = batch.to(device)

        # Forward pass
        reconstruction, hidden = model(x)

        # Calcular pérdida con sparsity_coef actualizado
        total_loss, mse_loss, sparsity_loss = loss_function(
            x, reconstruction, hidden, current_sparsity_coef
        )

        # Métricas
        l0 = (hidden > 0).float().sum(dim=1).mean().item()
        fraction_active = (hidden > 0).float().mean().item()

```

```

        # Acumular
        val_total_loss += total_loss.item()
        val_mse_loss += mse_loss.item()
        val_sparsity_loss += sparsity_loss.item()
        val_l0 += l0
        val_fraction_active += fraction_active

# Promedios de validación
num_val_batches = len(val_loader)
avg_val_total_loss = val_total_loss / num_val_batches
avg_val_mse_loss = val_mse_loss / num_val_batches
avg_val_sparsity_loss = val_sparsity_loss / num_val_batches
avg_val_l0 = val_l0 / num_val_batches
avg_val_fraction_active = val_fraction_active / num_val_batches

# Guardar en historial
history['train_total_loss'].append(avg_train_total_loss)
history['train_mse_loss'].append(avg_train_mse_loss)
history['train_sparsity_loss'].append(avg_train_sparsity_loss)
history['train_l0_sparsity'].append(avg_train_l0)
history['train_fraction_active'].append(avg_train_fraction_active)

history['val_total_loss'].append(avg_val_total_loss)
history['val_mse_loss'].append(avg_val_mse_loss)
history['val_sparsity_loss'].append(avg_val_sparsity_loss)
history['val_l0_sparsity'].append(avg_val_l0)
history['val_fraction_active'].append(avg_val_fraction_active)

# ===== EARLY STOPPING =====
if CONFIG['early_stopping']:
    if avg_val_mse_loss < best_val_mse:
        best_val_mse = avg_val_mse_loss
        best_epoch = epoch + 1
        epochs_no_improve = 0

    # Guardar mejor modelo (local + organizado)
    best_checkpoint = {
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'config': CONFIG,
        'history': history,
        'epoch': best_epoch,

```

```

        'val_mse': best_val_mse
    }

    # Guardado local
    best_model_name = get_model_name(
        variante=NAMING_META['variante'],
        tecnica=NAMING_META['tecnica'],
        capa=NAMING_META['capa'],
        juegos=NAMING_META['juegos'],
        estado='best'
    )
    torch.save(best_checkpoint, CONFIG['save_dir'] / best_model_name)

    # Guardado organizado
    save_checkpoint(
        model,
        base_path=NAMING_META['base_path'],
        variante=NAMING_META['variante'],
        tecnica=NAMING_META['tecnica'],
        capa=NAMING_META['capa'],
        juegos=NAMING_META['juegos'],
        estado='best',
        extras=NAMING_META['extras'],
        config=CONFIG,
        history=history,
        optimizer=optimizer
    )

    print(f' Mejor modelo guardado en época {best_epoch} (Val MSE: {best_val_mse:.4f})')
else:
    epochs_no_improve += 1
    if epochs_no_improve >= CONFIG['patience']:
        print(f'\n Early stopping activado en época {epoch+1}')
        print(f' Mejor modelo: época {best_epoch} con Val MSE = {best_val_mse:.4f}')
        break

# Log cada 5 épocas
if (epoch + 1) % 5 == 0 or epoch == 0:
    print(f'\n Época {epoch+1}/{CONFIG["num_epochs"]} | L1 coef: {current_sparsity_coef:.4f}')
    print(f' Train - Total: {avg_train_total_loss:.4f} | MSE: {avg_train_mse_loss:.4f}')
    print(f' Val - Total: {avg_val_total_loss:.4f} | MSE: {avg_val_mse_loss:.4f} | L0')
    if CONFIG['early_stopping']:

```



```

        print(f'    Best Val MSE: {best_val_mse:.4f} @ época {best_epoch} | Sin mejora: {best_val_mse:.4f}')

print('\n' + '='*70)
print(' RESUMEN DEL ENTRENAMIENTO')
print('='*70)

if CONFIG['early_stopping']:
    print(f'Mejor modelo en época {best_epoch} con Val MSE = {best_val_mse:.4f}')
    print(f'Early stopping tras {len(history["train_total_loss"])} épocas')
else:
    print(f'Entrenamiento completo: {CONFIG["num_epochs"]} épocas')

# Tabla de métricas finales
final_metrics = pd.DataFrame({
    'Métrica': ['Total Loss', 'MSE Loss', 'L0 Sparsity', 'Fraction Active'],
    'Train': [
        f"{history['train_total_loss'][-1]:.4f}",
        f"{history['train_mse_loss'][-1]:.4f}",
        f"{history['train_l0_sparsity'][-1]:.1f}",
        f"{history['train_fraction_active'][-1]:.2%}"
    ],
    'Validation': [
        f"{history['val_total_loss'][-1]:.4f}",
        f"{history['val_mse_loss'][-1]:.4f}",
        f"{history['val_l0_sparsity'][-1]:.1f}",
        f"{history['val_fraction_active'][-1]:.2%}"
    ]
})

print('\nMétricas Finales (Última Época):')
print(final_metrics.to_string(index=False))
print('='*70)

```

Iniciando entrenamiento: 200 épocas máximo
 L1 decay: $5.00e-04 \times 0.95^{\text{época}}$
 Early stopping: patience=15 épocas

=====

Época 1/200 [Train] L1=5.00e-04: 100% | 185/185 [00:01<00:00, 94.67it/s, loss=0.2294, m

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 1 (Val MSE: 0.1145)

Época 1/200 | L1 coef: 5.00e-04
Train - Total: 0.3985 | MSE: 0.2853 | L0: 1065.4
Val - Total: 0.2207 | MSE: 0.1145 | L0: 827.8
Best Val MSE: 0.1145 @ época 1 | Sin mejora: 0/15

Época 2/200 [Train] L1=4.75e-04: 100% | 185/185 [00:01<00:00, 108.06it/s, loss=0.1541,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 2 (Val MSE: 0.0805)

Época 3/200 [Train] L1=4.51e-04: 100% | 185/185 [00:01<00:00, 106.92it/s, loss=0.1404,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 3 (Val MSE: 0.0753)

Época 4/200 [Train] L1=4.29e-04: 100% | 185/185 [00:01<00:00, 114.14it/s, loss=0.1300,
Época 5/200 [Train] L1=4.07e-04: 100% | 185/185 [00:01<00:00, 110.10it/s, loss=0.1234,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 5 (Val MSE: 0.0742)

Época 5/200 | L1 coef: 4.07e-04
Train - Total: 0.1155 | MSE: 0.0566 | L0: 625.6
Val - Total: 0.1318 | MSE: 0.0742 | L0: 645.7
Best Val MSE: 0.0742 @ época 5 | Sin mejora: 0/15

Época 6/200 [Train] L1=3.87e-04: 100% | 185/185 [00:01<00:00, 110.23it/s, loss=0.1088,
Época 7/200 [Train] L1=3.68e-04: 100% | 185/185 [00:01<00:00, 121.22it/s, loss=0.0992,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 7 (Val MSE: 0.0738)

Época 8/200 [Train] L1=3.49e-04: 100%| | 185/185 [00:01<00:00, 126.04it/s, loss=0.0946,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 8 (Val MSE: 0.0734)

Época 9/200 [Train] L1=3.32e-04: 100%| | 185/185 [00:01<00:00, 105.72it/s, loss=0.0897,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 9 (Val MSE: 0.0726)

Época 10/200 [Train] L1=3.15e-04: 100%| | 185/185 [00:01<00:00, 99.79it/s, loss=0.0880,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 10 (Val MSE: 0.0722)

Época 10/200 | L1 coef: 3.15e-04

Train - Total: 0.0841 | MSE: 0.0452 | L0: 587.8

Val - Total: 0.1104 | MSE: 0.0722 | L0: 607.3

Best Val MSE: 0.0722 @ época 10 | Sin mejora: 0/15

Época 11/200 [Train] L1=2.99e-04: 100%| | 185/185 [00:01<00:00, 102.88it/s, loss=0.0828,

Época 12/200 [Train] L1=2.84e-04: 100%| | 185/185 [00:01<00:00, 102.76it/s, loss=0.0753,

Época 13/200 [Train] L1=2.70e-04: 100%| | 185/185 [00:01<00:00, 101.48it/s, loss=0.0770,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 13 (Val MSE: 0.0717)

Época 14/200 [Train] L1=2.57e-04: 100%| | 185/185 [00:01<00:00, 108.36it/s, loss=0.0748,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 14 (Val MSE: 0.0709)

Época 15/200 [Train] L1=2.44e-04: 100%| | 185/185 [00:01<00:00, 105.06it/s, loss=0.0726,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 15 (Val MSE: 0.0697)

Época 15/200 | L1 coef: 2.44e-04
Train - Total: 0.0697 | MSE: 0.0408 | L0: 568.8
Val - Total: 0.0984 | MSE: 0.0697 | L0: 589.3
Best Val MSE: 0.0697 @ época 15 | Sin mejora: 0/15

Época 16/200 [Train] L1=2.32e-04: 100%| | 185/185 [00:01<00:00, 111.62it/s, loss=0.0662,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 16 (Val MSE: 0.0689)

Época 17/200 [Train] L1=2.20e-04: 100%| | 185/185 [00:01<00:00, 103.57it/s, loss=0.0642,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 17 (Val MSE: 0.0669)

Época 18/200 [Train] L1=2.09e-04: 100%| | 185/185 [00:01<00:00, 106.13it/s, loss=0.0644,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 18 (Val MSE: 0.0663)

Época 19/200 [Train] L1=1.99e-04: 100%| | 185/185 [00:01<00:00, 104.61it/s, loss=0.0610,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 19 (Val MSE: 0.0656)

Época 20/200 [Train] L1=1.89e-04: 100%| | 185/185 [00:01<00:00, 101.85it/s, loss=0.0574,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 20 (Val MSE: 0.0636)

Época 20/200 | L1 coef: 1.89e-04
Train - Total: 0.0574 | MSE: 0.0346 | L0: 565.9
Val - Total: 0.0861 | MSE: 0.0636 | L0: 584.2
Best Val MSE: 0.0636 @ época 20 | Sin mejora: 0/15

Época 21/200 [Train] L1=1.79e-04: 100%| | 185/185 [00:01<00:00, 103.89it/s, loss=0.0565,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 21 (Val MSE: 0.0625)

Época 22/200 [Train] L1=1.70e-04: 100%| | 185/185 [00:01<00:00, 104.32it/s, loss=0.0544,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 22 (Val MSE: 0.0623)

Época 23/200 [Train] L1=1.62e-04: 100%| | 185/185 [00:01<00:00, 99.78it/s, loss=0.0510,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 23 (Val MSE: 0.0597)

Época 24/200 [Train] L1=1.54e-04: 100%| | 185/185 [00:01<00:00, 101.72it/s, loss=0.0525,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 24 (Val MSE: 0.0573)

Época 25/200 [Train] L1=1.46e-04: 100%| | 185/185 [00:01<00:00, 102.23it/s, loss=0.0491,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 25 (Val MSE: 0.0547)

Época 25/200 | L1 coef: 1.46e-04
Train - Total: 0.0477 | MSE: 0.0289 | L0: 579.9
Val - Total: 0.0736 | MSE: 0.0547 | L0: 597.0
Best Val MSE: 0.0547 @ época 25 | Sin mejora: 0/15

Época 26/200 [Train] L1=1.39e-04: 100%| | 185/185 [00:01<00:00, 100.78it/s, loss=0.0476,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 26 (Val MSE: 0.0532)

Época 27/200 [Train] L1=1.32e-04: 100%| | 185/185 [00:01<00:00, 103.49it/s, loss=0.0502,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 27 (Val MSE: 0.0515)

Época 28/200 [Train] L1=1.25e-04: 100%| | 185/185 [00:01<00:00, 104.60it/s, loss=0.0425,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 28 (Val MSE: 0.0507)

Época 29/200 [Train] L1=1.19e-04: 100%| | 185/185 [00:01<00:00, 108.16it/s, loss=0.0415,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 29 (Val MSE: 0.0476)

Época 30/200 [Train] L1=1.13e-04: 100%| | 185/185 [00:01<00:00, 104.89it/s, loss=0.0422,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 30 (Val MSE: 0.0452)

Época 30/200 | L1 coef: 1.13e-04
Train - Total: 0.0401 | MSE: 0.0241 | L0: 600.4
Val - Total: 0.0613 | MSE: 0.0452 | L0: 611.8
Best Val MSE: 0.0452 @ época 30 | Sin mejora: 0/15

Época 31/200 [Train] L1=1.07e-04: 100%| | 185/185 [00:01<00:00, 102.81it/s, loss=0.0366,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 31 (Val MSE: 0.0427)

Época 32/200 [Train] L1=1.02e-04: 100%| | 185/185 [00:01<00:00, 106.04it/s, loss=0.0381,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 32 (Val MSE: 0.0414)

Época 33/200 [Train] L1=9.69e-05: 100%| | 185/185 [00:01<00:00, 107.58it/s, loss=0.0364,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 33 (Val MSE: 0.0387)

Época 34/200 [Train] L1=9.20e-05: 100%| | 185/185 [00:01<00:00, 101.54it/s, loss=0.0328,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 34 (Val MSE: 0.0364)

Época 35/200 [Train] L1=8.74e-05: 100%| | 185/185 [00:01<00:00, 104.31it/s, loss=0.0342,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 35 (Val MSE: 0.0358)

Época 35/200 | L1 coef: 8.74e-05
Train - Total: 0.0329 | MSE: 0.0190 | L0: 620.1
Val - Total: 0.0498 | MSE: 0.0358 | L0: 624.2
Best Val MSE: 0.0358 @ época 35 | Sin mejora: 0/15

Época 36/200 [Train] L1=8.30e-05: 100%| | 185/185 [00:01<00:00, 105.43it/s, loss=0.0313,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 36 (Val MSE: 0.0319)

Época 37/200 [Train] L1=7.89e-05: 100%| | 185/185 [00:01<00:00, 107.01it/s, loss=0.0325,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 37 (Val MSE: 0.0295)

Época 38/200 [Train] L1=7.49e-05: 100%| | 185/185 [00:01<00:00, 106.86it/s, loss=0.0290,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 38 (Val MSE: 0.0278)

Época 39/200 [Train] L1=7.12e-05: 100%| | 185/185 [00:01<00:00, 124.82it/s, loss=0.0278,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 39 (Val MSE: 0.0257)

Época 40/200 [Train] L1=6.76e-05: 100%| | 185/185 [00:01<00:00, 120.22it/s, loss=0.0257,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 40 (Val MSE: 0.0235)

Época 40/200 | L1 coef: 6.76e-05
Train - Total: 0.0258 | MSE: 0.0137 | L0: 628.6
Val - Total: 0.0359 | MSE: 0.0235 | L0: 629.2
Best Val MSE: 0.0235 @ época 40 | Sin mejora: 0/15

Época 41/200 [Train] L1=6.43e-05: 100%| | 185/185 [00:01<00:00, 120.66it/s, loss=0.0243,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 41 (Val MSE: 0.0216)

Época 42/200 [Train] L1=6.10e-05: 100%| | 185/185 [00:01<00:00, 120.96it/s, loss=0.0231,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 42 (Val MSE: 0.0202)

Época 43/200 [Train] L1=5.80e-05: 100%| | 185/185 [00:01<00:00, 123.50it/s, loss=0.0240,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 43 (Val MSE: 0.0187)

Época 44/200 [Train] L1=5.51e-05: 100%| | 185/185 [00:01<00:00, 124.19it/s, loss=0.0195,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 44 (Val MSE: 0.0172)

Época 45/200 [Train] L1=5.23e-05: 100%| | 185/185 [00:01<00:00, 121.51it/s, loss=0.0188,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 45 (Val MSE: 0.0155)

Época 45/200 | L1 coef: 5.23e-05
Train - Total: 0.0199 | MSE: 0.0095 | L0: 623.9
Val - Total: 0.0261 | MSE: 0.0155 | L0: 621.4
Best Val MSE: 0.0155 @ época 45 | Sin mejora: 0/15

Época 46/200 [Train] L1=4.97e-05: 100%| | 185/185 [00:01<00:00, 104.25it/s, loss=0.0188,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 46 (Val MSE: 0.0148)

Época 47/200 [Train] L1=4.72e-05: 100%| | 185/185 [00:01<00:00, 109.23it/s, loss=0.0170,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 47 (Val MSE: 0.0137)

Época 48/200 [Train] L1=4.49e-05: 100%| | 185/185 [00:01<00:00, 109.53it/s, loss=0.0169,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 48 (Val MSE: 0.0130)

Época 49/200 [Train] L1=4.26e-05: 100%| | 185/185 [00:01<00:00, 107.21it/s, loss=0.0169,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 49 (Val MSE: 0.0123)

Época 50/200 [Train] L1=4.05e-05: 100%| | 185/185 [00:01<00:00, 107.55it/s, loss=0.0145,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 50 (Val MSE: 0.0116)

Época 50/200 | L1 coef: 4.05e-05
Train - Total: 0.0162 | MSE: 0.0075 | L0: 615.2
Val - Total: 0.0203 | MSE: 0.0116 | L0: 612.1
Best Val MSE: 0.0116 @ época 50 | Sin mejora: 0/15

Época 51/200 [Train] L1=3.85e-05: 100%| | 185/185 [00:01<00:00, 110.11it/s, loss=0.0147,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 51 (Val MSE: 0.0113)

Época 52/200 [Train] L1=3.65e-05: 100%| | 185/185 [00:01<00:00, 107.38it/s, loss=0.0178,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 52 (Val MSE: 0.0107)

Época 53/200 [Train] L1=3.47e-05: 100%| | 185/185 [00:01<00:00, 104.80it/s, loss=0.0133,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 53 (Val MSE: 0.0097)

Época 54/200 [Train] L1=3.30e-05: 100%| | 185/185 [00:01<00:00, 105.68it/s, loss=0.0143,
Época 55/200 [Train] L1=3.13e-05: 100%| | 185/185 [00:01<00:00, 103.47it/s, loss=0.0121,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 55 (Val MSE: 0.0087)

Época 55/200 | L1 coef: 3.13e-05
Train - Total: 0.0131 | MSE: 0.0059 | L0: 612.5
Val - Total: 0.0159 | MSE: 0.0087 | L0: 610.4
Best Val MSE: 0.0087 @ época 55 | Sin mejora: 0/15

Época 56/200 [Train] L1=2.98e-05: 100%| | 185/185 [00:01<00:00, 104.00it/s, loss=0.0129,
Época 57/200 [Train] L1=2.83e-05: 100%| | 185/185 [00:01<00:00, 104.59it/s, loss=0.0122,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 57 (Val MSE: 0.0086)

Época 58/200 [Train] L1=2.69e-05: 100%| | 185/185 [00:01<00:00, 106.22it/s, loss=0.0121,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 58 (Val MSE: 0.0077)

Época 59/200 [Train] L1=2.55e-05: 100%| | 185/185 [00:01<00:00, 106.21it/s, loss=0.0108,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 59 (Val MSE: 0.0076)

Época 60/200 [Train] L1=2.42e-05: 100%| | 185/185 [00:01<00:00, 102.33it/s, loss=0.0107,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 60 (Val MSE: 0.0074)

Época 60/200 | L1 coef: 2.42e-05

Train - Total: 0.0107 | MSE: 0.0049 | L0: 611.5

Val - Total: 0.0133 | MSE: 0.0074 | L0: 608.9

Best Val MSE: 0.0074 @ época 60 | Sin mejora: 0/15

Época 61/200 [Train] L1=2.30e-05: 100%| | 185/185 [00:01<00:00, 99.37it/s, loss=0.0110,

Época 62/200 [Train] L1=2.19e-05: 100%| | 185/185 [00:01<00:00, 94.94it/s, loss=0.0104,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_st
decay_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 62 (Val MSE: 0.0069)

Época 63/200 [Train] L1=2.08e-05: 100%| | 185/185 [00:01<00:00, 112.16it/s, loss=0.0095,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 63 (Val MSE: 0.0065)

Época 64/200 [Train] L1=1.97e-05: 100%| | 185/185 [00:01<00:00, 109.13it/s, loss=0.0090,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 64 (Val MSE: 0.0061)

Época 65/200 [Train] L1=1.88e-05: 100%| | 185/185 [00:01<00:00, 107.14it/s, loss=0.0098,

Época 65/200 | L1 coef: 1.88e-05
Train - Total: 0.0090 | MSE: 0.0043 | L0: 613.0
Val - Total: 0.0114 | MSE: 0.0066 | L0: 611.0
Best Val MSE: 0.0061 @ época 64 | Sin mejora: 1/15

Época 66/200 [Train] L1=1.78e-05: 100%| | 185/185 [00:01<00:00, 109.12it/s, loss=0.0096,
Época 67/200 [Train] L1=1.69e-05: 100%| | 185/185 [00:01<00:00, 96.72it/s, loss=0.0081,
Época 68/200 [Train] L1=1.61e-05: 100%| | 185/185 [00:01<00:00, 104.54it/s, loss=0.0089,
Época 69/200 [Train] L1=1.53e-05: 100%| | 185/185 [00:01<00:00, 106.87it/s, loss=0.0070,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 69 (Val MSE: 0.0056)

Época 70/200 [Train] L1=1.45e-05: 100%| | 185/185 [00:01<00:00, 109.28it/s, loss=0.0083,

Época 70/200 | L1 coef: 1.45e-05
Train - Total: 0.0076 | MSE: 0.0038 | L0: 613.2
Val - Total: 0.0095 | MSE: 0.0057 | L0: 610.8
Best Val MSE: 0.0056 @ época 69 | Sin mejora: 1/15

Época 71/200 [Train] L1=1.38e-05: 100%| | 185/185 [00:01<00:00, 105.53it/s, loss=0.0081,
Época 72/200 [Train] L1=1.31e-05: 100%| | 185/185 [00:01<00:00, 103.96it/s, loss=0.0074,
Época 73/200 [Train] L1=1.24e-05: 100%| | 185/185 [00:01<00:00, 107.57it/s, loss=0.0061,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 73 (Val MSE: 0.0047)

Época 74/200 [Train] L1=1.18e-05: 100%| | 185/185 [00:01<00:00, 107.72it/s, loss=0.0062,

Época 75/200 [Train] L1=1.12e-05: 100%| | 185/185 [00:01<00:00, 109.86it/s, loss=0.0062,

Época 75/200 | L1 coef: 1.12e-05

Train - Total: 0.0067 | MSE: 0.0037 | L0: 614.9

Val - Total: 0.0080 | MSE: 0.0050 | L0: 613.3

Best Val MSE: 0.0047 @ época 73 | Sin mejora: 2/15

Época 76/200 [Train] L1=1.07e-05: 100%| | 185/185 [00:01<00:00, 108.12it/s, loss=0.0058,

Época 77/200 [Train] L1=1.01e-05: 100%| | 185/185 [00:01<00:00, 112.06it/s, loss=0.0066,

Época 78/200 [Train] L1=9.63e-06: 100%| | 185/185 [00:01<00:00, 108.19it/s, loss=0.0062,

Época 79/200 [Train] L1=9.15e-06: 100%| | 185/185 [00:01<00:00, 106.45it/s, loss=0.0069,

Época 80/200 [Train] L1=8.69e-06: 100%| | 185/185 [00:01<00:00, 109.34it/s, loss=0.0062,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 80 (Val MSE: 0.0041)

Época 80/200 | L1 coef: 8.69e-06

Train - Total: 0.0059 | MSE: 0.0034 | L0: 616.5

Val - Total: 0.0066 | MSE: 0.0041 | L0: 615.6

Best Val MSE: 0.0041 @ época 80 | Sin mejora: 0/15

Época 81/200 [Train] L1=8.26e-06: 100%| | 185/185 [00:01<00:00, 106.13it/s, loss=0.0064,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 81 (Val MSE: 0.0041)

Época 82/200 [Train] L1=7.84e-06: 100%| | 185/185 [00:01<00:00, 105.22it/s, loss=0.0063,

Época 83/200 [Train] L1=7.45e-06: 100%| | 185/185 [00:01<00:00, 117.32it/s, loss=0.0054,

Época 84/200 [Train] L1=7.08e-06: 100%| | 185/185 [00:01<00:00, 121.43it/s, loss=0.0049,

Época 85/200 [Train] L1=6.73e-06: 100%| | 185/185 [00:01<00:00, 121.22it/s, loss=0.0051,

Época 85/200 | L1 coef: 6.73e-06
Train - Total: 0.0047 | MSE: 0.0027 | L0: 617.6
Val - Total: 0.0071 | MSE: 0.0051 | L0: 614.9
Best Val MSE: 0.0041 @ época 81 | Sin mejora: 4/15

Época 86/200 [Train] L1=6.39e-06: 100%| | 185/185 [00:01<00:00, 120.00it/s, loss=0.0050,
Época 87/200 [Train] L1=6.07e-06: 100%| | 185/185 [00:01<00:00, 123.80it/s, loss=0.0039,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 87 (Val MSE: 0.0036)

Época 88/200 [Train] L1=5.77e-06: 100%| | 185/185 [00:01<00:00, 121.06it/s, loss=0.0045,
Época 89/200 [Train] L1=5.48e-06: 100%| | 185/185 [00:01<00:00, 119.14it/s, loss=0.0053,
Época 90/200 [Train] L1=5.20e-06: 100%| | 185/185 [00:01<00:00, 121.16it/s, loss=0.0039,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 90 (Val MSE: 0.0034)

Época 90/200 | L1 coef: 5.20e-06
Train - Total: 0.0047 | MSE: 0.0031 | L0: 616.5
Val - Total: 0.0050 | MSE: 0.0034 | L0: 615.0
Best Val MSE: 0.0034 @ época 90 | Sin mejora: 0/15

Época 91/200 [Train] L1=4.94e-06: 100%| | 185/185 [00:01<00:00, 123.63it/s, loss=0.0036,
Época 92/200 [Train] L1=4.70e-06: 100%| | 185/185 [00:01<00:00, 123.68it/s, loss=0.0045,
Época 93/200 [Train] L1=4.46e-06: 100%| | 185/185 [00:01<00:00, 122.53it/s, loss=0.0040,
Época 94/200 [Train] L1=4.24e-06: 100%| | 185/185 [00:01<00:00, 123.60it/s, loss=0.0041,
Época 95/200 [Train] L1=4.03e-06: 100%| | 185/185 [00:01<00:00, 121.98it/s, loss=0.0038,

Época 95/200 | L1 coef: 4.03e-06
Train - Total: 0.0042 | MSE: 0.0030 | L0: 614.0
Val - Total: 0.0056 | MSE: 0.0043 | L0: 611.1
Best Val MSE: 0.0034 @ época 90 | Sin mejora: 5/15

Época 96/200 [Train] L1=3.83e-06: 100%| | 185/185 [00:01<00:00, 123.90it/s, loss=0.0043,
Época 97/200 [Train] L1=3.63e-06: 100%| | 185/185 [00:01<00:00, 121.86it/s, loss=0.0033,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 97 (Val MSE: 0.0033)

Época 98/200 [Train] L1=3.45e-06: 100%| | 185/185 [00:01<00:00, 116.56it/s, loss=0.0042,
Época 99/200 [Train] L1=3.28e-06: 100%| | 185/185 [00:01<00:00, 120.07it/s, loss=0.0042,
Época 100/200 [Train] L1=3.12e-06: 100%| | 185/185 [00:01<00:00, 124.46it/s, loss=0.0039

Época 100/200 | L1 coef: 3.12e-06

Train - Total: 0.0048 | MSE: 0.0038 | L0: 610.1

Val - Total: 0.0052 | MSE: 0.0042 | L0: 608.3

Best Val MSE: 0.0033 @ época 97 | Sin mejora: 3/15

Época 101/200 [Train] L1=2.96e-06: 100%| | 185/185 [00:01<00:00, 121.85it/s, loss=0.0028

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 101 (Val MSE: 0.0033)

Época 102/200 [Train] L1=2.81e-06: 100%| | 185/185 [00:01<00:00, 122.45it/s, loss=0.0030

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 102 (Val MSE: 0.0032)

Época 103/200 [Train] L1=2.67e-06: 100%| | 185/185 [00:01<00:00, 124.78it/s, loss=0.0037
Época 104/200 [Train] L1=2.54e-06: 100%| | 185/185 [00:01<00:00, 121.99it/s, loss=0.0035
Época 105/200 [Train] L1=2.41e-06: 100%| | 185/185 [00:01<00:00, 123.81it/s, loss=0.0028

Época 105/200 | L1 coef: 2.41e-06

Train - Total: 0.0033 | MSE: 0.0025 | L0: 608.6

Val - Total: 0.0044 | MSE: 0.0036 | L0: 606.3

Best Val MSE: 0.0032 @ época 102 | Sin mejora: 3/15

Época 106/200 [Train] L1=2.29e-06: 100%| | 185/185 [00:01<00:00, 124.44it/s, loss=0.0032

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 106 (Val MSE: 0.0031)

Época 107/200 [Train] L1=2.18e-06: 100%| | 185/185 [00:01<00:00, 121.01it/s, loss=0.0032

Época 108/200 [Train] L1=2.07e-06: 100%| | 185/185 [00:01<00:00, 119.33it/s, loss=0.0032

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 108 (Val MSE: 0.0031)

Época 109/200 [Train] L1=1.96e-06: 100%| | 185/185 [00:01<00:00, 124.27it/s, loss=0.0028

Época 110/200 [Train] L1=1.87e-06: 100%| | 185/185 [00:01<00:00, 121.71it/s, loss=0.0034

Época 110/200 | L1 coef: 1.87e-06

Train - Total: 0.0037 | MSE: 0.0031 | L0: 606.9

Val - Total: 0.0054 | MSE: 0.0047 | L0: 603.8

Best Val MSE: 0.0031 @ época 108 | Sin mejora: 2/15

Época 111/200 [Train] L1=1.77e-06: 100%| | 185/185 [00:01<00:00, 123.45it/s, loss=0.0034

Época 112/200 [Train] L1=1.68e-06: 100%| | 185/185 [00:01<00:00, 119.28it/s, loss=0.0023

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 112 (Val MSE: 0.0027)

Época 113/200 [Train] L1=1.60e-06: 100%| | 185/185 [00:01<00:00, 117.51it/s, loss=0.0024

Época 114/200 [Train] L1=1.52e-06: 100%| | 185/185 [00:01<00:00, 124.69it/s, loss=0.0023

Época 115/200 [Train] L1=1.44e-06: 100%| | 185/185 [00:01<00:00, 123.60it/s, loss=0.0038

Época 115/200 | L1 coef: 1.44e-06

Train - Total: 0.0028 | MSE: 0.0023 | L0: 605.6

Val - Total: 0.0043 | MSE: 0.0038 | L0: 602.8

Best Val MSE: 0.0027 @ época 112 | Sin mejora: 3/15

Época 116/200 [Train] L1=1.37e-06: 100%| | 185/185 [00:01<00:00, 123.84it/s, loss=0.0031

Época 117/200 [Train] L1=1.30e-06: 100%| | 185/185 [00:01<00:00, 123.00it/s, loss=0.0023

Época 118/200 [Train] L1=1.24e-06: 100%| | 185/185 [00:01<00:00, 123.17it/s, loss=0.0027

Época 119/200 [Train] L1=1.18e-06: 100%| | 185/185 [00:01<00:00, 121.49it/s, loss=0.0030

Época 120/200 [Train] L1=1.12e-06: 100%| | 185/185 [00:01<00:00, 125.36it/s, loss=0.0028

Época 120/200 | L1 coef: 1.12e-06
Train - Total: 0.0025 | MSE: 0.0021 | L0: 604.5
Val - Total: 0.0033 | MSE: 0.0029 | L0: 602.7
Best Val MSE: 0.0027 @ época 112 | Sin mejora: 8/15

Época 121/200 [Train] L1=1.06e-06: 100%| | 185/185 [00:01<00:00, 122.33it/s, loss=0.0026
Época 122/200 [Train] L1=1.01e-06: 100%| | 185/185 [00:01<00:00, 123.95it/s, loss=0.0030
Época 123/200 [Train] L1=9.58e-07: 100%| | 185/185 [00:01<00:00, 123.34it/s, loss=0.0027
Época 124/200 [Train] L1=9.10e-07: 100%| | 185/185 [00:01<00:00, 118.93it/s, loss=0.0024
Época 125/200 [Train] L1=8.64e-07: 100%| | 185/185 [00:01<00:00, 122.85it/s, loss=0.0026

Época 125/200 | L1 coef: 8.64e-07
Train - Total: 0.0027 | MSE: 0.0024 | L0: 604.5
Val - Total: 0.0033 | MSE: 0.0029 | L0: 603.4
Best Val MSE: 0.0027 @ época 112 | Sin mejora: 13/15

Época 126/200 [Train] L1=8.21e-07: 100%| | 185/185 [00:01<00:00, 122.74it/s, loss=0.0018
Época 127/200 [Train] L1=7.80e-07: 100%| | 185/185 [00:01<00:00, 121.81it/s, loss=0.0020

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt
Mejor modelo guardado en época 127 (Val MSE: 0.0026)

Época 128/200 [Train] L1=7.41e-07: 100%| | 185/185 [00:01<00:00, 118.73it/s, loss=0.0022
Época 129/200 [Train] L1=7.04e-07: 100%| | 185/185 [00:01<00:00, 121.53it/s, loss=0.0026
Época 130/200 [Train] L1=6.69e-07: 100%| | 185/185 [00:01<00:00, 122.41it/s, loss=0.0018

Época 130/200 | L1 coef: 6.69e-07
Train - Total: 0.0027 | MSE: 0.0025 | L0: 603.7
Val - Total: 0.0034 | MSE: 0.0031 | L0: 602.3
Best Val MSE: 0.0026 @ época 127 | Sin mejora: 3/15

Época 131/200 [Train] L1=6.35e-07: 100%| | 185/185 [00:01<00:00, 122.04it/s, loss=0.0020
Época 132/200 [Train] L1=6.04e-07: 100%| | 185/185 [00:01<00:00, 126.87it/s, loss=0.0020
Época 133/200 [Train] L1=5.73e-07: 100%| | 185/185 [00:01<00:00, 124.34it/s, loss=0.0036
Época 134/200 [Train] L1=5.45e-07: 100%| | 185/185 [00:01<00:00, 126.05it/s, loss=0.0021
Época 135/200 [Train] L1=5.18e-07: 100%| | 185/185 [00:01<00:00, 120.84it/s, loss=0.0016

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_best.pt

Mejor modelo guardado en época 135 (Val MSE: 0.0022)

Época 135/200 | L1 coef: 5.18e-07

Train - Total: 0.0021 | MSE: 0.0019 | L0: 602.8

Val - Total: 0.0024 | MSE: 0.0022 | L0: 600.6

Best Val MSE: 0.0022 @ época 135 | Sin mejora: 0/15

Época 136/200 [Train] L1=4.92e-07: 100%| | 185/185 [00:01<00:00, 125.59it/s, loss=0.0020

Época 137/200 [Train] L1=4.67e-07: 100%| | 185/185 [00:01<00:00, 121.70it/s, loss=0.0020

Época 138/200 [Train] L1=4.44e-07: 100%| | 185/185 [00:01<00:00, 118.57it/s, loss=0.0020

Época 139/200 [Train] L1=4.22e-07: 100%| | 185/185 [00:01<00:00, 123.56it/s, loss=0.0018

Época 140/200 [Train] L1=4.00e-07: 100%| | 185/185 [00:01<00:00, 124.80it/s, loss=0.0028

Época 140/200 | L1 coef: 4.00e-07

Train - Total: 0.0023 | MSE: 0.0022 | L0: 600.4

Val - Total: 0.0027 | MSE: 0.0026 | L0: 599.2

Best Val MSE: 0.0022 @ época 135 | Sin mejora: 5/15

Época 141/200 [Train] L1=3.80e-07: 100%| | 185/185 [00:01<00:00, 123.71it/s, loss=0.0025

Época 142/200 [Train] L1=3.61e-07: 100%| | 185/185 [00:01<00:00, 110.61it/s, loss=0.0024

Época 143/200 [Train] L1=3.43e-07: 100%| | 185/185 [00:01<00:00, 104.40it/s, loss=0.0016

Época 144/200 [Train] L1=3.26e-07: 100%| | 185/185 [00:01<00:00, 108.41it/s, loss=0.0021

Época 145/200 [Train] L1=3.10e-07: 100%| | 185/185 [00:01<00:00, 110.53it/s, loss=0.0020

Época 145/200 | L1 coef: 3.10e-07

Train - Total: 0.0022 | MSE: 0.0021 | L0: 599.8

Val - Total: 0.0037 | MSE: 0.0035 | L0: 598.8

Best Val MSE: 0.0022 @ época 135 | Sin mejora: 10/15

Época 146/200 [Train] L1=2.94e-07: 100%| | 185/185 [00:01<00:00, 111.15it/s, loss=0.0017

Época 147/200 [Train] L1=2.80e-07: 100%| | 185/185 [00:01<00:00, 106.02it/s, loss=0.0022

Época 148/200 [Train] L1=2.66e-07: 100%| | 185/185 [00:01<00:00, 108.16it/s, loss=0.0025

Época 149/200 [Train] L1=2.52e-07: 100%| | 185/185 [00:01<00:00, 108.55it/s, loss=0.0017

Época 150/200 [Train] L1=2.40e-07: 100%| | 185/185 [00:01<00:00, 110.83it/s, loss=0.0017

Early stopping activado en época 150
Mejor modelo: época 135 con Val MSE = 0.0022

```
=====
RESUMEN DEL ENTRENAMIENTO
=====
Mejor modelo en época 135 con Val MSE = 0.0022
Early stopping tras 150 épocas

Métricas Finales (Última Época):
      Métrica  Train Validation
Total Loss 0.0020      0.0031
MSE Loss 0.0019      0.0030
L0 Sparsity 600.5      598.8
Fraction Active 3.67%    3.65%
=====
```

8. Guardar Modelo

```
# Guardar checkpoint final (local + organizado)
checkpoint = {
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'config': CONFIG,
    'history': history
}

# Guardado local
final_model_name = get_model_name(
    variante=NAMING_META['variante'],
    tecnica=NAMING_META['tecnica'],
    capa=NAMING_META['capa'],
    juegos=NAMING_META['juegos'],
    estado='final'
)
checkpoint_path = CONFIG['save_dir'] / final_model_name
torch.save(checkpoint, checkpoint_path)

# Guardado organizado
save_checkpoint(
```

```

    model,
    base_path=NAMING_META['base_path'],
    variante=NAMING_META['variante'],
    tecnica=NAMING_META['tecnica'],
    capa=NAMING_META['capa'],
    juegos=NAMING_META['juegos'],
    estado='final',
    extras=NAMING_META['extras'],
    config=CONFIG,
    history=history,
    optimizer=optimizer
)

# Resumen de archivos guardados
saved_files = pd.DataFrame({
    'Tipo': ['Local', 'Organizado'],
    'Ubicación': [
        str(checkpoint_path),
        f"layer_{NAMING_META['capa']:02d}/models/..."
    ]
})

print('\n' + '='*70)
print(' CHECKPOINTS GUARDADOS')
print('='*70)
print(saved_files.to_string(index=False))
print('='*70)

```

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\500games\sae_standard_l1-decay_l6_500g_ex3_final.pt

```

=====
CHECKPOINTS GUARDADOS
=====

```

Tipo	Ubicación
Local	saved_model\sae_standard_l1-decay_l6_500g_final.pt
Organizado	layer_06/models/...

```

=====

```

9. Métricas básicas del modelo

```
# Guardar metricas de entrenamiento
metrics_dir = get_metrics_dir(
    NAMING_META["experiment_base_path"],
    NAMING_META["variante"],
    NAMING_META["tecnica"],
    NAMING_META["capa"],
    NAMING_META["juegos"],
    extras_id=extras_id
)

training_metrics = {
    # Metricas de calidad
    "val_mse_final": history["val_mse_loss"][-1],
    "val_mse_best": best_val_mse,
    "val_l0_final": history["val_l0_sparsity"][-1],
    "val_fraction_active_final": history["val_fraction_active"][-1],
    # Contexto del entrenamiento
    "best_epoch": best_epoch,
    "total_epochs": len(history["train_total_loss"]),
    "train_mse_final": history["train_mse_loss"][-1],
    "train_l0_final": history["train_l0_sparsity"][-1],
    # Hiperparametros
    "hidden_dim": CONFIG["hidden_dim"],
    "expansion_factor": CONFIG["hidden_dim"] / CONFIG["input_dim"],
    "sparsity_coef": CONFIG["sparsity_coef"],
    "sparsity_decay": CONFIG["sparsity_decay"],
    "learning_rate": CONFIG["learning_rate"],
    "num_games": NAMING_META["juegos"],
}

metrics_path = metrics_dir / "training_metrics.json"
with open(metrics_path, "w") as f:
    _json.dump(training_metrics, f, indent=2)

# Resumen
metrics_df = pd.DataFrame({
    "Metrica": ["Val MSE (final)", "Val MSE (mejor)", "Val L0 (final)", "Fraccion Activa"],
    "Valor": [
        f"{training_metrics['val_mse_final']:.6f}",
        f"{training_metrics['val_mse_best']:.6f}",

```

```

        f"{training_metrics['val_l0_final']:.1f}",
        f"{training_metrics['val_fraction_active_final']:.2%}"
    ]
})
print(" " + "="*50)
print(" METRICAS DE ENTRENAMIENTO GUARDADAS")
print("="*50)
print(metrics_df.to_string(index=False))
print(f"Archivo: {metrics_path}")
print("="*50)

```

```

=====
METRICAS DE ENTRENAMIENTO GUARDADAS
=====

```

	Metrica	Valor
Val MSE (final)	0.002974	
Val MSE (mejor)	0.002162	
Val L0 (final)	598.8	
Fraccion Activa	3.65%	
Archivo:	c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\decay_standard\500games\ex3\metrics\training_metrics.json	

```

=====

```

10. Visualización de Resultados

```

# Gráficas de pérdidas (Train vs Validation) + L1 Decay
fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# Total loss
axes[0, 0].plot(history['train_total_loss'], label='Train', alpha=0.8)
axes[0, 0].plot(history['val_total_loss'], label='Validation', alpha=0.8)
axes[0, 0].set_title('Total Loss')
axes[0, 0].set_xlabel('Época')
axes[0, 0].set_ylabel('Loss')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)

# MSE loss
axes[0, 1].plot(history['train_mse_loss'], label='Train', color='orange', alpha=0.8)
axes[0, 1].plot(history['val_mse_loss'], label='Validation', color='red', alpha=0.8)

```

```

if CONFIG['early_stopping'] and best_epoch > 0:
    axes[0, 1].axvline(best_epoch - 1, color='green', linestyle='--', alpha=0.7, label=f'Best')
axes[0, 1].set_title('MSE Loss (Reconstrucción)')
axes[0, 1].set_xlabel('Época')
axes[0, 1].set_ylabel('MSE')
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)

# L0 sparsity
axes[0, 2].plot(history['train_l0_sparsity'], label='Train', color='green', alpha=0.8)
axes[0, 2].plot(history['val_l0_sparsity'], label='Validation', color='darkgreen', alpha=0.8)
axes[0, 2].set_title('L0 Sparsity (Features Activas)')
axes[0, 2].set_xlabel('Época')
axes[0, 2].set_ylabel('Número de features')
axes[0, 2].legend()
axes[0, 2].grid(True, alpha=0.3)

# Fraction active
axes[1, 0].plot(history['train_fraction_active'], label='Train', color='purple', alpha=0.8)
axes[1, 0].plot(history['val_fraction_active'], label='Validation', color='magenta', alpha=0.8)
axes[1, 0].set_title('Fracción de Features Activas')
axes[1, 0].set_xlabel('Época')
axes[1, 0].set_ylabel('Fracción')
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)

# Sparsity Coefficient (L1 Decay)
axes[1, 1].plot(history['sparsity_coef_history'], color='brown', linewidth=2)
axes[1, 1].set_title('L1 Coefficient (Decay)')
axes[1, 1].set_xlabel('Época')
axes[1, 1].set_ylabel('Sparsity Coef')
axes[1, 1].set_yscale('log')
axes[1, 1].grid(True, alpha=0.3)

# Train vs Val MSE Gap
mse_gap = [v - t for v, t in zip(history['val_mse_loss'], history['train_mse_loss'])]
axes[1, 2].plot(mse_gap, color='crimson', linewidth=2)
axes[1, 2].axhline(0, color='black', linestyle='--', alpha=0.5)
if CONFIG['early_stopping'] and best_epoch > 0:
    axes[1, 2].axvline(best_epoch - 1, color='green', linestyle='--', alpha=0.7, label=f'Best')
axes[1, 2].set_title('Overfitting Gap (Val MSE - Train MSE)')
axes[1, 2].set_xlabel('Época')

```



```


axes[1, 2].set_ylabel('Gap MSE')
axes[1, 2].legend()
axes[1, 2].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(CONFIG['plots_dir'] / 'training_curves.png', dpi=300, bbox_inches='tight')
plt.show()

print('\n' + '='*50)
print('  Gráficas guardadas exitosamente')

print('='*50)

```



=====
 Gráficas guardadas exitosamente
 =====

11. Resumen Final

```

# Resumen final completo
print('\n' + '='*70)
print(' RESUMEN FINAL DEL ENTRENAMIENTO')
print('='*70)

# Arquitectura
arch_summary = pd.DataFrame({
    'Componente': ['Input', 'Hidden', 'Parámetros'],
    'Valor': [
        f"{CONFIG['input_dim']}",
        f"{CONFIG['hidden_dim']} ({CONFIG['hidden_dim']/CONFIG['input_dim']:.1f}x)",
        f"{num_params:}"
    ]
})
print('\nArquitectura:')
print(arch_summary.to_string(index=False))

# Datos
data_summary = pd.DataFrame({
    'Conjunto': ['Train', 'Validation', 'Total'],
    'Muestras': [
        f"{len(train_dataset):}",
        f"{len(val_dataset):}",
        f"{len(full_dataset):}"
    ]
})
print('\nDatos:')
print(data_summary.to_string(index=False))

# Métricas finales
final_summary = pd.DataFrame({
    'Métrica': ['Train Loss', 'Val Loss', 'Train MSE', 'Val MSE', 'Train L0', 'Val L0'],
    'Valor': [
        f"{history['train_total_loss'][-1]:.4f}",
        f"{history['val_total_loss'][-1]:.4f}",
        f"{history['train_mse_loss'][-1]:.4f}",
        f"{history['val_mse_loss'][-1]:.4f}",
        f"{history['train_l0_sparsity'][-1]:.1f}",
        f"{history['val_l0_sparsity'][-1]:.1f}"
    ]
})
print('\nMétricas Finales:')

```

```

print(final_summary.to_string(index=False))

# Archivos
files_summary = pd.DataFrame({
    'Tipo': ['Modelo', 'Gráficas'],
    'Ubicación': [
        str(checkpoint_path),
        str(CONFIG['plots_dir'] / 'training_curves.png')
    ]
})
print('\nArchivos Guardados:')
print(files_summary.to_string(index=False))

print('\n' + '='*70)
print('  ENTRENAMIENTO COMPLETADO EXITOSAMENTE')
print('='*70)

```

```

=====
RESUMEN FINAL DEL ENTRENAMIENTO
=====

```

Arquitectura:

Componente	Valor
Input	512
Hidden	16384 (32.0x)
Parámetros	16,794,112

Datos:

Conjunto	Muestras
Train	23,600
Validation	5,900
Total	29,500

Métricas Finales:

Métrica	Valor
Train Loss	0.0020
Val Loss	0.0031
Train MSE	0.0019
Val MSE	0.0030
Train L0	600.5
Val L0	598.8

Archivos Guardados:

Tipo

Modelo

decay_l6_500g_final.pt

Gráficas c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard

decay_standard\500games\ex3\plots\training_curves.png

```
=====
ENTRENAMIENTO COMPLETADO EXITOSAMENTE
=====
```

12. Generar PDF con Quarto

Una vez ejecutado todo el notebook, ejecutar el siguiente comando en la terminal para generar el PDF con Quarto:

```
import subprocess

# Obtener directorio del experimento y nombre del PDF
exp_dir = get_experiment_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

pdf_name = get_report_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    'training',
    extras_id=extras_id
)

notebook_name = 'train_sae_standard_l1-decay.ipynb'
output_path = exp_dir / pdf_name

print(f' Generando PDF con Quarto...')
```

```

print(f' Archivo de salida: {output_path}')

# Quarto no acepta rutas en --output, solo nombre de archivo
# Se usa --output-dir para el directorio y --output solo para el nombre
result = subprocess.run(
    f'quarto render {notebook_name} --to pdf --output-dir "{exp_dir}" --output {pdf_name}',
    shell=True,
    capture_output=True,
    text=True
)

if result.returncode == 0:
    print(f' PDF generado exitosamente: {output_path}')
else:
    print(f' Error al generar PDF:')
    print(result.stderr)

```

Generando PDF con Quarto...

Archivo de salida: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\500games\ex3\sae_standard_l1-decay_l6_500g_ex3_training.pdf

PDF generado exitosamente: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\500games\ex3\sae_standard_l1-decay_l6_500g_ex3_training.pdf