

1. Setup e Imports

```
import json as _json
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import os
import sys
from pathlib import Path

# Agregar path del proyecto para imports
sys.path.append(os.path.abspath('../..'))
from sae.models.sae import SparseAutoencoder
from sae.tools.naming_utils import save_checkpoint, get_model_name, get_extras_id, get_checkpoint_dir, get_model_dir
from sae.tools.experiment_utils import get_experiment_dir, get_plots_dir, get_metrics_dir, get_logs_dir

# Configuración de visualización para Quarto/PDF
pd.set_option('display.width', 100)
pd.set_option('display.max_columns', None)
np.set_printoptions(linewidth=100, precision=4)
os.environ['COLUMNS'] = '100'

# Configuración de reproducibilidad
torch.manual_seed(42)
np.random.seed(42)

# Device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
# Print config con formato profesional
config_df = pd.DataFrame({
    'Configuración': ['Device', 'Seed', 'Output Width'],
    'Valor': [str(device), '42', '100']
})

print('\n' + '='*50)
print(' CONFIGURACIÓN INICIAL')
print('='*50)
print(config_df.to_string(index=False))
print('='*50)
```

```
=====
CONFIGURACIÓN INICIAL
=====
Configuración Valor
      Device  cuda
      Seed    42
Output Width  100
=====
```

2 Metadata para Naming de Checkpoints

```
# Metadata para naming de checkpoints (separada de hiperparámetros)
NAMING_META = {
    'variante': 'standard',          # Arquitectura del SAE
    'tecnica': 'l1-decay',           # L1 con penalty annealing
    'capa': 6,                       # Layer del modelo OthelloGPT
    'juegos': 1500,                  # Número de partidas
    'base_path': 'C:\\Users\\Esposa\\Documents\\Repos\\sae-othello-gpt', # Path base para cl
    'experiment_base_path': os.path.abspath('../../experiments') ,      # sae/experiments,
    'extras': {
        'expansion': 32,
        'sparsity_coef': 1e-2,
        'epochs': 1000,
        'patience': 20,
        'learning_rate': 1e-3,
        'sparsity_decay': 0.98,
```

```

        'batch_size': 256
    }
}

print('\n Metadata de naming:')
for key, value in NAMING_META.items():
    print(f' {key}: {value}')

```

```

Metadata de naming:
variante: standard
tecnica: l1-decay
capa: 6
juegos: 1500
base_path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt
experiment_base_path: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments
extras: {'expansion': 32, 'sparsity_coef': 0.01, 'epochs': 1000, 'patience': 20, 'learning

```

```

extras_id = get_extras_id(
    get_checkpoint_dir(
        NAMING_META['base_path'],
        NAMING_META['variante'],
        NAMING_META['tecnica'],
        NAMING_META['capa'],
        NAMING_META['juegos']
    ),
    NAMING_META['extras']
) if NAMING_META.get('extras') else None

print(f'extras_id: {extras_id}')

```

Nuevo extras registrado: ex1 → {'expansion': 32, 'sparsity_coef': 0.01, 'epochs': 1000, 'pa
 extras_id: ex1

3. Configuración del SAE

Las activaciones ya fueron extraídas previamente con `sae/activations/run_extraction.py`. Aquí especificamos qué archivo cargar y los hiperparámetros del SAE.

```

# Configuración del SAE
CONFIG = {
    # Arquitectura
    'input_dim': 512,
    'hidden_dim': 16384,
    'tied_weights': False,

    # Entrenamiento
    'batch_size': 256,
    'num_epochs': 1000,
    'learning_rate': 1e-3,
    'sparsity_coef': 1e-2,
    'sparsity_decay': 0.98,

    # Early Stopping
    'early_stopping': True,
    'patience': 20,

    # Datos
    'activations_file': Path(f"../../activations/data/layer{NAMING_META['capa']}-{NAMING_META['juegos']}.npy"),
    'save_dir': Path('./saved_model'),
    'plots_dir': get_plots_dir(
        NAMING_META['experiment_base_path'],
        NAMING_META['variante'],
        NAMING_META['tecnica'],
        NAMING_META['capa'],
        NAMING_META['juegos'],
        extras_id=extras_id
    )
}

# Crear directorio local del modelo (plots_dir ya fue creado por get_plots_dir)
CONFIG['save_dir'].mkdir(parents=True, exist_ok=True)

# Mostrar configuración con formato profesional
config_data = {
    'Categoría': [],
    'Parámetro': [],
    'Valor': []
}

```

```

# Arquitectura
for key in ['input_dim', 'hidden_dim', 'tied_weights']:
    config_data['Categoría'].append('Arquitectura')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

# Entrenamiento
for key in ['batch_size', 'num_epochs', 'learning_rate', 'sparsity_coef', 'sparsity_decay']:
    config_data['Categoría'].append('Entrenamiento')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

# Early Stopping
for key in ['early_stopping', 'patience']:
    config_data['Categoría'].append('Early Stopping')
    config_data['Parámetro'].append(key)
    config_data['Valor'].append(str(CONFIG[key]))

config_table = pd.DataFrame(config_data)

print('\n' + '='*70)
print(' CONFIGURACIÓN DEL SAE')
print('='*70)
print(config_table.to_string(index=False))
print('='*70)

print(f'\nArchivo de activaciones: {CONFIG["activations_file"]}')
print(f'Expansión: {CONFIG["hidden_dim"]/CONFIG["input_dim"]:.1f}x')
print(f'Plots dir: {CONFIG["plots_dir"]}')
print('='*70)

```

```

=====
CONFIGURACIÓN DEL SAE
=====

```

Categoría	Parámetro	Valor
Arquitectura	input_dim	512
Arquitectura	hidden_dim	16384
Arquitectura	tied_weights	False
Entrenamiento	batch_size	256
Entrenamiento	num_epochs	1000
Entrenamiento	learning_rate	0.001

```

Entrenamiento sparsity_coef 0.01
Entrenamiento sparsity_decay 0.98
Early Stopping early_stopping True
Early Stopping patience 20
=====

Archivo de activaciones: ..\..\..\activations\data\layer6_1500games.npy
Expansión: 32.0x
Plots dir: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\1500games\ex1\plots
=====

```

4. Dataset de Activaciones (Train/Validation Split 80-20%)

```

class ActivationsDataset(Dataset):
    """Dataset para cargar activaciones extraídas"""

    def __init__(self, activations_path):
        self.activations = np.load(activations_path)

    def __len__(self):
        return len(self.activations)

    def __getitem__(self, idx):
        return torch.tensor(self.activations[idx], dtype=torch.float32)

# Cargar dataset completo
full_dataset = ActivationsDataset(CONFIG['activations_file'])

# Información del dataset con formato profesional
dataset_info = pd.DataFrame({
    'Propiedad': ['Total muestras', 'Shape', 'Dtype', 'Tamaño en memoria'],
    'Valor': [
        f"{len(full_dataset):,}",
        str(full_dataset.activations.shape),
        str(full_dataset.activations.dtype),
        f"{full_dataset.activations.nbytes / 1024**2:.2f} MB"
    ]
})

print('\n' + '='*50)
print(' DATASET DE ACTIVACIONES')

```

```

print('='*50)
print(dataset_info.to_string(index=False))
print('='*50)

# Split train/validation (80-20)
train_size = int(0.8 * len(full_dataset))
val_size = len(full_dataset) - train_size
train_dataset, val_dataset = torch.utils.data.random_split(
    full_dataset,
    [train_size, val_size],
    generator=torch.Generator().manual_seed(42)
)

# Resumen del split con formato profesional
split_info = pd.DataFrame({
    'Conjunto': ['Train', 'Validation', 'Total'],
    'Muestras': [
        f"{len(train_dataset):,}",
        f"{len(val_dataset):,}",
        f"{len(full_dataset):,}"
    ],
    'Porcentaje': [
        f"{len(train_dataset)/len(full_dataset):.1%}",
        f"{len(val_dataset)/len(full_dataset):.1%}",
        "100.0%"
    ]
})
print('\n' + '='*50)
print(' SPLIT DE DATOS (80-20)')
print('='*50)
print(split_info.to_string(index=False))
print('='*50)

# Crear DataLoaders
train_loader = DataLoader(
    train_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=True,
    num_workers=0,
    pin_memory=(device.type == 'cuda')
)

```

```

val_loader = DataLoader(
    val_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=False, # No shuffle para validación
    num_workers=0,
    pin_memory=(device.type == 'cuda')
)

# Resumen de DataLoaders
loader_info = pd.DataFrame({
    'DataLoader': ['Train', 'Validation'],
    'Batches': [len(train_loader), len(val_loader)],
    'Batch Size': [CONFIG['batch_size'], CONFIG['batch_size']],
    'Shuffle': ['Sí', 'No']
})

print('\n' + '='*50)
print(' DATALOADERS')
print('='*50)
print(loader_info.to_string(index=False))
print('='*50)

```

```

=====
DATASET DE ACTIVACIONES
=====

```

	Propiedad	Valor
	Total muestras	88,500
	Shape	(88500, 512)
	Dtype	float32
Tamaño en memoria		172.85 MB

```

=====

=====
SPLIT DE DATOS (80-20)
=====

```

	Conjunto	Muestras	Porcentaje
	Train	70,800	80.0%
Validation		17,700	20.0%
	Total	88,500	100.0%

```

=====

```


=====				
DATALOADERS				
=====				
Loader	Batches	Batch Size	Shuffle	
Train	277	256	Sí	
Validation	70	256	No	
=====				

5. Modelo: Sparse Autoencoder

Modelo importado desde `sae.models.sae.SparseAutoencoder`

Arquitectura:

- **Encoder:** Linear $\rightarrow$ ReLU (sparsity natural)- **Decoder:** Linear (reconstrucción)

```
# Crear modelo desde el módulo
model = SparseAutoencoder(
    input_dim=CONFIG['input_dim'],
    hidden_dim=CONFIG['hidden_dim'],
    tied_weights=CONFIG['tied_weights']
).to(device)

# Calcular parámetros
num_params = sum(p.numel() for p in model.parameters())
encoder_params = sum(p.numel() for p in model.encoder.parameters())
if CONFIG['tied_weights']:
    decoder_params = CONFIG['input_dim'] # Solo bias
else:
    decoder_params = sum(p.numel() for p in model.decoder.parameters())

# Resumen del modelo con formato profesional
model_info = pd.DataFrame({
    'Componente': ['Input', 'Encoder', 'Hidden (Latent)', 'Decoder', 'Output', 'Total'],
    'Dimensión': [
        CONFIG['input_dim'],
        f"{CONFIG['input_dim']} → {CONFIG['hidden_dim']}",
        CONFIG['hidden_dim'],
        f"{CONFIG['hidden_dim']} → {CONFIG['input_dim']}",
        CONFIG['input_dim'],
        '_'
    ],
    'Parámetros': [
```

```

        '-',
        f"{encoder_params: ,}",
        '-',
        f"{decoder_params: ,}",
        '-',
        f"{num_params: ,}"
    ]
})

print('\n' + '='*70)
print(' ARQUITECTURA DEL MODELO')
print('='*70)
print(model_info.to_string(index=False))
print('='*70)
print(f'\n Factor de expansión: {CONFIG["hidden_dim"]/CONFIG["input_dim"]:.1f}x')
print(f' Tied weights: {"Sí" if CONFIG["tied_weights"] else "No"}')
print(f' Device: {device}')
print('='*70)

```

```

=====
ARQUITECTURA DEL MODELO
=====

```

	Componente	Dimensión	Parámetros
	Input	512	-
	Encoder	512 → 16384	8,404,992
Hidden	(Latent)	16384	-
	Decoder	16384 → 512	8,389,120
	Output	512	-
	Total	-	16,794,112

```

=====

Factor de expansión: 32.0x
Tied weights: No
Device: cuda
=====

```

6. Función de Pérdida

Loss = MSE(reconstrucción) + λ × L1(activaciones)

- **MSE:** Mide qué tan bien reconstruimos las activaciones originales
- **L1:** Penaliza activaciones grandes → fuerza sparsity

```
def loss_function(x, reconstruction, hidden, sparsity_coef):
    """
    Pérdida compuesta: MSE + L1 penalty

    Args:
        x: Activaciones originales
        reconstruction: Activaciones reconstruidas
        hidden: Activaciones latentes (sparse)
        sparsity_coef: Peso de la penalización L1
    """
    # Error de reconstrucción
    mse_loss = F.mse_loss(reconstruction, x)

    # Sumar sobre las features (dim=1), promediar sobre el batch
    sparsity_loss = hidden.abs().sum(dim=1).mean()

    # Pérdida total
    total_loss = mse_loss + sparsity_coef * sparsity_loss

    return total_loss, mse_loss, sparsity_loss

# Optimizador
optimizer = torch.optim.Adam(model.parameters(), lr=CONFIG['learning_rate'])

# Resumen del optimizador
optim_info = pd.DataFrame({
    'Parámetro': ['Optimizador', 'Learning Rate', 'Sparsity Coef (inicial)', 'Sparsity Decay'],
    'Valor': [
        'Adam',
        f"{CONFIG['learning_rate']:.0e}",
        f"{CONFIG['sparsity_coef']:.0e}",
        f"{CONFIG['sparsity_decay']:.2f}"
    ]
})

print('\n' + '='*50)
print(' OPTIMIZACIÓN')
print('='*50)
print(optim_info.to_string(index=False))
print('='*50)
```

```
=====
OPTIMIZACIÓN
=====
                Parámetro Valor
            Optimizador  Adam
        Learning Rate 1e-03
Sparsity Coef (inicial) 1e-02
        Sparsity Decay  0.98
=====
```

7. Loop de Entrenamiento

```
# Historial de entrenamiento
history = {
    'train_total_loss': [],
    'train_mse_loss': [],
    'train_sparsity_loss': [],
    'train_l0_sparsity': [],
    'train_fraction_active': [],
    'val_total_loss': [],
    'val_mse_loss': [],
    'val_sparsity_loss': [],
    'val_l0_sparsity': [],
    'val_fraction_active': [],
    'sparsity_coef_history': [] # Guardar evolución del coeficiente L1
}

# Early stopping variables
best_val_mse = float('inf')
best_epoch = 0
epochs_no_improve = 0

print(f'\n Iniciando entrenamiento: {CONFIG["num_epochs"]} épocas máximo')
print(f' L1 decay: {CONFIG["sparsity_coef"]:.2e} × {CONFIG["sparsity_decay"]:.2f}^época')
print(f' Early stopping: patience={CONFIG["patience"]} épocas')
print('=' * 70)

for epoch in range(CONFIG['num_epochs']):
```

```

# ===== L1 DECAY =====
# Reducir sparsity_coef exponencialmente
current_sparsity_coef = CONFIG['sparsity_coef'] * (CONFIG['sparsity_decay'] ** epoch)
history['sparsity_coef_history'].append(current_sparsity_coef)

# ===== ENTRENAMIENTO =====
model.train()

epoch_total_loss = 0
epoch_mse_loss = 0
epoch_sparsity_loss = 0
epoch_l0 = 0
epoch_fraction_active = 0

# Barra de progreso
pbar = tqdm(train_loader, desc=f'Época {epoch+1}/{CONFIG["num_epochs"]} [Train] L1={current_sparsity_coef}')

for batch in pbar:
    # Mover al device
    x = batch.to(device)

    # Forward pass
    reconstruction, hidden = model(x)

    # Calcular pérdida con sparsity_coef actualizado
    total_loss, mse_loss, sparsity_loss = loss_function(
        x, reconstruction, hidden, current_sparsity_coef
    )

    # Backward pass
    optimizer.zero_grad()
    total_loss.backward()

    optimizer.step()

    # Métricas
    with torch.no_grad():
        l0 = (hidden > 0).float().sum(dim=1).mean().item()
        fraction_active = (hidden > 0).float().mean().item()

    # Acumular
    epoch_total_loss += total_loss.item()

```

```

epoch_mse_loss += mse_loss.item()
epoch_sparsity_loss += sparsity_loss.item()
epoch_l0 += l0
epoch_fraction_active += fraction_active

# Actualizar barra
pbar.set_postfix({
    'loss': f'{total_loss.item():.4f}',
    'mse': f'{mse_loss.item():.4f}',
    'l0': f'{l0:.1f}'
})

# Promedios de entrenamiento
num_train_batches = len(train_loader)
avg_train_total_loss = epoch_total_loss / num_train_batches
avg_train_mse_loss = epoch_mse_loss / num_train_batches
avg_train_sparsity_loss = epoch_sparsity_loss / num_train_batches
avg_train_l0 = epoch_l0 / num_train_batches
avg_train_fraction_active = epoch_fraction_active / num_train_batches

# ===== VALIDACIÓN =====
model.eval()

val_total_loss = 0
val_mse_loss = 0
val_sparsity_loss = 0
val_l0 = 0
val_fraction_active = 0

with torch.no_grad():
    for batch in val_loader:
        x = batch.to(device)

        # Forward pass
        reconstruction, hidden = model(x)

        # Calcular pérdida con sparsity_coef actualizado
        total_loss, mse_loss, sparsity_loss = loss_function(
            x, reconstruction, hidden, current_sparsity_coef
        )

# Métricas

```

```

10 = (hidden > 0).float().sum(dim=1).mean().item()
fraction_active = (hidden > 0).float().mean().item()

# Acumular
val_total_loss += total_loss.item()
val_mse_loss += mse_loss.item()
val_sparsity_loss += sparsity_loss.item()
val_l0 += l0
val_fraction_active += fraction_active

# Promedios de validación
num_val_batches = len(val_loader)
avg_val_total_loss = val_total_loss / num_val_batches
avg_val_mse_loss = val_mse_loss / num_val_batches
avg_val_sparsity_loss = val_sparsity_loss / num_val_batches
avg_val_l0 = val_l0 / num_val_batches
avg_val_fraction_active = val_fraction_active / num_val_batches

# Guardar en historial
history['train_total_loss'].append(avg_train_total_loss)
history['train_mse_loss'].append(avg_train_mse_loss)
history['train_sparsity_loss'].append(avg_train_sparsity_loss)
history['train_l0_sparsity'].append(avg_train_l0)
history['train_fraction_active'].append(avg_train_fraction_active)

history['val_total_loss'].append(avg_val_total_loss)
history['val_mse_loss'].append(avg_val_mse_loss)
history['val_sparsity_loss'].append(avg_val_sparsity_loss)
history['val_l0_sparsity'].append(avg_val_l0)
history['val_fraction_active'].append(avg_val_fraction_active)

# ===== EARLY STOPPING =====
if CONFIG['early_stopping']:
    if avg_val_mse_loss < best_val_mse:
        best_val_mse = avg_val_mse_loss
        best_epoch = epoch + 1
        epochs_no_improve = 0

    # Guardar mejor modelo (local + organizado)
    best_checkpoint = {
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),

```

```

        'config': CONFIG,
        'history': history,
        'epoch': best_epoch,
        'val_mse': best_val_mse
    }

    # Guardado local
    best_model_name = get_model_name(
        variante=NAMING_META['variante'],
        tecnica=NAMING_META['tecnica'],
        capa=NAMING_META['capa'],
        juegos=NAMING_META['juegos'],
        estado='best'
    )
    torch.save(best_checkpoint, CONFIG['save_dir'] / best_model_name)

    # Guardado organizado
    save_checkpoint(
        model,
        base_path=NAMING_META['base_path'],
        variante=NAMING_META['variante'],
        tecnica=NAMING_META['tecnica'],
        capa=NAMING_META['capa'],
        juegos=NAMING_META['juegos'],
        estado='best',
        extras=NAMING_META['extras'],
        config=CONFIG,
        history=history,
        optimizer=optimizer
    )

    print(f' Mejor modelo guardado en época {best_epoch} (Val MSE: {best_val_mse:.4f})')
else:
    epochs_no_improve += 1
    if epochs_no_improve >= CONFIG['patience']:
        print(f'\n Early stopping activado en época {epoch+1}')
        print(f' Mejor modelo: época {best_epoch} con Val MSE = {best_val_mse:.4f}')
        break

# Log cada 5 épocas
if (epoch + 1) % 5 == 0 or epoch == 0:
    print(f'\n Época {epoch+1}/{CONFIG["num_epochs"]} | L1 coef: {current_sparsity_coef:}')

```



```

        print(f'  Train - Total: {avg_train_total_loss:.4f} | MSE: {avg_train_mse_loss:.4f}')
        print(f'  Val   - Total: {avg_val_total_loss:.4f} | MSE: {avg_val_mse_loss:.4f} | L0')
        if CONFIG['early_stopping']:
            print(f'    Best Val MSE: {best_val_mse:.4f} @ época {best_epoch} | Sin mejora: {best_val_mse:.4f}')

print('\n' + '='*70)
print(' RESUMEN DEL ENTRENAMIENTO')
print('='*70)

if CONFIG['early_stopping']:
    print(f'Mejor modelo en época {best_epoch} con Val MSE = {best_val_mse:.4f}')
    print(f'Early stopping tras {len(history["train_total_loss"])} épocas')
else:
    print(f'Entrenamiento completo: {CONFIG["num_epochs"]} épocas')

# Tabla de métricas finales
final_metrics = pd.DataFrame({
    'Métrica': ['Total Loss', 'MSE Loss', 'L0 Sparsity', 'Fraction Active'],
    'Train': [
        f"{history['train_total_loss'][-1]:.4f}",
        f"{history['train_mse_loss'][-1]:.4f}",
        f"{history['train_l0_sparsity'][-1]:.1f}",
        f"{history['train_fraction_active'][-1]:.2%}"
    ],
    'Validation': [
        f"{history['val_total_loss'][-1]:.4f}",
        f"{history['val_mse_loss'][-1]:.4f}",
        f"{history['val_l0_sparsity'][-1]:.1f}",
        f"{history['val_fraction_active'][-1]:.2%}"
    ]
})

print('\nMétricas Finales (Última Época):')
print(final_metrics.to_string(index=False))
print('='*70)

```

Iniciando entrenamiento: 1000 épocas máximo
 L1 decay: $1.00e-02 \times 0.98^{\text{época}}$
 Early stopping: patience=20 épocas

=====

Época 1/1000 [Train] L1=1.00e-02: 100%| | 277/277 [00:03<00:00, 85.38it/s, loss=0.6419,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 1 (Val MSE: 0.5908)

Época 1/1000 | L1 coef: 1.00e-02

Train - Total: 0.8375 | MSE: 0.6668 | L0: 102.4

Val - Total: 0.6321 | MSE: 0.5908 | L0: 4.1

Best Val MSE: 0.5908 @ época 1 | Sin mejora: 0/20

Época 2/1000 [Train] L1=9.80e-03: 100%| | 277/277 [00:02<00:00, 94.09it/s, loss=0.5819,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 2 (Val MSE: 0.5616)

Época 3/1000 [Train] L1=9.60e-03: 100%| | 277/277 [00:03<00:00, 92.23it/s, loss=0.5813,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 3 (Val MSE: 0.5467)

Época 4/1000 [Train] L1=9.41e-03: 100%| | 277/277 [00:03<00:00, 90.45it/s, loss=0.5663,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 4 (Val MSE: 0.5335)

Época 5/1000 [Train] L1=9.22e-03: 100%| | 277/277 [00:03<00:00, 91.21it/s, loss=0.5384,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 5 (Val MSE: 0.5218)

Época 5/1000 | L1 coef: 9.22e-03

Train - Total: 0.5628 | MSE: 0.5279 | L0: 8.6

Val - Total: 0.5579 | MSE: 0.5218 | L0: 11.2

Best Val MSE: 0.5218 @ época 5 | Sin mejora: 0/20

Época 6/1000 [Train] L1=9.04e-03: 100%| | 277/277 [00:03<00:00, 83.87it/s, loss=0.5653,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 6 (Val MSE: 0.5150)

Época 7/1000 [Train] L1=8.86e-03: 100%| | 277/277 [00:03<00:00, 87.07it/s, loss=0.5524,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 7 (Val MSE: 0.5103)

Época 8/1000 [Train] L1=8.68e-03: 100%| | 277/277 [00:03<00:00, 91.83it/s, loss=0.5452,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 8 (Val MSE: 0.5070)

Época 9/1000 [Train] L1=8.51e-03: 100%| | 277/277 [00:02<00:00, 92.74it/s, loss=0.5404,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 9 (Val MSE: 0.5052)

Época 10/1000 [Train] L1=8.34e-03: 100%| | 277/277 [00:03<00:00, 92.29it/s, loss=0.5287,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 10 (Val MSE: 0.5019)

Época 10/1000 | L1 coef: 8.34e-03

Train - Total: 0.5346 | MSE: 0.5044 | L0: 14.4

Val - Total: 0.5328 | MSE: 0.5019 | L0: 16.6

Best Val MSE: 0.5019 @ época 10 | Sin mejora: 0/20

Época 11/1000 [Train] L1=8.17e-03: 100%| | 277/277 [00:02<00:00, 92.71it/s, loss=0.5206,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 11 (Val MSE: 0.5012)

Época 12/1000 [Train] L1=8.01e-03: 100%| | 277/277 [00:02<00:00, 92.55it/s, loss=0.5277,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 12 (Val MSE: 0.4986)

Época 13/1000 [Train] L1=7.85e-03: 100%| | 277/277 [00:02<00:00, 92.68it/s, loss=0.5252,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 13 (Val MSE: 0.4980)

Época 14/1000 [Train] L1=7.69e-03: 100%| | 277/277 [00:02<00:00, 93.57it/s, loss=0.5098,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 14 (Val MSE: 0.4960)

Época 15/1000 [Train] L1=7.54e-03: 100%| | 277/277 [00:02<00:00, 92.63it/s, loss=0.5222,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 15 (Val MSE: 0.4954)

Época 15/1000 | L1 coef: 7.54e-03

Train - Total: 0.5221 | MSE: 0.4966 | L0: 17.2

Val - Total: 0.5217 | MSE: 0.4954 | L0: 19.1

Best Val MSE: 0.4954 @ época 15 | Sin mejora: 0/20

Época 16/1000 [Train] L1=7.39e-03: 100%| | 277/277 [00:03<00:00, 92.17it/s, loss=0.5175,

Época 17/1000 [Train] L1=7.24e-03: 100%| | 277/277 [00:02<00:00, 92.89it/s, loss=0.5130,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 17 (Val MSE: 0.4942)

Época 18/1000 [Train] L1=7.09e-03: 100%| | 277/277 [00:02<00:00, 92.67it/s, loss=0.5171,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 18 (Val MSE: 0.4929)

Época 19/1000 [Train] L1=6.95e-03: 100% | 277/277 [00:03<00:00, 92.27it/s, loss=0.5203,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 19 (Val MSE: 0.4928)

Época 20/1000 [Train] L1=6.81e-03: 100% | 277/277 [00:02<00:00, 93.38it/s, loss=0.5166,

Época 20/1000 | L1 coef: 6.81e-03
Train - Total: 0.5149 | MSE: 0.4931 | L0: 18.8
Val - Total: 0.5163 | MSE: 0.4940 | L0: 20.3
Best Val MSE: 0.4928 @ época 19 | Sin mejora: 1/20

Época 21/1000 [Train] L1=6.68e-03: 100% | 277/277 [00:02<00:00, 93.27it/s, loss=0.5097,
Época 22/1000 [Train] L1=6.54e-03: 100% | 277/277 [00:02<00:00, 93.42it/s, loss=0.5202,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 22 (Val MSE: 0.4903)

Época 23/1000 [Train] L1=6.41e-03: 100% | 277/277 [00:02<00:00, 93.06it/s, loss=0.5357,
Época 24/1000 [Train] L1=6.28e-03: 100% | 277/277 [00:02<00:00, 92.92it/s, loss=0.5076,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 24 (Val MSE: 0.4889)

Época 25/1000 [Train] L1=6.16e-03: 100% | 277/277 [00:02<00:00, 93.19it/s, loss=0.5107,

Época 25/1000 | L1 coef: 6.16e-03
Train - Total: 0.5110 | MSE: 0.4921 | L0: 19.7
Val - Total: 0.5107 | MSE: 0.4915 | L0: 21.2
Best Val MSE: 0.4889 @ época 24 | Sin mejora: 1/20

Época 26/1000 [Train] L1=6.03e-03: 100%| | 277/277 [00:03<00:00, 91.31it/s, loss=0.5031,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 26 (Val MSE: 0.4887)

Época 27/1000 [Train] L1=5.91e-03: 100%| | 277/277 [00:03<00:00, 91.98it/s, loss=0.5070,

Época 28/1000 [Train] L1=5.80e-03: 100%| | 277/277 [00:02<00:00, 92.39it/s, loss=0.5174,

Época 29/1000 [Train] L1=5.68e-03: 100%| | 277/277 [00:02<00:00, 93.00it/s, loss=0.4977,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 29 (Val MSE: 0.4871)

Época 30/1000 [Train] L1=5.57e-03: 100%| | 277/277 [00:02<00:00, 92.44it/s, loss=0.4994,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 30 (Val MSE: 0.4869)

Época 30/1000 | L1 coef: 5.57e-03

Train - Total: 0.5048 | MSE: 0.4879 | L0: 20.8

Val - Total: 0.5043 | MSE: 0.4869 | L0: 22.6

Best Val MSE: 0.4869 @ época 30 | Sin mejora: 0/20

Época 31/1000 [Train] L1=5.45e-03: 100%| | 277/277 [00:03<00:00, 91.92it/s, loss=0.4818,

Época 32/1000 [Train] L1=5.35e-03: 100%| | 277/277 [00:02<00:00, 93.89it/s, loss=0.5039,

Época 33/1000 [Train] L1=5.24e-03: 100%| | 277/277 [00:02<00:00, 93.11it/s, loss=0.4973,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 33 (Val MSE: 0.4862)

Época 34/1000 [Train] L1=5.13e-03: 100%| | 277/277 [00:02<00:00, 93.12it/s, loss=0.5100,

Época 35/1000 [Train] L1=5.03e-03: 100%| | 277/277 [00:02<00:00, 93.20it/s, loss=0.4980,

Época 35/1000 | L1 coef: 5.03e-03

Train - Total: 0.5027 | MSE: 0.4878 | L0: 21.2

Val - Total: 0.5031 | MSE: 0.4879 | L0: 23.0

Best Val MSE: 0.4862 @ época 33 | Sin mejora: 2/20

Época 36/1000 [Train] L1=4.93e-03: 100%| | 277/277 [00:02<00:00, 92.87it/s, loss=0.5087,
Época 37/1000 [Train] L1=4.83e-03: 100%| | 277/277 [00:02<00:00, 93.24it/s, loss=0.5016,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 37 (Val MSE: 0.4861)

Época 38/1000 [Train] L1=4.74e-03: 100%| | 277/277 [00:02<00:00, 92.90it/s, loss=0.4783,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 38 (Val MSE: 0.4855)

Época 39/1000 [Train] L1=4.64e-03: 100%| | 277/277 [00:03<00:00, 91.66it/s, loss=0.4828,

Época 40/1000 [Train] L1=4.55e-03: 100%| | 277/277 [00:02<00:00, 92.77it/s, loss=0.5093,

Época 40/1000 | L1 coef: 4.55e-03

Train - Total: 0.4999 | MSE: 0.4865 | L0: 21.7

Val - Total: 0.5012 | MSE: 0.4876 | L0: 23.4

Best Val MSE: 0.4855 @ época 38 | Sin mejora: 2/20

Época 41/1000 [Train] L1=4.46e-03: 100%| | 277/277 [00:02<00:00, 93.37it/s, loss=0.5002,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 41 (Val MSE: 0.4852)

Época 42/1000 [Train] L1=4.37e-03: 100%| | 277/277 [00:02<00:00, 93.29it/s, loss=0.4924,

Época 43/1000 [Train] L1=4.28e-03: 100%| | 277/277 [00:02<00:00, 94.01it/s, loss=0.4983,

Época 44/1000 [Train] L1=4.19e-03: 100%| | 277/277 [00:02<00:00, 93.06it/s, loss=0.4929,

Época 45/1000 [Train] L1=4.11e-03: 100%| | 277/277 [00:02<00:00, 92.72it/s, loss=0.4887,

Época 45/1000 | L1 coef: 4.11e-03

Train - Total: 0.4982 | MSE: 0.4862 | L0: 22.1

Val - Total: 0.4986 | MSE: 0.4863 | L0: 23.8

Best Val MSE: 0.4852 @ época 41 | Sin mejora: 4/20

Época 46/1000 [Train] L1=4.03e-03: 100%| | 277/277 [00:03<00:00, 92.27it/s, loss=0.4923,
Época 47/1000 [Train] L1=3.95e-03: 100%| | 277/277 [00:02<00:00, 92.90it/s, loss=0.5013,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 47 (Val MSE: 0.4851)

Época 48/1000 [Train] L1=3.87e-03: 100%| | 277/277 [00:02<00:00, 92.50it/s, loss=0.4890,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 48 (Val MSE: 0.4847)

Época 49/1000 [Train] L1=3.79e-03: 100%| | 277/277 [00:02<00:00, 92.74it/s, loss=0.4982,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 49 (Val MSE: 0.4841)

Época 50/1000 [Train] L1=3.72e-03: 100%| | 277/277 [00:02<00:00, 93.53it/s, loss=0.4981,

Época 50/1000 | L1 coef: 3.72e-03

Train - Total: 0.4973 | MSE: 0.4864 | L0: 22.4

Val - Total: 0.4987 | MSE: 0.4875 | L0: 24.0

Best Val MSE: 0.4841 @ época 49 | Sin mejora: 1/20

Época 51/1000 [Train] L1=3.64e-03: 100%| | 277/277 [00:02<00:00, 92.76it/s, loss=0.5212,

Época 52/1000 [Train] L1=3.57e-03: 100%| | 277/277 [00:02<00:00, 93.39it/s, loss=0.5021,

Época 53/1000 [Train] L1=3.50e-03: 100%| | 277/277 [00:02<00:00, 92.50it/s, loss=0.4989,

Época 54/1000 [Train] L1=3.43e-03: 100%| | 277/277 [00:02<00:00, 93.53it/s, loss=0.4956,

Época 55/1000 [Train] L1=3.36e-03: 100%| | 277/277 [00:02<00:00, 93.66it/s, loss=0.5028,

Época 55/1000 | L1 coef: 3.36e-03

Train - Total: 0.4945 | MSE: 0.4846 | L0: 22.8

Val - Total: 0.4952 | MSE: 0.4849 | L0: 24.4

Best Val MSE: 0.4841 @ época 49 | Sin mejora: 6/20

Época 56/1000 [Train] L1=3.29e-03: 100%| | 277/277 [00:02<00:00, 92.62it/s, loss=0.5211,
Época 57/1000 [Train] L1=3.23e-03: 100%| | 277/277 [00:02<00:00, 93.42it/s, loss=0.4827,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 57 (Val MSE: 0.4834)

Época 58/1000 [Train] L1=3.16e-03: 100%| | 277/277 [00:02<00:00, 92.65it/s, loss=0.4811,
Época 59/1000 [Train] L1=3.10e-03: 100%| | 277/277 [00:02<00:00, 93.51it/s, loss=0.4812,
Época 60/1000 [Train] L1=3.04e-03: 100%| | 277/277 [00:02<00:00, 92.73it/s, loss=0.4854,

Época 60/1000 | L1 coef: 3.04e-03

Train - Total: 0.4931 | MSE: 0.4840 | L0: 23.1

Val - Total: 0.4946 | MSE: 0.4851 | L0: 24.7

Best Val MSE: 0.4834 @ época 57 | Sin mejora: 3/20

Época 61/1000 [Train] L1=2.98e-03: 100%| | 277/277 [00:02<00:00, 94.17it/s, loss=0.4925,
Época 62/1000 [Train] L1=2.92e-03: 100%| | 277/277 [00:02<00:00, 93.15it/s, loss=0.4996,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 62 (Val MSE: 0.4830)

Época 63/1000 [Train] L1=2.86e-03: 100%| | 277/277 [00:03<00:00, 92.27it/s, loss=0.4827,
Época 64/1000 [Train] L1=2.80e-03: 100%| | 277/277 [00:03<00:00, 92.10it/s, loss=0.4998,
Época 65/1000 [Train] L1=2.74e-03: 100%| | 277/277 [00:02<00:00, 93.97it/s, loss=0.4866,

Época 65/1000 | L1 coef: 2.74e-03

Train - Total: 0.4924 | MSE: 0.4841 | L0: 23.2

Val - Total: 0.4923 | MSE: 0.4837 | L0: 25.0

Best Val MSE: 0.4830 @ época 62 | Sin mejora: 3/20

Época 66/1000 [Train] L1=2.69e-03: 100%| | 277/277 [00:02<00:00, 92.63it/s, loss=0.4887,
Época 67/1000 [Train] L1=2.64e-03: 100%| | 277/277 [00:02<00:00, 93.80it/s, loss=0.4807,
Época 68/1000 [Train] L1=2.58e-03: 100%| | 277/277 [00:03<00:00, 92.14it/s, loss=0.4804,
Época 69/1000 [Train] L1=2.53e-03: 100%| | 277/277 [00:02<00:00, 93.13it/s, loss=0.4839,
Época 70/1000 [Train] L1=2.48e-03: 100%| | 277/277 [00:02<00:00, 93.32it/s, loss=0.4784,

Época 70/1000 | L1 coef: 2.48e-03
Train - Total: 0.4913 | MSE: 0.4836 | L0: 23.5
Val - Total: 0.4913 | MSE: 0.4834 | L0: 25.1
Best Val MSE: 0.4830 @ época 62 | Sin mejora: 8/20

Época 71/1000 [Train] L1=2.43e-03: 100%| | 277/277 [00:03<00:00, 91.56it/s, loss=0.4855,
Época 72/1000 [Train] L1=2.38e-03: 100%| | 277/277 [00:02<00:00, 93.85it/s, loss=0.4993,
Época 73/1000 [Train] L1=2.33e-03: 100%| | 277/277 [00:02<00:00, 93.26it/s, loss=0.4984,
Época 74/1000 [Train] L1=2.29e-03: 100%| | 277/277 [00:02<00:00, 92.52it/s, loss=0.5051,
Época 75/1000 [Train] L1=2.24e-03: 100%| | 277/277 [00:02<00:00, 93.13it/s, loss=0.4834,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 75 (Val MSE: 0.4828)

Época 75/1000 | L1 coef: 2.24e-03
Train - Total: 0.4911 | MSE: 0.4841 | L0: 23.6
Val - Total: 0.4901 | MSE: 0.4828 | L0: 25.4
Best Val MSE: 0.4828 @ época 75 | Sin mejora: 0/20

Época 76/1000 [Train] L1=2.20e-03: 100%| | 277/277 [00:02<00:00, 92.36it/s, loss=0.4903,
Época 77/1000 [Train] L1=2.15e-03: 100%| | 277/277 [00:02<00:00, 92.62it/s, loss=0.4988,
Época 78/1000 [Train] L1=2.11e-03: 100%| | 277/277 [00:02<00:00, 93.49it/s, loss=0.4892,
Época 79/1000 [Train] L1=2.07e-03: 100%| | 277/277 [00:02<00:00, 93.45it/s, loss=0.4723,
Época 80/1000 [Train] L1=2.03e-03: 100%| | 277/277 [00:03<00:00, 92.16it/s, loss=0.4843,

Época 80/1000 | L1 coef: 2.03e-03
Train - Total: 0.4910 | MSE: 0.4845 | L0: 23.7
Val - Total: 0.4919 | MSE: 0.4853 | L0: 25.4
Best Val MSE: 0.4828 @ época 75 | Sin mejora: 5/20

Época 81/1000 [Train] L1=1.99e-03: 100%| | 277/277 [00:02<00:00, 93.23it/s, loss=0.4815,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 81 (Val MSE: 0.4821)

Época 82/1000 [Train] L1=1.95e-03: 100%| | 277/277 [00:03<00:00, 92.33it/s, loss=0.4850,
Época 83/1000 [Train] L1=1.91e-03: 100%| | 277/277 [00:02<00:00, 93.58it/s, loss=0.4891,
Época 84/1000 [Train] L1=1.87e-03: 100%| | 277/277 [00:02<00:00, 93.42it/s, loss=0.4886,
Época 85/1000 [Train] L1=1.83e-03: 100%| | 277/277 [00:02<00:00, 93.40it/s, loss=0.4862,

Época 85/1000 | L1 coef: 1.83e-03

Train - Total: 0.4891 | MSE: 0.4831 | L0: 23.9

Val - Total: 0.4884 | MSE: 0.4822 | L0: 25.7

Best Val MSE: 0.4821 @ época 81 | Sin mejora: 4/20

Época 86/1000 [Train] L1=1.80e-03: 100%| | 277/277 [00:02<00:00, 92.72it/s, loss=0.4824,
Época 87/1000 [Train] L1=1.76e-03: 100%| | 277/277 [00:02<00:00, 93.97it/s, loss=0.4990,
Época 88/1000 [Train] L1=1.72e-03: 100%| | 277/277 [00:02<00:00, 93.47it/s, loss=0.4827,
Época 89/1000 [Train] L1=1.69e-03: 100%| | 277/277 [00:03<00:00, 91.37it/s, loss=0.4969,
Época 90/1000 [Train] L1=1.66e-03: 100%| | 277/277 [00:02<00:00, 93.43it/s, loss=0.4987,

Época 90/1000 | L1 coef: 1.66e-03

Train - Total: 0.4907 | MSE: 0.4852 | L0: 23.8

Val - Total: 0.4892 | MSE: 0.4835 | L0: 25.6

Best Val MSE: 0.4821 @ época 81 | Sin mejora: 9/20

Época 91/1000 [Train] L1=1.62e-03: 100%| | 277/277 [00:02<00:00, 93.49it/s, loss=0.4909,
Época 92/1000 [Train] L1=1.59e-03: 100%| | 277/277 [00:02<00:00, 92.94it/s, loss=0.4905,
Época 93/1000 [Train] L1=1.56e-03: 100%| | 277/277 [00:02<00:00, 94.01it/s, loss=0.4987,
Época 94/1000 [Train] L1=1.53e-03: 100%| | 277/277 [00:02<00:00, 92.76it/s, loss=0.4744,
Época 95/1000 [Train] L1=1.50e-03: 100%| | 277/277 [00:02<00:00, 92.84it/s, loss=0.4740,

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt

Mejor modelo guardado en época 95 (Val MSE: 0.4817)

Época 95/1000 | L1 coef: 1.50e-03

Train - Total: 0.4892 | MSE: 0.4841 | L0: 24.0

Val - Total: 0.4869 | MSE: 0.4817 | L0: 25.9

Best Val MSE: 0.4817 @ época 95 | Sin mejora: 0/20

Época 96/1000 [Train] L1=1.47e-03: 100%| | 277/277 [00:03<00:00, 91.75it/s, loss=0.4898,
Época 97/1000 [Train] L1=1.44e-03: 100%| | 277/277 [00:02<00:00, 93.37it/s, loss=0.4775,
Época 98/1000 [Train] L1=1.41e-03: 100%| | 277/277 [00:02<00:00, 92.86it/s, loss=0.4953,

Época 99/1000 [Train] L1=1.38e-03: 100%| | 277/277 [00:02<00:00, 93.33it/s, loss=0.4977,
Época 100/1000 [Train] L1=1.35e-03: 100%| | 277/277 [00:02<00:00, 93.23it/s, loss=0.4858

Época 100/1000 | L1 coef: 1.35e-03
Train - Total: 0.4891 | MSE: 0.4844 | L0: 24.2
Val - Total: 0.4896 | MSE: 0.4848 | L0: 25.8
Best Val MSE: 0.4817 @ época 95 | Sin mejora: 5/20

Época 101/1000 [Train] L1=1.33e-03: 100%| | 277/277 [00:02<00:00, 92.78it/s, loss=0.4972
Época 102/1000 [Train] L1=1.30e-03: 100%| | 277/277 [00:03<00:00, 90.87it/s, loss=0.4867
Época 103/1000 [Train] L1=1.27e-03: 100%| | 277/277 [00:03<00:00, 91.52it/s, loss=0.5145
Época 104/1000 [Train] L1=1.25e-03: 100%| | 277/277 [00:03<00:00, 92.12it/s, loss=0.4840

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 104 (Val MSE: 0.4809)

Época 105/1000 [Train] L1=1.22e-03: 100%| | 277/277 [00:03<00:00, 91.99it/s, loss=0.4953

Época 105/1000 | L1 coef: 1.22e-03
Train - Total: 0.4869 | MSE: 0.4826 | L0: 24.4
Val - Total: 0.4877 | MSE: 0.4832 | L0: 26.0
Best Val MSE: 0.4809 @ época 104 | Sin mejora: 1/20

Época 106/1000 [Train] L1=1.20e-03: 100%| | 277/277 [00:02<00:00, 93.36it/s, loss=0.4785
Época 107/1000 [Train] L1=1.17e-03: 100%| | 277/277 [00:02<00:00, 93.05it/s, loss=0.4727
Época 108/1000 [Train] L1=1.15e-03: 100%| | 277/277 [00:03<00:00, 91.51it/s, loss=0.4912
Época 109/1000 [Train] L1=1.13e-03: 100%| | 277/277 [00:03<00:00, 91.81it/s, loss=0.5012
Época 110/1000 [Train] L1=1.11e-03: 100%| | 277/277 [00:02<00:00, 92.42it/s, loss=0.4644

Época 110/1000 | L1 coef: 1.11e-03
Train - Total: 0.4865 | MSE: 0.4825 | L0: 24.5
Val - Total: 0.4868 | MSE: 0.4827 | L0: 26.1
Best Val MSE: 0.4809 @ época 104 | Sin mejora: 6/20

Época 111/1000 [Train] L1=1.08e-03: 100%| | 277/277 [00:03<00:00, 91.91it/s, loss=0.4872
Época 112/1000 [Train] L1=1.06e-03: 100%| | 277/277 [00:02<00:00, 92.41it/s, loss=0.4824
Época 113/1000 [Train] L1=1.04e-03: 100%| | 277/277 [00:03<00:00, 91.73it/s, loss=0.5049
Época 114/1000 [Train] L1=1.02e-03: 100%| | 277/277 [00:03<00:00, 91.98it/s, loss=0.4800
Época 115/1000 [Train] L1=9.99e-04: 100%| | 277/277 [00:03<00:00, 92.33it/s, loss=0.4888

Época 115/1000 | L1 coef: 9.99e-04
Train - Total: 0.4850 | MSE: 0.4812 | L0: 24.6
Val - Total: 0.4854 | MSE: 0.4815 | L0: 26.2
Best Val MSE: 0.4809 @ época 104 | Sin mejora: 11/20

Época 116/1000 [Train] L1=9.79e-04: 100%| | 277/277 [00:03<00:00, 91.97it/s, loss=0.4865

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
Mejor modelo guardado en época 116 (Val MSE: 0.4807)

Época 117/1000 [Train] L1=9.60e-04: 100%| | 277/277 [00:02<00:00, 92.73it/s, loss=0.4837
Época 118/1000 [Train] L1=9.41e-04: 100%| | 277/277 [00:02<00:00, 93.57it/s, loss=0.4733
Época 119/1000 [Train] L1=9.22e-04: 100%| | 277/277 [00:02<00:00, 93.44it/s, loss=0.4854
Época 120/1000 [Train] L1=9.03e-04: 100%| | 277/277 [00:02<00:00, 92.99it/s, loss=0.4933

Época 120/1000 | L1 coef: 9.03e-04
Train - Total: 0.4860 | MSE: 0.4825 | L0: 24.6
Val - Total: 0.4848 | MSE: 0.4812 | L0: 26.3
Best Val MSE: 0.4807 @ época 116 | Sin mejora: 4/20

Época 121/1000 [Train] L1=8.85e-04: 100%| | 277/277 [00:02<00:00, 92.81it/s, loss=0.4792
Época 122/1000 [Train] L1=8.68e-04: 100%| | 277/277 [00:03<00:00, 92.01it/s, loss=0.4875
Época 123/1000 [Train] L1=8.50e-04: 100%| | 277/277 [00:03<00:00, 92.32it/s, loss=0.4968
Época 124/1000 [Train] L1=8.33e-04: 100%| | 277/277 [00:02<00:00, 92.48it/s, loss=0.4867
Época 125/1000 [Train] L1=8.17e-04: 100%| | 277/277 [00:02<00:00, 92.62it/s, loss=0.4789

Época 125/1000 | L1 coef: 8.17e-04
Train - Total: 0.4870 | MSE: 0.4838 | L0: 24.6
Val - Total: 0.4855 | MSE: 0.4823 | L0: 26.3
Best Val MSE: 0.4807 @ época 116 | Sin mejora: 9/20

Época 126/1000 [Train] L1=8.00e-04: 100%| | 277/277 [00:02<00:00, 93.00it/s, loss=0.4943
 Época 127/1000 [Train] L1=7.84e-04: 100%| | 277/277 [00:02<00:00, 93.54it/s, loss=0.4744
 Época 128/1000 [Train] L1=7.69e-04: 100%| | 277/277 [00:02<00:00, 92.36it/s, loss=0.4980
 Época 129/1000 [Train] L1=7.53e-04: 100%| | 277/277 [00:02<00:00, 92.62it/s, loss=0.5032
 Época 130/1000 [Train] L1=7.38e-04: 100%| | 277/277 [00:02<00:00, 92.97it/s, loss=0.4756

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_best.pt
 Mejor modelo guardado en época 130 (Val MSE: 0.4804)

Época 130/1000 | L1 coef: 7.38e-04
 Train - Total: 0.4859 | MSE: 0.4829 | L0: 24.6
 Val - Total: 0.4835 | MSE: 0.4804 | L0: 26.4
 Best Val MSE: 0.4804 @ época 130 | Sin mejora: 0/20

Época 131/1000 [Train] L1=7.23e-04: 100%| | 277/277 [00:03<00:00, 91.65it/s, loss=0.4819
 Época 132/1000 [Train] L1=7.09e-04: 100%| | 277/277 [00:02<00:00, 92.42it/s, loss=0.4947
 Época 133/1000 [Train] L1=6.95e-04: 100%| | 277/277 [00:02<00:00, 92.48it/s, loss=0.4909
 Época 134/1000 [Train] L1=6.81e-04: 100%| | 277/277 [00:03<00:00, 92.09it/s, loss=0.4704
 Época 135/1000 [Train] L1=6.67e-04: 100%| | 277/277 [00:02<00:00, 92.77it/s, loss=0.4928

Época 135/1000 | L1 coef: 6.67e-04
 Train - Total: 0.4860 | MSE: 0.4832 | L0: 24.7
 Val - Total: 0.4838 | MSE: 0.4811 | L0: 26.5
 Best Val MSE: 0.4804 @ época 130 | Sin mejora: 5/20

Época 136/1000 [Train] L1=6.54e-04: 100%| | 277/277 [00:02<00:00, 93.18it/s, loss=0.4866
 Época 137/1000 [Train] L1=6.41e-04: 100%| | 277/277 [00:02<00:00, 92.65it/s, loss=0.4925
 Época 138/1000 [Train] L1=6.28e-04: 100%| | 277/277 [00:03<00:00, 90.72it/s, loss=0.4773
 Época 139/1000 [Train] L1=6.15e-04: 100%| | 277/277 [00:03<00:00, 92.07it/s, loss=0.4808
 Época 140/1000 [Train] L1=6.03e-04: 100%| | 277/277 [00:03<00:00, 92.20it/s, loss=0.4948

Época 140/1000 | L1 coef: 6.03e-04
 Train - Total: 0.4864 | MSE: 0.4839 | L0: 24.8
 Val - Total: 0.4851 | MSE: 0.4825 | L0: 26.5
 Best Val MSE: 0.4804 @ época 130 | Sin mejora: 10/20

```

Época 141/1000 [Train] L1=5.91e-04: 100%|      | 277/277 [00:03<00:00, 92.30it/s, loss=0.4968
Época 142/1000 [Train] L1=5.79e-04: 100%|      | 277/277 [00:03<00:00, 90.98it/s, loss=0.4836
Época 143/1000 [Train] L1=5.68e-04: 100%|      | 277/277 [00:03<00:00, 92.08it/s, loss=0.4853
Época 144/1000 [Train] L1=5.56e-04: 100%|      | 277/277 [00:03<00:00, 92.21it/s, loss=0.5091
Época 145/1000 [Train] L1=5.45e-04: 100%|      | 277/277 [00:03<00:00, 91.77it/s, loss=0.4974

```

```

Época 145/1000 | L1 coef: 5.45e-04
Train - Total: 0.4844 | MSE: 0.4821 | L0: 24.9
Val   - Total: 0.4831 | MSE: 0.4807 | L0: 26.7
Best Val MSE: 0.4804 @ época 130 | Sin mejora: 15/20

```

```

Época 146/1000 [Train] L1=5.34e-04: 100%|      | 277/277 [00:03<00:00, 92.21it/s, loss=0.4675
Época 147/1000 [Train] L1=5.24e-04: 100%|      | 277/277 [00:03<00:00, 92.20it/s, loss=0.4828
Época 148/1000 [Train] L1=5.13e-04: 100%|      | 277/277 [00:02<00:00, 92.64it/s, loss=0.5010
Época 149/1000 [Train] L1=5.03e-04: 100%|      | 277/277 [00:03<00:00, 91.89it/s, loss=0.4937
Época 150/1000 [Train] L1=4.93e-04: 100%|      | 277/277 [00:03<00:00, 92.19it/s, loss=0.4773

```

```

Early stopping activado en época 150
Mejor modelo: época 130 con Val MSE = 0.4804

```

```

=====
RESUMEN DEL ENTRENAMIENTO
=====
Mejor modelo en época 130 con Val MSE = 0.4804
Early stopping tras 150 épocas

Métricas Finales (Última Época):
      Métrica  Train Validation
Total Loss 0.4848      0.4828
MSE Loss 0.4827      0.4806
L0 Sparsity 24.9       26.6
Fraction Active 0.15%    0.16%
=====

```

8. Guardar Modelo

```

# Guardar checkpoint final (local + organizado)
checkpoint = {
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'config': CONFIG,
    'history': history
}

# Guardado local
final_model_name = get_model_name(
    variante=NAMING_META['variante'],
    tecnica=NAMING_META['tecnica'],
    capa=NAMING_META['capa'],
    juegos=NAMING_META['juegos'],
    estado='final'
)
checkpoint_path = CONFIG['save_dir'] / final_model_name
torch.save(checkpoint, checkpoint_path)

# Guardado organizado
save_checkpoint(
    model,
    base_path=NAMING_META['base_path'],
    variante=NAMING_META['variante'],
    tecnica=NAMING_META['tecnica'],
    capa=NAMING_META['capa'],
    juegos=NAMING_META['juegos'],
    estado='final',
    extras=NAMING_META['extras'],
    config=CONFIG,
    history=history,
    optimizer=optimizer
)

# Resumen de archivos guardados
saved_files = pd.DataFrame({
    'Tipo': ['Local', 'Organizado'],
    'Ubicación': [
        str(checkpoint_path),
        f"layer_{NAMING_META['capa']:02d}/models/..."
    ]
})

```



```

print('\n' + '='*70)
print(' CHECKPOINTS GUARDADOS')
print('='*70)
print(saved_files.to_string(index=False))
print('='*70)

```

Checkpoint guardado: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\1500games\sae_standard_l1-decay_l6_1500g_ex1_final.pt

```

=====
CHECKPOINTS GUARDADOS
=====

```

Tipo	Ubicación
Local	saved_model\sae_standard_l1-decay_l6_1500g_final.pt
Organizado	layer_06/models/...

```

=====

```

9. Métricas básicas del modelo

```

# Guardar metricas de entrenamiento
metrics_dir = get_metrics_dir(
    NAMING_META["experiment_base_path"],
    NAMING_META["variante"],
    NAMING_META["tecnica"],
    NAMING_META["capa"],
    NAMING_META["juegos"],
    extras_id=extras_id
)

training_metrics = {
    # Metricas de calidad
    "val_mse_final": history["val_mse_loss"][-1],
    "val_mse_best": best_val_mse,
    "val_l0_final": history["val_l0_sparsity"][-1],
    "val_fraction_active_final": history["val_fraction_active"][-1],
    # Contexto del entrenamiento
    "best_epoch": best_epoch,
    "total_epochs": len(history["train_total_loss"]),
    "train_mse_final": history["train_mse_loss"][-1],

```

```

    "train_l0_final": history["train_l0_sparsity"][-1],
    # Hiperparametros
    "hidden_dim": CONFIG["hidden_dim"],
    "expansion_factor": CONFIG["hidden_dim"] / CONFIG["input_dim"],
    "sparsity_coef": CONFIG["sparsity_coef"],
    "sparsity_decay": CONFIG["sparsity_decay"],
    "learning_rate": CONFIG["learning_rate"],
    "num_games": NAMING_META["juegos"],
}

metrics_path = metrics_dir / "training_metrics.json"
with open(metrics_path, "w") as f:
    _json.dump(training_metrics, f, indent=2)

# Resumen
metrics_df = pd.DataFrame({
    "Metrica": ["Val MSE (final)", "Val MSE (mejor)", "Val L0 (final)", "Fraccion Activa"],
    "Valor": [
        f"{training_metrics['val_mse_final']:.6f}",
        f"{training_metrics['val_mse_best']:.6f}",
        f"{training_metrics['val_l0_final']:.1f}",
        f"{training_metrics['val_fraction_active_final']:.2%}"
    ]
})

print("=" * 50)
print(" METRICAS DE ENTRENAMIENTO GUARDADAS")
print("=" * 50)
print(metrics_df.to_string(index=False))
print(f"Archivo: {metrics_path}")
print("=" * 50)

```

```

=====
METRICAS DE ENTRENAMIENTO GUARDADAS
=====

```

Metrica	Valor
Val MSE (final)	0.480590
Val MSE (mejor)	0.480440
Val L0 (final)	26.6
Fraccion Activa	0.16%

```

Archivo: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\
decay_standard\1500games\ex1\metrics\training_metrics.json
=====

```

10. Visualización de Resultados

```
# Gráficas de pérdidas (Train vs Validation) + L1 Decay
fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# Total loss
axes[0, 0].plot(history['train_total_loss'], label='Train', alpha=0.8)
axes[0, 0].plot(history['val_total_loss'], label='Validation', alpha=0.8)
axes[0, 0].set_title('Total Loss')
axes[0, 0].set_xlabel('Época')
axes[0, 0].set_ylabel('Loss')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)

# MSE loss
axes[0, 1].plot(history['train_mse_loss'], label='Train', color='orange', alpha=0.8)
axes[0, 1].plot(history['val_mse_loss'], label='Validation', color='red', alpha=0.8)
if CONFIG['early_stopping'] and best_epoch > 0:
    axes[0, 1].axvline(best_epoch - 1, color='green', linestyle='--', alpha=0.7, label=f'Best Epoch')
axes[0, 1].set_title('MSE Loss (Reconstrucción)')
axes[0, 1].set_xlabel('Época')
axes[0, 1].set_ylabel('MSE')
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)

# L0 sparsity
axes[0, 2].plot(history['train_l0_sparsity'], label='Train', color='green', alpha=0.8)
axes[0, 2].plot(history['val_l0_sparsity'], label='Validation', color='darkgreen', alpha=0.8)
axes[0, 2].set_title('L0 Sparsity (Features Activas)')
axes[0, 2].set_xlabel('Época')
axes[0, 2].set_ylabel('Número de features')
axes[0, 2].legend()
axes[0, 2].grid(True, alpha=0.3)

# Fraction active
axes[1, 0].plot(history['train_fraction_active'], label='Train', color='purple', alpha=0.8)
axes[1, 0].plot(history['val_fraction_active'], label='Validation', color='magenta', alpha=0.8)
axes[1, 0].set_title('Fracción de Features Activas')
axes[1, 0].set_xlabel('Época')
axes[1, 0].set_ylabel('Fracción')
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)
```

```

# Sparsity Coefficient (L1 Decay)
axes[1, 1].plot(history['sparsity_coef_history'], color='brown', linewidth=2)
axes[1, 1].set_title('L1 Coefficient (Decay)')
axes[1, 1].set_xlabel('Época')
axes[1, 1].set_ylabel('Sparsity Coef')
axes[1, 1].set_yscale('log')
axes[1, 1].grid(True, alpha=0.3)


# Train vs Val MSE Gap
mse_gap = [v - t for v, t in zip(history['val_mse_loss'], history['train_mse_loss'])]
axes[1, 2].plot(mse_gap, color='crimson', linewidth=2)
axes[1, 2].axhline(0, color='black', linestyle='--', alpha=0.5)
if CONFIG['early_stopping'] and best_epoch > 0:
    axes[1, 2].axvline(best_epoch - 1, color='green', linestyle='--', alpha=0.7, label=f'Best')
axes[1, 2].set_title('Overfitting Gap (Val MSE - Train MSE)')
axes[1, 2].set_xlabel('Época')
axes[1, 2].set_ylabel('Gap MSE')
axes[1, 2].legend()
axes[1, 2].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(CONFIG['plots_dir'] / 'training_curves.png', dpi=300, bbox_inches='tight')
plt.show()

print('\n' + '='*50)
print('  Gráficas guardadas exitosamente')

print('='*50)

```



=====
 Gráficas guardadas exitosamente
 =====

11. Resumen Final

```
# Resumen final completo
print('\n' + '='*70)
print(' RESUMEN FINAL DEL ENTRENAMIENTO')
print('='*70)

# Arquitectura
arch_summary = pd.DataFrame({
    'Componente': ['Input', 'Hidden', 'Parámetros'],
    'Valor': [
        f"{CONFIG['input_dim']}",
        f"{CONFIG['hidden_dim']} ({CONFIG['hidden_dim']/CONFIG['input_dim']:.1f}x)",
        f"{num_params:}"
    ]
})
```

```

print('\nArquitectura:')
print(arch_summary.to_string(index=False))

# Datos
data_summary = pd.DataFrame({
    'Conjunto': ['Train', 'Validation', 'Total'],
    'Muestras': [
        f"{len(train_dataset):,}",
        f"{len(val_dataset):,}",
        f"{len(full_dataset):,}"
    ]
})
print('\nDatos:')
print(data_summary.to_string(index=False))

# Métricas finales
final_summary = pd.DataFrame({
    'Métrica': ['Train Loss', 'Val Loss', 'Train MSE', 'Val MSE', 'Train L0', 'Val L0'],
    'Valor': [
        f"{history['train_total_loss'][-1]:.4f}",
        f"{history['val_total_loss'][-1]:.4f}",
        f"{history['train_mse_loss'][-1]:.4f}",
        f"{history['val_mse_loss'][-1]:.4f}",
        f"{history['train_l0_sparsity'][-1]:.1f}",
        f"{history['val_l0_sparsity'][-1]:.1f}"
    ]
})
print('\nMétricas Finales:')
print(final_summary.to_string(index=False))

# Archivos
files_summary = pd.DataFrame({
    'Tipo': ['Modelo', 'Gráficas'],
    'Ubicación': [
        str(checkpoint_path),
        str(CONFIG['plots_dir'] / 'training_curves.png')
    ]
})
print('\nArchivos Guardados:')
print(files_summary.to_string(index=False))

print('\n' + '='*70)

```

```
print('  ENTRENAMIENTO COMPLETADO EXITOSAMENTE')
print('='*70)
```

```
=====
RESUMEN FINAL DEL ENTRENAMIENTO
=====
```

Arquitectura:

Componente	Valor
Input	512
Hidden	16384 (32.0x)
Parámetros	16,794,112

Datos:

Conjunto	Muestras
Train	70,800
Validation	17,700
Total	88,500

Métricas Finales:

Métrica	Valor
Train Loss	0.4848
Val Loss	0.4828
Train MSE	0.4827
Val MSE	0.4806
Train L0	24.9
Val L0	26.6

Archivos Guardados:

Tipo	Modelo
decay_l6_1500g_final.pt	
Gráficas	c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\1500games\ex1\plots\training_curves.png

```
=====
ENTRENAMIENTO COMPLETADO EXITOSAMENTE
=====
```

12. Generar PDF con Quarto

Una vez ejecutado todo el notebook, ejecutar el siguiente comando en la terminal para generar el PDF con Quarto:

```
import subprocess

# Obtener directorio del experimento y nombre del PDF
exp_dir = get_experiment_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

pdf_name = get_report_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    'training',
    extras_id=extras_id
)

notebook_name = 'train_sae_standard_l1-decay.ipynb'
output_path = exp_dir / pdf_name

print(f' Generando PDF con Quarto...')
print(f' Archivo de salida: {output_path}')

# Quarto no acepta rutas en --output, solo nombre de archivo
# Se usa --output-dir para el directorio y --output solo para el nombre
result = subprocess.run(
    f'quarto render {notebook_name} --to pdf --output-dir "{exp_dir}" --output {pdf_name}',
    shell=True,
    capture_output=True,
    text=True
)

if result.returncode == 0:
    print(f' PDF generado exitosamente: {output_path}')
```



```
else:
    print(f' Error al generar PDF:')
    print(result.stderr)
```

```
Generando PDF con Quarto...
Archivo de salida: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\1500games\ex1\sae_standard_l1-decay_l6_1500g_ex1_training.pdf
PDF generado exitosamente: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\1500games\ex1\sae_standard_l1-decay_l6_1500g_ex1_training.pdf
```