

Reconstruction Metric - GPU Optimized

Karina

2026-05-20

Table of contents

0.1	1. Setup e Imports	1
0.2	2. Cargar Datos	2
0.3	3. Cargar SAE y Extraer Features	3
0.4	4. Split Train/Test y Filtrar BSPs de Piezas	4
0.5	5. Funciones GPU-Optimizadas	5
0.6	6. Fase 1: Identificar Features de Alta Precisión (GPU)	6
0.7	7. Fase 2: Reconstruir Tableros en Test Set (GPU)	9
0.8	8. Análisis de Resultados	12
0.9	9. Guardar Resultados	14

0.1 1. Setup e Imports

```
import numpy as np
import pandas as pd
import sys
import os
import subprocess
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt

from pathlib import Path
from tqdm import tqdm
from sae.models.sae import SparseAutoencoder
from sae.tools.experiment_utils import get_experiment_dir, get_metrics_dir, get_report_name
from sae.tools.naming_utils import get_checkpoint_dir, get_model_name, get_extras_id

project_root = Path('...').resolve()
sys.path.insert(0, str(project_root))

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Device: {device}")
if torch.cuda.is_available():
    print(f"GPU: {torch.cuda.get_device_name(0)}")
    print(f"Memoria: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
```

```
Device: cuda
GPU: NVIDIA GeForce RTX 5060
Memoria: 8.55 GB
```

```
# Configuración para visualización en PDF
pd.set_option('display.max_colwidth', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 100)
```

```

np.set_printoptions(linewidth=100, edgeitems=3)

os.environ['COLUMNS'] = '100'

print(" Configuración para PDF lista")

```

Configuración para PDF lista

```

NAMING_META= {
    'variante': 'standard', #ArquitecturadelSAE
    'tecnica': 'p-annealing', #P-annealing(Lp~pconpdecreciente)
    'capa': 6, #LayerdelmodeloOthelloGPT
    'juegos': 12500, #Númerodepartidas
    'base_path': 'C:\\Users\\Esposa\\Documents\\Repos\\sae-othello-gpt', #Pathbaseparacheckpoints
    'experiment_base_path': os.path.abspath('../../experiments'), #sae/experiments/
    'extras': {
        'expansion': 16,
        'sparsity_coef': 5e-1,
        'epochs': 200,
        'early_stopping': False,
        'learning_rate': 1e-3,
        'p_start': 1.0,
        'p_end': 0.2,
        'n_sparsity_updates': 10,
        'anneal_start': 5000,
        'batch_size': 2048,
        'warmup_steps': 1000,
        'sparsity_warmup_steps': 2000,
        'estimated_epochs': 200
    }
}

print('\n Metadata de naming:')
for key, value in NAMING_META.items():
    print(f' {key}: {value}')

```

Metadata de naming:

```

variante: standard
tecnica: p-annealing
capa: 6
juegos: 12500
base_path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt
experiment_base_path: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments
extras: {'expansion': 16, 'sparsity_coef': 0.5, 'epochs': 200, 'early_stopping': False, 'learning_rate': 0.001, 'p_st

```

0.2 2. Cargar Datos

```

# Cargar activaciones
activations_path = project_root / "activations" / "data" / "layer6_2000games.npy"
activations = np.load(activations_path)

print(f"Activaciones del modelo:")
print(f" Shape: {activations.shape}")
print(f" Memoria: {activations.nbytes / (1024**2):.2f} MB")

```

Activaciones del modelo:

```

Shape: (118000, 512)
Memoria: 230.47 MB

```

```
# Cargar ground truth de BSPs
bsp_gt_path = project_root / "metrics" / "02_data" / "bsp_ground_truth_2000games.npy"
bsp_names_path = project_root / "metrics" / "02_data" / "bsp_ground_truth_2000games.names.npy"

bsp_ground_truth = np.load(bsp_gt_path)
bsp_names = np.load(bsp_names_path, allow_pickle=True)

print(f"\nGround truth BSPs:")
print(f"  Shape: {bsp_ground_truth.shape}")
print(f"  Total BSPs: {len(bsp_names)}")
```

Ground truth BSPs:
 Shape: (118000, 326)
 Total BSPs: 326

0.3 3. Cargar SAE y Extraer Features

```
# Configuración del SAE
input_dim = 512
hidden_dim = 8192

# Construir ruta del modelo usando naming_utils
checkpoint_dir = get_checkpoint_dir(
    NAMING_META['base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos']
)
extras_id = get_extras_id(checkpoint_dir, NAMING_META['extras']) if NAMING_META.get('extras') else None

model_filename = get_model_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    'final',
    extras_id=extras_id
)
model_path = checkpoint_dir / model_filename

# Cargar modelo SAE estándar
sae = SparseAutoencoder(input_dim, hidden_dim).to(device)

checkpoint = torch.load(model_path, map_location=device, weights_only=False)
sae.load_state_dict(checkpoint['model_state_dict'])
sae.eval()

print(f"SAE cargado:")
print(f"  Path: {model_path}")
print(f"  Input: {input_dim}, Hidden: {hidden_dim}")
print(f"  Expansion: {hidden_dim/input_dim}x")
```

SAE cargado:
 Path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_standard\sae_p-annealing_standard\12500game
 annealing_16_12500g_ex5_final.pt
 Input: 512, Hidden: 8192
 Expansion: 16.0x

```
# Extraer features del SAE en GPU
print("Extrayendo features del SAE...")
activations_tensor = torch.from_numpy(activations).float().to(device)

with torch.no_grad():
    sae_features = sae.encode(activations_tensor)

print(f"\nFeatures SAE:")
print(f"  Shape: {sae_features.shape}")
print(f"  Device: {sae_features.device}")
sparsity = (sae_features == 0).float().mean().item()
print(f"  Sparsity: {sparsity:.2%}")
print(f"  Activaciones promedio: {sae_features.shape[1] * (1 - sparsity):.1f}")
```

Extrayendo features del SAE...

```
Features SAE:
  Shape: torch.Size([118000, 8192])
  Device: cuda:0
  Sparsity: 97.90%
  Activaciones promedio: 172.0
```

```
del activations_tensor
torch.cuda.empty_cache()
```

0.4 4. Split Train/Test y Filtrar BSPs de Piezas

```
# Split 50/50:
n_moves = 59
n_games = 2000 # número de partidas para prueba de métricas
train_games = n_games // 2
split_idx = train_games * n_moves

sae_features_train = sae_features[:split_idx]
sae_features_test = sae_features[split_idx:]

print(f"Total partidas: {n_games}")
print(f"Train: {train_games} partidas ({split_idx} activaciones)")
print(f"Test: {n_games - train_games} partidas ({len(activations) - split_idx} activaciones)")
print(f"Train set: {sae_features_train.shape}")
print(f"Test set: {sae_features_test.shape}")
```

```
Total partidas: 2000
Train: 1000 partidas (59000 activaciones)
Test: 1000 partidas (59000 activaciones)
Train set: torch.Size([59000, 8192])
Test set: torch.Size([59000, 8192])
```

```
# =====
# CONFIGURACIÓN DE FILTRADO DE BSPs
# =====
USE_PLAYER = False # BSPs perspectiva del jugador (1 y 2)
USE_ABSOLUTE = True # BSPs absolutas negro/blanco (B y W)
USE_STRATEGIC = False # BSPs especiales (BSP_)

bsp_piezas_indices = []
bsp_piezas_names = []

for i, name in enumerate(bsp_names):
```

```

include = False
if USE_PLAYER and len(name) == 6 and name.startswith('BSP') and (name.endswith('1') or name.endswith('2')):
    include = True
if USE_ABSOLUTE and (name.endswith('B') or name.endswith('W')):
    include = True
if USE_STRATEGIC and name.startswith('BSP_'):
    include = True
if include:
    bsp_pieces_indices.append(i)
    bsp_pieces_names.append(name)

bsp_pieces_indices = np.array(bsp_pieces_indices)

# Para reconstruction necesitas split train/test
bsp_gt_train_pieces = torch.from_numpy(bsp_ground_truth[:split_idx, bsp_pieces_indices]).bool().to(device)
bsp_gt_test_pieces = torch.from_numpy(bsp_ground_truth[split_idx:, bsp_pieces_indices]).bool().to(device)

# Convertir a tensor en GPU
# Identificador del filtrado (para naming de archivos de salida)
parts = []
if USE_PLAYER: parts.append("player")
if USE_ABSOLUTE: parts.append("absolute")
if USE_STRATEGIC: parts.append("strategic")
filter_suffix = "_".join(parts) if parts else "none"

print(f"BSPs seleccionadas para Reconstruction:")
print(f"  Player: {USE_PLAYER} | Absolute: {USE_ABSOLUTE} | Strategic: {USE_STRATEGIC}")
print(f"  Total: {len(bsp_pieces_indices)}")
print(f"  Train shape: {bsp_gt_train_pieces.shape}")
print(f"  Test shape: {bsp_gt_test_pieces.shape}")
print(f"  Device: {bsp_gt_train_pieces.device}")
print(f"  Primeras 5: {'', '.join(bsp_pieces_names[:5])}")
print(f"  Últimas 5: {'', '.join(bsp_pieces_names[-5:])}")
print(f"  Filter suffix: {filter_suffix}")

```

```

BSPs seleccionadas para Reconstruction:
Player: False | Absolute: True | Strategic: False
Total: 128
Train shape: torch.Size([59000, 128])
Test shape: torch.Size([59000, 128])
Device: cuda:0
Primeras 5: BSPA1B, BSPA1W, BSPA2B, BSPA2W, BSPA3B
Últimas 5: BSPH6W, BSPH7B, BSPH7W, BSPH8B, BSPH8W
Filter suffix: absolute

```

0.5 5. Funciones GPU-Optimizadas

```

def fast_precision_score_gpu(y_true, y_pred, eps=1e-8):
    """
    Precisión vectorizada en GPU.

    Args:
        y_true: Tensor (n_positions,) booleano
        y_pred: Tensor (n_positions, n_candidates) booleano

    Returns:
        precision_scores: Tensor (n_candidates,)
    """
    y_true_expanded = y_true.unsqueeze(1)

```

```

tp = (y_true_expanded & y_pred).sum(dim=0).float()
fp = (~y_true_expanded & y_pred).sum(dim=0).float()

precision = tp / (tp + fp + eps)

return precision

def fast_f1_score_gpu(y_true, y_pred, eps=1e-8):
    """
    F1 score vectorizado en GPU.

    Args:
        y_true: Tensor (n_positions,) booleano
        y_pred: Tensor (n_positions,) booleano

    Returns:
        f1: Float
    """
    tp = (y_true & y_pred).sum().float()
    fp = (~y_true & y_pred).sum().float()
    fn = (y_true & ~y_pred).sum().float()

    precision = tp / (tp + fp + eps)
    recall = tp / (tp + fn + eps)

    f1 = 2 * (precision * recall) / (precision + recall + eps)

    return f1.item()

print(" Funciones GPU definidas")

```

Funciones GPU definidas

0.6 6. Fase 1: Identificar Features de Alta Precisión (GPU)

Para cada BSP, encontrar features con precisión 0.95 en el train set.

```

def identify_high_precision_features_gpu(
    sae_features_train,
    bsp_gt_train,
    f_max,
    precision_threshold=0.95,
    significance_threshold=10,
    thresholds=[0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9],
    feature_batch_size=512,
    verbose=True
):
    """
    Para cada (threshold t, feature f, BSP b): determina si f es clasificador
    de alta precisión para b al threshold t.

    Criterios (igual que implementación oficial):
        - precision(f, b, t) >= precision_threshold (default 0.95)
        - on_count(f, t) >= significance_threshold (default 10)

    Args:
        sae_features_train: Tensor (P, n_features) GPU
        bsp_gt_train: Tensor (P, n_bsps) GPU, bool
        f_max: Tensor (n_features,) GPU - calculado FUERA para
        garantizar identidad exacta con Fase 2
    """

```

```

precision_threshold: float, default 0.95
significance_threshold: int, mín. activaciones requeridas (default 10)
thresholds: lista de fracciones t [0, 1)
feature_batch_size: int, features por mini-batch interno

Returns:
    hp_mask_TFB: BoolTensor (T, n_active, n_bsps)
    active_feature_indices: LongTensor (n_active,)
"""
n_positions, n_features = sae_features_train.shape
n_bsps = bsp_gt_train.shape[1]
device = sae_features_train.device

thresholds_T = torch.tensor(thresholds, device=device)
n_thresholds = len(thresholds_T)

# Features vivas (f_max > 0)
active_feature_indices = (f_max > 0).nonzero(as_tuple=True)[0]
n_active = len(active_feature_indices)

if verbose:
    print("=" * 60)
    print("FASE 1: IDENTIFICAR FEATURES DE ALTA PRECISIÓN")
    print("=" * 60)
    print(f"Precision threshold: {precision_threshold}")
    print(f"Significance threshold: >= {significance_threshold} activaciones")
    print(f"Features activas: {n_active}/{n_features} ({n_active/n_features*100:.1f}%)")
    print(f"BSPs: {n_bsps}")
    print(f"Thresholds: {thresholds}")
    print("=" * 60)

# Thresholds absolutos: thresh_TF[t, f] = t * f_max[f]
active_f_max = f_max[active_feature_indices] # (n_active,)
thresholds_TF = thresholds_T[:, None] * active_f_max # (T, n_active)

# Acumuladores sobre todo el train set
on_count_TF = torch.zeros(n_thresholds, n_active, device=device)
tp_count_TFB = torch.zeros(n_thresholds, n_active, n_bsps, device=device)

bsp_float_PB = bsp_gt_train.float() # (P, B)

n_iters = (n_active + feature_batch_size - 1) // feature_batch_size

for fi in tqdm(range(n_iters), desc="Fase 1"):
    f_start = fi * feature_batch_size
    f_end = min(f_start + feature_batch_size, n_active)

    global_idx = active_feature_indices[f_start:f_end]
    acts_PF = sae_features_train[:, global_idx] # (P, bs)
    thresh_TF = thresholds_T[:, f_start:f_end] # (T, bs)

    # active_TFP[t, f, p] = True si acts[p,f] > thresh[t,f]
    active_TFP = acts_PF.T.unsqueeze(0) > thresh_TF.unsqueeze(2) # (T, bs, P)

    # Conteo de activaciones por (threshold, feature)
    on_count_TF[:, f_start:f_end] += active_TFP.sum(dim=2).float()

    # True positives: f activa Y BSP verdadera
    # tp[t, f, b] = Σ_p active[t,f,p] * bsp[p,b]
    tp_count_TFB[:, f_start:f_end, :] += torch.einsum(
        'tfp,pb->tfb', active_TFP.float(), bsp_float_PB
    )

```

```

# Precisión: tp / on_count
eps = 1e-8
precision_TFB = tp_count_TFB / (on_count_TF.unsqueeze(2) + eps) # (T, F, B)

# Máscara: alta precisión AND significancia suficiente
sig_mask_TF1 = (on_count_TF >= significance_threshold).unsqueeze(2) # (T, F, 1)
hp_mask_TFB = (precision_TFB >= precision_threshold) & sig_mask_TF1 # (T, F, B)

if verbose:
    hp_per_t = hp_mask_TFB.sum(dim=(1, 2)) # (T,)
    hp_per_b = hp_mask_TFB.sum(dim=1) # (T, B)
    best_t = hp_per_t.argmax().item()
    print(f"\nPares (feature, BSP) de alta precisión por threshold:")
    for i, t in enumerate(thresholds):
        marker = " ← más pares" if i == best_t else ""
        print(f" t={t:.1f}: {hp_per_t[i].item():5d} pares{marker}")
    print(f"\nFeatures HP por BSP en t={thresholds[best_t]:.1f}:")
    print(f" Media: {hp_per_b[best_t].float().mean():.1f}")
    print(f" Máx: {hp_per_b[best_t].max().item()}")
    print(f" BSPs sin features: {(hp_per_b[best_t] == 0).sum().item()}")

return hp_mask_TFB, active_feature_indices

print(" Fase 1 rediseñada: umbral de significancia (10) + estructura (T, F, B)")

```

Fase 1 rediseñada: umbral de significancia (10) + estructura (T, F, B)

```

THRESHOLDS = torch.arange(0.0, 1.05, 0.05).tolist()

# f_max calculado UNA SOLA VEZ aquí - se pasa a Fase 1 y Fase 2
f_max = sae_features_train.max(dim=0)[0] # (n_features,)
print(f"f_max calculado: {(f_max > 0).sum().item()} features activas de {f_max.shape[0]}")

start_time = time.time()

hp_mask_TFB, active_feature_indices = identify_high_precision_features_gpu(
    sae_features_train,
    bsp_gt_train_pieces,
    f_max,
    precision_threshold=0.95,
    significance_threshold=10,
    thresholds=THRESHOLDS,
    feature_batch_size=512,
    verbose=True
)

phase1_time = time.time() - start_time

print(f"\nTiempo Fase 1: {phase1_time:.2f} seg")
print(f"hp_mask_TFB shape: {hp_mask_TFB.shape} (T, n_active, n_bsps)")

```

f_max calculado: 8192 features activas de 8192

=====

FASE 1: IDENTIFICAR FEATURES DE ALTA PRECISIÓN

=====

Precision threshold: 0.95

Significance threshold: >= 10 activaciones

Features activas: 8192/8192 (100.0%)

BSPs: 128

Thresholds: [0.0, 0.05000000074505806, 0.10000000149011612, 0.15000000596046448, 0.20000000298023224, 0.25000000596046448, 0.30000000796046448, 0.35000001196046448, 0.40000001596046448, 0.45000001996046448, 0.50000002396046448, 0.55000002796046448, 0.60000003196046448, 0.65000003596046448, 0.70000003996046448, 0.75000004396046448, 0.80000004796046448, 0.85000005196046448, 0.90000005596046448, 0.95000005996046448, 1.00000006396046448]

=====

Fase 1: 100% | 16/16 [01:07<00:00, 4.24s/it]

Pares (feature, BSP) de alta precisión por threshold:

```
t=0.0: 20 pares
t=0.1: 69 pares
t=0.1: 174 pares
t=0.2: 417 pares
t=0.2: 735 pares
t=0.2: 1051 pares
t=0.3: 1293 pares
t=0.3: 1475 pares
t=0.4: 1561 pares
t=0.5: 1603 pares
t=0.5: 1656 pares ← más pares
t=0.6: 1604 pares
t=0.6: 1501 pares
t=0.6: 1344 pares
t=0.7: 1291 pares
t=0.8: 1179 pares
t=0.8: 1078 pares
t=0.9: 954 pares
t=0.9: 864 pares
t=0.9: 842 pares
t=1.0: 0 pares
```

Features HP por BSP en t=0.5:

```
Media: 12.9
Máx: 186
BSPs sin features: 52
```

Tiempo Fase 1: 70.03 seg

hp_mask_TFB shape: torch.Size([21, 8192, 128]) (T, n_active, n_bsps)

0.7 7. Fase 2: Reconstruir Tableros en Test Set (GPU)

```
def reconstruct_boards_gpu(
    sae_features_test,
    bsp_gt_test,
    hp_mask_TFB,
    active_feature_indices,
    f_max,
    thresholds=[0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9],
    verbose=True
):
    """
    Fase 2: barrido simultáneo de thresholds T → toma el mejor F1.

    Para cada threshold t:
    1. Binarizar features de test: active[p, f] = acts[p, f] > t * f_max[f]
    2. Predicción por (posición, BSP):
        predict[p, b] = OR_f ( active[p,f] AND hp_mask[t, f, b] )
        implementado como: (active_PF @ hp_FB) > 0
    3. F1 global micro-averaged (acumula TP/FP/FN sobre todas las posiciones y BSPs)

    Retorna el max F1 sobre todos los thresholds (Eq. 5 del paper: max_t).

    Args:
        sae_features_test: Tensor (P, n_features) GPU
        bsp_gt_test: Tensor (P, n_bsps) GPU, bool
```

```

hp_mask_TFB:          BoolTensor (T, n_active, n_bsps) de Fase 1
active_feature_indices: LongTensor (n_active,) de Fase 1
f_max:                Tensor (n_features,) GPU - MISMO objeto que Fase 1
thresholds:            lista de fracciones (debe coincidir con Fase 1)

Returns:
    best_f1:            float, max F1 sobre thresholds
    best_threshold:     float, threshold óptimo
    f1_per_threshold:   np.array (T,), F1 por threshold
"""
device = sae_features_test.device
n_thresholds = len(thresholds)
thresholds_T = torch.tensor(thresholds, device=device)

# Activaciones del test solo para features activas
test_active_PF = sae_features_test[:, active_feature_indices] # (P, n_active)
active_f_max = f_max[active_feature_indices] # (n_active,)

gt_PB = bsp_gt_test.bool() # (P, B)
eps = 1e-8

f1_scores = torch.zeros(n_thresholds, device=device)
precision_T = torch.zeros(n_thresholds, device=device)
recall_T = torch.zeros(n_thresholds, device=device)

if verbose:
    print("=" * 60)
    print("FASE 2: RECONSTRUCCIÓN - BARRIDO DE THRESHOLDS")
    print("=" * 60)
    print(f"{'t':>5}  {'F1':>8}  {'P':>8}  {'R':>8}  {'TP':>8}  {'FP':>8}  {'FN':>8}")
    print("-" * 60)

for t_idx in range(n_thresholds):
    t = thresholds[t_idx]

    # Threshold absoluto por feature: t * f_max[f]
    abs_thresh_F = thresholds_T[t_idx] * active_f_max # (n_active,)

    # Binarizar features en test
    active_PF = test_active_PF > abs_thresh_F.unsqueeze(0) # (P, n_active)

    # predict[p, b] = any( active[p,f] AND hp[t,f,b] )
    # = (active_PF.float() @ hp_mask[t].float()) > 0
    hp_FB = hp_mask_TFB[t_idx].float() # (n_active, B)
    predictions_PB = (active_PF.float() @ hp_FB) > 0 # (P, B)

    # Métricas globales micro-averaged
    tp = (predictions_PB & gt_PB).sum().float()
    fp = (predictions_PB & ~gt_PB).sum().float()
    fn = (~predictions_PB & gt_PB).sum().float()

    p = tp / (tp + fp + eps)
    r = tp / (tp + fn + eps)
    f1 = 2 * p * r / (p + r + eps)

    f1_scores[t_idx] = f1
    precision_T[t_idx] = p
    recall_T[t_idx] = r

    if verbose:
        print(f"{'t':>5.1f}  {'f1.item():>8.4f}  {'p.item():>8.4f}  {'r.item():>8.4f}"
              f"  {'tp.item():>8.0f}  {'fp.item():>8.0f}  {'fn.item():>8.0f}")

```

```

best_idx      = f1_scores.argmax().item()
best_f1       = f1_scores[best_idx].item()
best_threshold = thresholds[best_idx]

if verbose:
    print("=" * 60)
    print(f"Mejor threshold: t={best_threshold:.1f}")
    print(f"Reconstruction Score (max_t F1): {best_f1:.4f}")
    print("=" * 60)

return best_f1, best_threshold, f1_scores.cpu().numpy()

print(" Fase 2 rediseñada: barrido T simultáneo + F1 global micro-averaged + max_t")

```

Fase 2 rediseñada: barrido T simultáneo + F1 global micro-averaged + max_t

```

start_time = time.time()

reconstruction_score, best_threshold, f1_per_threshold = reconstruct_boards_gpu(
    sae_features_test,
    bsp_gt_test_pieces,
    hp_mask_TFB,
    active_feature_indices,
    f_max,          # mismo tensor calculado antes de Fase 1
    thresholds=THRESHOLDS,
    verbose=True
)

phase2_time = time.time() - start_time
total_time  = phase1_time + phase2_time

print(f"\nTiempo Fase 2: {phase2_time:.2f} seg")
print(f"Tiempo total:  {total_time:.2f} seg ({total_time/60:.2f} min)")

```

```

=====
FASE 2: RECONSTRUCCIÓN - BARRIDO DE THRESHOLDS
=====

```

t	F1	P	R	TP	FP	FN
0.0	0.0026	0.9643	0.0013	2569	95	2003431
0.1	0.0100	0.9746	0.0050	10097	263	1995903
0.1	0.0151	0.9688	0.0076	15266	492	1990734
0.2	0.0216	0.9647	0.0109	21952	804	1984048
0.2	0.0266	0.9668	0.0135	27102	931	1978898
0.2	0.0306	0.9639	0.0156	31237	1170	1974763
0.3	0.0321	0.9675	0.0163	32767	1101	1973233
0.3	0.0329	0.9707	0.0168	33609	1013	1972391
0.4	0.0323	0.9725	0.0164	32897	929	1973103
0.5	0.0317	0.9745	0.0161	32348	845	1973652
0.5	0.0310	0.9711	0.0157	31549	940	1974451
0.6	0.0291	0.9750	0.0148	29619	758	1976381
0.6	0.0272	0.9774	0.0138	27628	639	1978372
0.6	0.0250	0.9796	0.0127	25380	529	1980620
0.7	0.0237	0.9829	0.0120	24099	419	1981901
0.8	0.0222	0.9876	0.0112	22531	284	1983469
0.8	0.0211	0.9896	0.0107	21404	224	1984596
0.9	0.0200	0.9955	0.0101	20253	91	1985747
0.9	0.0191	0.9987	0.0096	19313	25	1986687
0.9	0.0187	0.9989	0.0094	18941	20	1987059
1.0	0.0000	0.0000	0.0000	0	0	2006000

```

=====

```

```
Mejor threshold: t=0.3
Reconstruction Score (max_t F1): 0.0329
```

```
=====

Tiempo Fase 2: 10.33 seg
Tiempo total: 80.36 seg (1.34 min)
```

0.8 8. Análisis de Resultados


```
print("=" * 60)
print("BARRIDO DE THRESHOLDS - F1 POR THRESHOLD")
print("=" * 60)
for i, (t, f1) in enumerate(zip(THRESHOLDS, f1_per_threshold)):
    marker = " ← MEJOR" if i == f1_per_threshold.argmax() else ""
    print(f"t={t:.1f}: F1={f1:.4f}{marker}")
print("=" * 60)

plt.figure(figsize=(10, 6))
plt.plot(THRESHOLDS, f1_per_threshold, marker='o', linewidth=2, markersize=8)
plt.axhline(y=reconstruction_score, color='r', linestyle='--',
            label=f'Best F1={reconstruction_score:.4f} (t={best_threshold:.1f})')
plt.axhline(y=0.95, color='g', linestyle=':', label='Target Othello SAE (paper)=0.95')
plt.xlabel('Threshold t', fontsize=12)
plt.ylabel('F1 Score (micro-averaged global)', fontsize=12)
plt.title('Board Reconstruction: F1 por Threshold', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

print(f"\n Threshold óptimo: t={best_threshold:.1f}")
print(f" Reconstruction Score: {reconstruction_score:.4f}")
```

```
=====
BARRIDO DE THRESHOLDS - F1 POR THRESHOLD
=====
```

```
t=0.0: F1=0.0026
t=0.1: F1=0.0100
t=0.1: F1=0.0151
t=0.2: F1=0.0216
t=0.2: F1=0.0266
t=0.2: F1=0.0306
t=0.3: F1=0.0321
t=0.3: F1=0.0329 ← MEJOR
t=0.4: F1=0.0323
t=0.5: F1=0.0317
t=0.5: F1=0.0310
t=0.6: F1=0.0291
t=0.6: F1=0.0272
t=0.6: F1=0.0250
t=0.7: F1=0.0237
t=0.8: F1=0.0222
t=0.8: F1=0.0211
t=0.9: F1=0.0200
t=0.9: F1=0.0191
t=0.9: F1=0.0187
t=1.0: F1=0.0000
=====
```



Threshold óptimo: t=0.3
Reconstruction Score: 0.0329

```
# Estadísticas sobre features de alta precisión por BSP
hp_per_bsp = hp_mask_TFB.sum(dim=1) # (T, n_bsps)
best_t_idx = torch.tensor(f1_per_threshold).argmax().item()
n_features_per_bsp = hp_per_bsp[best_t_idx].cpu().numpy()

print("="*60)
print("ESTADÍSTICAS - FEATURES POR BSP")
print("="*60)
print(f"Total BSPs: {len(n_features_per_bsp)}")
print(f"Features promedio por BSP: {np.mean(n_features_per_bsp):.1f}")
print(f"Features mediana por BSP: {np.median(n_features_per_bsp):.1f}")
print(f"BSPs con 0 features: {np.sum(n_features_per_bsp == 0)}")
print(f"BSPs con 1+ features: {np.sum(n_features_per_bsp > 0)}")
print(f"BSPs con 5+ features: {np.sum(n_features_per_bsp >= 5)}")
print(f"BSPs con 10+ features: {np.sum(n_features_per_bsp >= 10)}")
print(f"Max features en una BSP: {np.max(n_features_per_bsp)}")
print("="*60)
```

```
=====
ESTADÍSTICAS - FEATURES POR BSP
=====
Total BSPs: 128
Features promedio por BSP: 11.5
Features mediana por BSP: 1.0
BSPs con 0 features: 58
BSPs con 1+ features: 70
BSPs con 5+ features: 43
BSPs con 10+ features: 28
Max features en una BSP: 169
=====
```

```
# Comparación con resultados del paper
print("="*60)
print("COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)")
print("="*60)
print(f"{'Modelo':<20} {'Reconstruction':<15}")
print("-"*60)
print(f"{'SAE random':<20} {'0.08':<15}")
print(f"{'SAE trained (paper)':<20} {'0.95':<15} ← Objetivo")
print(f"{'Linear probe':<20} {'0.99':<15}")
print(f"{'Nuestro SAE':<20} {'reconstruction_score:<15.4f'} ← Resultado")
print("="*60)

if reconstruction_score >= 0.90:
    print("\n Excelente: Muy cercano al objetivo del paper")
elif reconstruction_score >= 0.80:
    print("\n Bueno: Por encima de random pero debajo del objetivo")
else:
    print("\n Bajo: Revisar implementación o entrenamiento del SAE")
```

```
=====
COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)
=====
Modelo                Reconstruction
-----
SAE random            0.08
SAE trained (paper)   0.95                ← Objetivo
Linear probe          0.99
Nuestro SAE           0.0329            ← Resultado
=====
```

Bajo: Revisar implementación o entrenamiento del SAE

0.9 9. Guardar Resultados

```
results = {
    'reconstruction_score': reconstruction_score,
    'best_threshold': best_threshold,
    'f1_per_threshold': f1_per_threshold,
    'thresholds': np.array(THRESHOLDS),
    'phase1_time_seconds': phase1_time,
    'phase2_time_seconds': phase2_time,
    'total_time_seconds': total_time,
    'bsp_names': bsp_pieces_names,
}

metrics_dir = get_metrics_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)
output_path = metrics_dir / f"reconstruction_{filter_suffix}.npz"
np.savez(output_path, **results)

print(f" Resultados guardados en: {output_path}")
print(f"\nContenido:")
print(f" reconstruction_score: {reconstruction_score:.4f}")
print(f" best_threshold: {best_threshold:.1f}")
```

```

print(f"  fl_per_threshold:      {fl_per_threshold}")

# Comparación con el paper
print()
print("=" * 60)
print("COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)")
print("=" * 60)
print(f"{'Modelo':<25} {'Reconstruction':>14}")
print("-" * 40)
print(f"{'SAE random':<25} {'0.08':>14}")
print(f"{'SAE trained (paper)':<25} {'0.95':>14} ← Objetivo Othello")
print(f"{'Linear probe':<25} {'0.99':>14}")
print(f"{'Nuestro SAE':<25} {reconstruction_score:>14.4f} ← Resultado")
print("=" * 60)

if reconstruction_score >= 0.90:
    print("\n Excelente: Muy cercano al objetivo del paper")
elif reconstruction_score >= 0.50:
    print("\n~ Moderado: Por encima de random, revisar SAE")
else:
    print("\n Bajo: Revisar entrenamiento del SAE o calidad de las activaciones")

```

Resultados guardados en: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\sae_p-annealing_standard\12500games\ex5\metrics\reconstruction_absolute.npz

Contenido:

```

reconstruction_score: 0.0329
best_threshold:      0.3
fl_per_threshold:    [0.00255792 0.01001508 0.01510171 0.02164085 0.02664853 0.03064844 0.03212659 0.03293996 0.0322
0.03172627 0.03095332 0.0290899  0.02716261 0.02498144 0.0237368  0.02221099 0.02111235 0.0199897
0.01907138 0.01870752 0.          ]

```

```

=====
COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)
=====
Modelo                Reconstruction
-----
SAE random            0.08
SAE trained (paper)   0.95 ← Objetivo Othello
Linear probe          0.99
Nuestro SAE           0.0329 ← Resultado
=====

```

Bajo: Revisar entrenamiento del SAE o calidad de las activaciones

```

import subprocess
from pathlib import Path
from sae.tools.experiment_utils import get_report_name

output_dir = get_metrics_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

pdf_name = get_report_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],

```

```

NAMING_META['juegos'],
f'reconstruction_{filter_suffix}',
extras_id=extras_id
)

notebook_name = 'reconstruction_sae_standar2.ipynb'
notebook_abs = str(Path('.').resolve() / notebook_name)
output_path = output_dir / pdf_name

print(f' Generando PDF con Quarto...')
print(f' Archivo de salida: {output_path}')

result = subprocess.run(
    f'quarto render {notebook_abs} --to pdf --output-dir "{output_dir}" --output {pdf_name}',
    shell=True,
    capture_output=True,
    text=True
)

if result.returncode == 0:
    print(f' PDF generado exitosamente: {output_path}')
else:
    print(f' Error al generar PDF:')
    print(result.stderr)

```

```

Generando PDF con Quarto...
Archivo de salida: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\sae_p-
annealing_standard\12500games\ex5\metrics\sae_standard_p-annealing_16_12500g_ex5_reconstruction_absolute.pdf
PDF generado exitosamente: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_standard\sae_p-
annealing_standard\12500games\ex5\metrics\sae_standard_p-annealing_16_12500g_ex5_reconstruction_absolute.pdf

```