

Reconstruction Metric - GPU Optimized

2026-05-19

Table of contents

0.1	1. Setup e Imports	1
0.2	2. Cargar Datos	3
0.3	3. Cargar SAE y Extraer Features	4
0.4	4. Split Train/Test y Filtrar BSPs de Piezas	5
0.5	5. Funciones GPU-Optimizadas	6
0.6	6. Fase 1: Identificar Features de Alta Precisión (GPU)	7
0.7	7. Fase 2: Reconstruir Tableros en Test Set (GPU)	10
0.8	8. Análisis de Resultados	13
0.9	9. Guardar Resultados	15
0.10	10. Generar PDF con Quarto	16

0.1 1. Setup e Imports

```
import numpy as np
import pandas as pd
import sys
import os
import subprocess
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt

from pathlib import Path
from tqdm import tqdm
from sae.tools.experiment_utils import get_experiment_dir, get_metrics_dir, get_report_name
from sae.tools.naming_utils import get_checkpoint_dir, get_model_name, get_extras_id

project_root = Path('.../.../...').resolve()
sys.path.insert(0, str(project_root))

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Device: {device}")
if torch.cuda.is_available():
    print(f"GPU: {torch.cuda.get_device_name(0)}")
    print(f"Memoria: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
```

Device: cuda
GPU: NVIDIA GeForce RTX 5060
Memoria: 8.55 GB

```
# Configuración para visualización en PDF
pd.set_option('display.max_colwidth', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 100)
```

```

np.set_printoptions(linewidth=100, edgeitems=3)

os.environ['COLUMNS'] = '100'

print(" Configuración para PDF lista")

```

Configuración para PDF lista

```

NAMING_META = {
    'variante': 'topk',
    'tecnica': 'batch-topk',
    'capa': 6,
    'juegos': 25000,
    'base_path': 'C:\\Users\\Esposa\\Documents\\Repos\\sae-othello-gpt',
    'experiment_base_path': os.path.abspath('../../../experiments'),
    'extras': {
        'expansion': 16,
        'k': 81,
        'epochs': 1000,
        'patience': 20,
        'batch_size_tokens': 1024,
    }
}

print('\n Metadata de naming:')
for key, value in NAMING_META.items():
    print(f' {key}: {value}')

```

Metadata de naming:

```

variante: topk
tecnica: batch-topk
capa: 6
juegos: 25000
base_path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt
experiment_base_path: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments
extras: {'expansion': 16, 'k': 81, 'epochs': 1000, 'patience': 20, 'batch_size_tokens': 1024}

```

```

class BatchTopKSAE(nn.Module):
    def __init__(self, d_in: int, d_sae: int, k: int, k_aux: int = 512):
        super().__init__()
        self.d_in = d_in
        self.d_sae = d_sae
        self.k = k
        self.k_aux = k_aux

        self.W_enc = nn.Parameter(torch.empty(d_in, d_sae))
        self.W_dec = nn.Parameter(torch.empty(d_sae, d_in))
        self.b_dec = nn.Parameter(torch.zeros(d_in))

        nn.init.kaiming_uniform_(self.W_enc)
        nn.init.kaiming_uniform_(self.W_dec)
        self.W_dec.data = self.W_dec.data / self.W_dec.data.norm(dim=-1, keepdim=True)
        self.register_buffer('num_batches_not_active', torch.zeros(d_sae))

    def encode(self, x: torch.Tensor, threshold=None) -> torch.Tensor:
        x_norm, _, _ = self.preprocess_input(x)
        x_cent = x_norm - self.b_dec
        pre_acts = F.relu(x_cent @ self.W_enc)
        return self._topk_activation(pre_acts, threshold=threshold)

```

```

def _topk_activation(self, x: torch.Tensor, threshold=None) -> torch.Tensor:
    if threshold is None:
        acts_topk = torch.topk(x.flatten(), self.k * x.shape[0], dim=-1)
        return (
            torch.zeros_like(x.flatten())
            .scatter(-1, acts_topk.indices, acts_topk.values)
            .reshape(x.shape)
        )
    return torch.where(x > threshold, x, torch.zeros_like(x))

def decode(self, hidden: torch.Tensor) -> torch.Tensor:
    return hidden @ self.W_dec + self.b_dec

def preprocess_input(self, x: torch.Tensor):
    x_mean = x.mean(dim=-1, keepdim=True)
    x = x - x_mean
    x_std = x.std(dim=-1, keepdim=True)
    x = x / (x_std + 1e-5)
    return x, x_mean, x_std

def postprocess_output(self, x_reconstruct, x_mean, x_std):
    return x_reconstruct * x_std + x_mean

def forward(self, x: torch.Tensor, threshold=None):
    x_norm, x_mean, x_std = self.preprocess_input(x)
    x_cent = x_norm - self.b_dec
    pre_acts = F.relu(x_cent @ self.W_enc)
    hidden = self._topk_activation(pre_acts, threshold=threshold)
    recon = self.postprocess_output(self.decode(hidden), x_mean, x_std)
    return recon, hidden, pre_acts

print("BatchTopKSAE definida")

```

BatchTopKSAE definida

0.2 2. Cargar Datos

```

# Cargar activaciones
activations_path = project_root / "activations" / "data" / "layer6_2000games.npy"
activations = np.load(activations_path)

print(f"Activaciones del modelo:")
print(f" Shape: {activations.shape}")
print(f" Memoria: {activations.nbytes / (1024**2):.2f} MB")

```

Activaciones del modelo:

```

Shape: (118000, 512)
Memoria: 230.47 MB

```

```

# Cargar ground truth de BSPs
bsp_gt_path = project_root / "metrics" / "02_data" / "bsp_ground_truth_2000games.npy"
bsp_names_path = project_root / "metrics" / "02_data" / "bsp_ground_truth_2000games.names.npy"

bsp_ground_truth = np.load(bsp_gt_path)
bsp_names = np.load(bsp_names_path, allow_pickle=True)

print(f"\nGround truth BSPs:")
print(f" Shape: {bsp_ground_truth.shape}")
print(f" Total BSPs: {len(bsp_names)}")

```

Ground truth BSPs:
Shape: (118000, 326)
Total BSPs: 326

0.3 3. Cargar SAE y Extraer Features

```
# Configuración del SAE
input_dim = 512
hidden_dim = 8192

# Construir ruta del modelo usando naming_utils
checkpoint_dir = get_checkpoint_dir(
    NAMING_META['base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos']
)
extras_id = get_extras_id(checkpoint_dir, NAMING_META['extras']) if NAMING_META.get('extras') else None

model_filename = get_model_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    'best',
    extras_id=extras_id
)
model_path = checkpoint_dir / model_filename

# Cargar modelo BatchTopK
sae = BatchTopKSAE(d_in=input_dim, d_sae=hidden_dim, k=NAMING_META['extras']['k']).to(device)

checkpoint = torch.load(model_path, map_location=device, weights_only=False)
sae.load_state_dict(checkpoint['model_state_dict'])
sae.eval()

print(f"SAE cargado:")
print(f"  Path: {model_path}")
print(f"  Input: {input_dim}, Hidden: {hidden_dim}, k: {NAMING_META['extras']['k']}")
print(f"  Expansion: {hidden_dim/input_dim}x")
```

SAE cargado:
Path: C:\Users\Esposa\Documents\Repos\sae-othello-gpt\layer_06\models\sae_topk\sae_batch-topk_topk\25000games\sae_topk_l6_25000g_ex5_best.pt
Input: 512, Hidden: 8192, k: 81
Expansion: 16.0x

```
# Extraer features del SAE en GPU
print("Extrayendo features del SAE...")
activations_tensor = torch.from_numpy(activations).float().to(device)

# Estimar threshold con mini-batch
with torch.no_grad():
    sample = activations_tensor[:4096]
    x_s, _, _ = sae.preprocess_input(sample)
    pre_s = F.relu(x_s - sae.b_dec) @ sae.W_enc
    threshold = torch.topk(pre_s.flatten(), sae.k * 4096).values.min().item()
print(f"Threshold estimado: {threshold:.6f}")
```

```
# Encoding completo con operación in-place
with torch.no_grad():
    x_norm, _, _ = sae.preprocess_input(activations_tensor)
    x_cent = x_norm - sae.b_dec
    pre_acts = F.relu(x_cent @ sae.W_enc)
    pre_acts.mul_(pre_acts > threshold)
    sae_features = pre_acts

del activations_tensor, x_norm, x_cent
torch.cuda.empty_cache()

print(f"\nFeatures SAE:")
print(f"  Shape: {sae_features.shape}")
print(f"  Device: {sae_features.device}")
print(f"  Sparsity: {(sae_features == 0).float().mean():.2%}")
print(f"  Activaciones promedio: {(sae_features > 0).sum(dim=1).float().mean():.1f}")
```

Extrayendo features del SAE...
Threshold estimado: 0.886375

Features SAE:
Shape: torch.Size([118000, 8192])
Device: cuda:0
Sparsity: 99.01%
Activaciones promedio: 81.0

0.4 4. Split Train/Test y Filtrar BSPs de Piezas

```
# Split 50/50:
n_moves = 59
n_games = 2000 # número de partidas para prueba de métricas
train_games = n_games // 2
split_idx = train_games * n_moves

sae_features_train = sae_features[:split_idx]
sae_features_test = sae_features[split_idx:]

print(f"Total partidas: {n_games}")
print(f"Train: {train_games} partidas ({split_idx} activaciones)")
print(f"Test: {n_games - train_games} partidas ({len(activations) - split_idx} activaciones)")
print(f"Train set: {sae_features_train.shape}")
print(f"Test set: {sae_features_test.shape}")
```

Total partidas: 2000
Train: 1000 partidas (59000 activaciones)
Test: 1000 partidas (59000 activaciones)
Train set: torch.Size([59000, 8192])
Test set: torch.Size([59000, 8192])

```
# =====
# CONFIGURACIÓN DE FILTRADO DE BSPs
# =====
USE_PLAYER = False # BSPs perspectiva del jugador (1 y 2)
USE_ABSOLUTE = True # BSPs absolutas negro/blanco (B y W)
USE_STRATEGIC = False # BSPs especiales (BSP_)

bsp_piezas_indices = []
bsp_piezas_names = []
```

```

for i, name in enumerate(bsp_names):
    include = False
    if USE_PLAYER and len(name) == 6 and name.startswith('BSP') and (name.endswith('1') or name.endswith('2')):
        include = True
    if USE_ABSOLUTE and (name.endswith('B') or name.endswith('W')):
        include = True
    if USE_STRATEGIC and name.startswith('BSP_'):
        include = True
    if include:
        bsp_pieces_indices.append(i)
        bsp_pieces_names.append(name)

bsp_pieces_indices = np.array(bsp_pieces_indices)

# Para reconstruction necesitas split train/test
bsp_gt_train_pieces = torch.from_numpy(bsp_ground_truth[:split_idx, bsp_pieces_indices]).bool().to(device)
bsp_gt_test_pieces = torch.from_numpy(bsp_ground_truth[split_idx:, bsp_pieces_indices]).bool().to(device)

# Convertir a tensor en GPU
# Identificador del filtrado (para naming de archivos de salida)
parts = []
if USE_PLAYER: parts.append("player")
if USE_ABSOLUTE: parts.append("absolute")
if USE_STRATEGIC: parts.append("strategic")
filter_suffix = "_".join(parts) if parts else "none"

print(f"BSPs seleccionadas para Reconstruction:")
print(f" Player: {USE_PLAYER} | Absolute: {USE_ABSOLUTE} | Strategic: {USE_STRATEGIC}")
print(f" Total: {len(bsp_pieces_indices)}")
print(f" Train shape: {bsp_gt_train_pieces.shape}")
print(f" Test shape: {bsp_gt_test_pieces.shape}")
print(f" Device: {bsp_gt_train_pieces.device}")
print(f" Primeras 5: {'', '.join(bsp_pieces_names[:5])}")
print(f" Últimas 5: {'', '.join(bsp_pieces_names[-5:])}")
print(f" Filter suffix: {filter_suffix}")

```

```

BSPs seleccionadas para Reconstruction:
Player: False | Absolute: True | Strategic: False
Total: 128
Train shape: torch.Size([59000, 128])
Test shape: torch.Size([59000, 128])
Device: cuda:0
Primeras 5: BSPA1B, BSPA1W, BSPA2B, BSPA2W, BSPA3B
Últimas 5: BSPH6W, BSPH7B, BSPH7W, BSPH8B, BSPH8W
Filter suffix: absolute

```

0.5 5. Funciones GPU-Optimizadas

```

def fast_precision_score_gpu(y_true, y_pred, eps=1e-8):
    """
    Precisión vectorizada en GPU.

    Args:
        y_true: Tensor (n_positions,) booleano
        y_pred: Tensor (n_positions, n_candidates) booleano

    Returns:
        precision_scores: Tensor (n_candidates,)
    """
    y_true_expanded = y_true.unsqueeze(1)

```

```

tp = (y_true_expanded & y_pred).sum(dim=0).float()
fp = (~y_true_expanded & y_pred).sum(dim=0).float()

precision = tp / (tp + fp + eps)

return precision

def fast_f1_score_gpu(y_true, y_pred, eps=1e-8):
    """
    F1 score vectorizado en GPU.

    Args:
        y_true: Tensor (n_positions,) booleano
        y_pred: Tensor (n_positions,) booleano

    Returns:
        f1: Float
    """
    tp = (y_true & y_pred).sum().float()
    fp = (~y_true & y_pred).sum().float()
    fn = (y_true & ~y_pred).sum().float()

    precision = tp / (tp + fp + eps)
    recall = tp / (tp + fn + eps)

    f1 = 2 * (precision * recall) / (precision + recall + eps)

    return f1.item()

print(" Funciones GPU definidas")

```

Funciones GPU definidas

0.6 6. Fase 1: Identificar Features de Alta Precisión (GPU)

Para cada BSP, encontrar features con precisión 0.95 en el train set.

```

def identify_high_precision_features_gpu(
    sae_features_train,
    bsp_gt_train,
    f_max,
    precision_threshold=0.95,
    significance_threshold=1,
    thresholds=[0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9],
    feature_batch_size=512,
    verbose=True
):
    """
    Para cada (threshold t, feature f, BSP b): determina si f es clasificador
    de alta precisión para b al threshold t.

    Criterios (igual que implementación oficial):
        - precision(f, b, t) >= precision_threshold (default 0.95)
        - on_count(f, t) >= significance_threshold (default 10)

    Args:
        sae_features_train: Tensor (P, n_features) GPU
        bsp_gt_train: Tensor (P, n_bsps) GPU, bool
        f_max: Tensor (n_features,) GPU - calculado FUERA para
    """

```

```

        garantizar identidad exacta con Fase 2
precision_threshold: float, default 0.95
significance_threshold: int, mín. activaciones requeridas (default 10)
thresholds: lista de fracciones t [0, 1)
feature_batch_size: int, features por mini-batch interno

Returns:
    hp_mask_TFB: BoolTensor (T, n_active, n_bsps)
    active_feature_indices: LongTensor (n_active,)
"""
n_positions, n_features = sae_features_train.shape
n_bsps = bsp_gt_train.shape[1]
device = sae_features_train.device

thresholds_T = torch.tensor(thresholds, device=device)
n_thresholds = len(thresholds_T)

# Features vivas (f_max > 0)
active_feature_indices = (f_max > 0).nonzero(as_tuple=True)[0]
n_active = len(active_feature_indices)

if verbose:
    print("=" * 60)
    print("FASE 1: IDENTIFICAR FEATURES DE ALTA PRECISIÓN")
    print("=" * 60)
    print(f"Precision threshold: {precision_threshold}")
    print(f"Significance threshold: >= {significance_threshold} activaciones")
    print(f"Features activas: {n_active}/{n_features} ({n_active/n_features*100:.1f}%)")
    print(f"BSPs: {n_bsps}")
    print(f"Thresholds: {thresholds}")
    print("=" * 60)

# Thresholds absolutos: thresh_TF[t, f] = t * f_max[f]
active_f_max = f_max[active_feature_indices] # (n_active,)
thresholds_TF = thresholds_T[:, None] * active_f_max # (T, n_active)

# Acumuladores sobre todo el train set
on_count_TF = torch.zeros(n_thresholds, n_active, device=device)
tp_count_TFB = torch.zeros(n_thresholds, n_active, n_bsps, device=device)

bsp_float_PB = bsp_gt_train.float() # (P, B)

n_iters = (n_active + feature_batch_size - 1) // feature_batch_size

for fi in tqdm(range(n_iters), desc="Fase 1"):
    f_start = fi * feature_batch_size
    f_end = min(f_start + feature_batch_size, n_active)

    global_idx = active_feature_indices[f_start:f_end]
    acts_PF = sae_features_train[:, global_idx] # (P, bs)
    thresh_TF = thresholds_TF[:, f_start:f_end] # (T, bs)

    # active_TFP[t, f, p] = True si acts[p,f] > thresh[t,f]
    active_TFP = acts_PF.T.unsqueeze(0) > thresh_TF.unsqueeze(2) # (T, bs, P)

    # Conteo de activaciones por (threshold, feature)
    on_count_TF[:, f_start:f_end] += active_TFP.sum(dim=2).float()

    # True positives: f activa Y BSP verdadera
    # tp[t, f, b] = Σ_p active[t,f,p] * bsp[p,b]
    tp_count_TFB[:, f_start:f_end, :] += torch.einsum(
        'tfp,pb->tfb', active_TFP.float(), bsp_float_PB

```



```

    )

    # Precisión: tp / on_count
    eps = 1e-8
    precision_TFB = tp_count_TFB / (on_count_TF.unsqueeze(2) + eps) # (T, F, B)

    # Máscara: alta precisión AND significancia suficiente
    sig_mask_TF1 = (on_count_TF >= significance_threshold).unsqueeze(2) # (T, F, 1)
    hp_mask_TFB = (precision_TFB >= precision_threshold) & sig_mask_TF1 # (T, F, B)

    if verbose:
        hp_per_t = hp_mask_TFB.sum(dim=(1, 2)) # (T,)
        hp_per_b = hp_mask_TFB.sum(dim=1) # (T, B)
        best_t = hp_per_t.argmax().item()
        print(f"\nPares (feature, BSP) de alta precisión por threshold:")
        for i, t in enumerate(thresholds):
            marker = " ← más pares" if i == best_t else ""
            print(f" t={t:.1f}: {hp_per_t[i].item():5d} pares{marker}")
        print(f"\nFeatures HP por BSP en t={thresholds[best_t]:.1f}:")
        print(f" Media: {hp_per_b[best_t].float().mean():.1f}")
        print(f" Máx: {hp_per_b[best_t].max().item()}")
        print(f" BSPs sin features: {(hp_per_b[best_t] == 0).sum().item()}")

    return hp_mask_TFB, active_feature_indices

print(" Fase 1 rediseñada: umbral de significancia (10) + estructura (T, F, B)")

```

Fase 1 rediseñada: umbral de significancia (10) + estructura (T, F, B)

```

THRESHOLDS = torch.arange(0.0, 1.05, 0.05).tolist()

# f_max calculado UNA SOLA VEZ aquí - se pasa a Fase 1 y Fase 2
f_max = sae_features_train.max(dim=0)[0] # (n_features,)
print(f"f_max calculado: {(f_max > 0).sum().item()} features activas de {f_max.shape[0]}")

start_time = time.time()

hp_mask_TFB, active_feature_indices = identify_high_precision_features_gpu(
    sae_features_train,
    bsp_gt_train_pieces,
    f_max,
    precision_threshold=0.95,
    significance_threshold=1,
    thresholds=THRESHOLDS,
    feature_batch_size=512,
    verbose=True
)

phase1_time = time.time() - start_time

print(f"\nTiempo Fase 1: {phase1_time:.2f} seg")
print(f"hp_mask_TFB shape: {hp_mask_TFB.shape} (T, n_active, n_bsps)")

```

```

f_max calculado: 6942 features activas de 8192
=====
FASE 1: IDENTIFICAR FEATURES DE ALTA PRECISIÓN
=====
Precision threshold:      0.95
Significance threshold:   >= 1 activaciones
Features activas:        6942/8192 (84.7%)

```

```

BSPs: 128
Thresholds: [0.0, 0.05000000074505806, 0.10000000149011612, 0.15000000596046448, 0.20000000298023224, 0.25000000799570312, 0.30000001594031525, 0.35000002389582736, 0.40000003185133948, 0.4500000398068516, 0.5000000477623637, 0.5500000557178758, 0.6000000636733879, 0.6500000716289001, 0.7000000795844122, 0.7500000875399243, 0.8000000954954364, 0.8500001034509485, 0.9000001114064606, 0.9500001193619727, 1.0000001273174848]
=====

```

```
Fase 1: 100%|          | 14/14 [00:01<00:00, 10.43it/s]
```

Pares (feature, BSP) de alta precisión por threshold:

```

t=0.0: 16233 pares
t=0.1: 16233 pares
t=0.1: 16423 pares
t=0.2: 17103 pares
t=0.2: 18197 pares
t=0.2: 19763 pares
t=0.3: 21604 pares
t=0.3: 23886 pares
t=0.4: 26636 pares
t=0.5: 29634 pares
t=0.5: 32836 pares
t=0.6: 36034 pares
t=0.6: 39910 pares
t=0.6: 44526 pares
t=0.7: 49080 pares
t=0.8: 54802 pares
t=0.8: 61520 pares
t=0.9: 70718 pares
t=0.9: 84658 pares
t=0.9: 104660 pares ← más pares
t=1.0: 0 pares

```

Features HP por BSP en t=0.9:

```

Media: 817.7
Máx: 3823
BSPs sin features: 0

```

Tiempo Fase 1: 33.26 seg

hp_mask_TFB shape: torch.Size([21, 6942, 128]) (T, n_active, n_bsps)

0.7 7. Fase 2: Reconstruir Tableros en Test Set (GPU)

```

def reconstruct_boards_gpu(
    sae_features_test,
    bsp_gt_test,
    hp_mask_TFB,
    active_feature_indices,
    f_max,
    thresholds=[0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9],
    verbose=True
):
    """
    Fase 2: barrido simultáneo de thresholds T → toma el mejor F1.

    Para cada threshold t:
    1. Binarizar features de test: active[p, f] = acts[p, f] > t * f_max[f]
    2. Predicción por (posición, BSP):
        predict[p, b] = OR_f ( active[p,f] AND hp_mask[t, f, b] )
        implementado como: (active_PF @ hp_FB) > 0
    3. F1 global micro-averaged (acumula TP/FP/FN sobre todas las posiciones y BSPs)

    Retorna el max F1 sobre todos los thresholds (Eq. 5 del paper: max_t).
    """

```

```

Args:
    sae_features_test:      Tensor (P, n_features) GPU
    bsp_gt_test:           Tensor (P, n_bsps) GPU, bool
    hp_mask_TFB:           BoolTensor (T, n_active, n_bsps) de Fase 1
    active_feature_indices: LongTensor (n_active,) de Fase 1
    f_max:                 Tensor (n_features,) GPU - MISMO objeto que Fase 1
    thresholds:            lista de fracciones (debe coincidir con Fase 1)

Returns:
    best_f1:               float, max F1 sobre thresholds
    best_threshold:        float, threshold óptimo
    f1_per_threshold:      np.array (T,), F1 por threshold
"""
device = sae_features_test.device
n_thresholds = len(thresholds)
thresholds_T = torch.tensor(thresholds, device=device)

# Activaciones del test solo para features activas
test_active_PF = sae_features_test[:, active_feature_indices] # (P, n_active)
active_f_max = f_max[active_feature_indices] # (n_active,)

gt_PB = bsp_gt_test.bool() # (P, B)
eps = 1e-8

f1_scores = torch.zeros(n_thresholds, device=device)
precision_T = torch.zeros(n_thresholds, device=device)
recall_T = torch.zeros(n_thresholds, device=device)

if verbose:
    print("=" * 60)
    print("FASE 2: RECONSTRUCCIÓN - BARRIDO DE THRESHOLDS")
    print("=" * 60)
    print(f"{'t':>5} {'F1':>8} {'P':>8} {'R':>8} {'TP':>8} {'FP':>8} {'FN':>8}")
    print("-" * 60)

for t_idx in range(n_thresholds):
    t = thresholds[t_idx]

    # Threshold absoluto por feature: t * f_max[f]
    abs_thresh_F = thresholds_T[t_idx] * active_f_max # (n_active,)

    # Binarizar features en test
    active_PF = test_active_PF > abs_thresh_F.unsqueeze(0) # (P, n_active)

    # predict[p, b] = any( active[p,f] AND hp[t,f,b] )
    # = (active_PF.float() @ hp_mask[t].float()) > 0
    hp_FB = hp_mask_TFB[t_idx].float() # (n_active, B)
    predictions_PB = (active_PF.float() @ hp_FB) > 0 # (P, B)

    # Métricas globales micro-averaged
    tp = (predictions_PB & gt_PB).sum().float()
    fp = (predictions_PB & ~gt_PB).sum().float()
    fn = (~predictions_PB & gt_PB).sum().float()

    p = tp / (tp + fp + eps)
    r = tp / (tp + fn + eps)
    f1 = 2 * p * r / (p + r + eps)

    f1_scores[t_idx] = f1
    precision_T[t_idx] = p
    recall_T[t_idx] = r

```

```

        if verbose:
            print(f"{t:>5.1f}  {f1.item():>8.4f}  {p.item():>8.4f}  {r.item():>8.4f}"
                  f"  {tp.item():>8.0f}  {fp.item():>8.0f}  {fn.item():>8.0f}")

    best_idx      = f1_scores.argmax().item()
    best_f1       = f1_scores[best_idx].item()
    best_threshold = thresholds[best_idx]

    if verbose:
        print("=" * 60)
        print(f"Mejor threshold: t={best_threshold:.1f}")
        print(f"Reconstruction Score (max_t F1): {best_f1:.4f}")
        print("=" * 60)

    return best_f1, best_threshold, f1_scores.cpu().numpy()

print(" Fase 2 rediseñada: barrido T simultáneo + F1 global micro-averaged + max_t")

```

Fase 2 rediseñada: barrido T simultáneo + F1 global micro-averaged + max_t

```

start_time = time.time()

reconstruction_score, best_threshold, f1_per_threshold = reconstruct_boards_gpu(
    sae_features_test,
    bsp_gt_test_pieces,
    hp_mask_TFB,
    active_feature_indices,
    f_max,          # mismo tensor calculado antes de Fase 1
    thresholds=THRESHOLDS,
    verbose=True
)

phase2_time = time.time() - start_time
total_time  = phase1_time + phase2_time

print(f"\nTiempo Fase 2: {phase2_time:.2f} seg")
print(f"Tiempo total:  {total_time:.2f} seg ({total_time/60:.2f} min)")

```

=====

FASE 2: RECONSTRUCCIÓN - BARRIDO DE THRESHOLDS

=====

t	F1	P	R	TP	FP	FN
0.0	0.0426	0.7040	0.0220	44104	18548	1961896
0.1	0.0426	0.7040	0.0220	44104	18548	1961896
0.1	0.0445	0.7116	0.0230	46087	18678	1959913
0.2	0.0487	0.7257	0.0252	50529	19099	1955471
0.2	0.0536	0.7405	0.0278	55762	19540	1950238
0.2	0.0579	0.7475	0.0301	60366	20390	1945634
0.3	0.0607	0.7503	0.0316	63487	21129	1942513
0.3	0.0634	0.7496	0.0331	66379	22171	1939621
0.4	0.0655	0.7438	0.0342	68687	23659	1937313
0.5	0.0671	0.7367	0.0351	70484	25191	1935516
0.5	0.0679	0.7263	0.0356	71392	26903	1934608
0.6	0.0681	0.7166	0.0357	71710	28366	1934290
0.6	0.0680	0.7025	0.0357	71680	30359	1934320
0.6	0.0682	0.6881	0.0359	71997	32628	1934003
0.7	0.0680	0.6691	0.0358	71861	35544	1934139
0.8	0.0683	0.6509	0.0360	72252	38758	1933748
0.8	0.0686	0.6321	0.0363	72800	42367	1933200
0.9	0.0702	0.6087	0.0372	74694	48024	1931306

0.9	0.0732	0.5805	0.0391	78398	56665	1927602
0.9	0.0752	0.5512	0.0404	80968	65938	1925032
1.0	0.0000	0.0000	0.0000	0	0	2006000

```
=====
Mejor threshold: t=0.9
Reconstruction Score (max_t F1): 0.0752
=====
```

```
Tiempo Fase 2: 0.90 seg
Tiempo total: 34.16 seg (0.57 min)
```

0.8 8. Análisis de Resultados

```
print("=" * 60)
print("BARRIDO DE THRESHOLDS - F1 POR THRESHOLD")
print("=" * 60)
for i, (t, f1) in enumerate(zip(THRESHOLDS, f1_per_threshold)):
    marker = " ← MEJOR" if i == f1_per_threshold.argmax() else ""
    print(f"t={t:.1f}: F1={f1:.4f}{marker}")
print("=" * 60)

plt.figure(figsize=(10, 6))
plt.plot(THRESHOLDS, f1_per_threshold, marker='o', linewidth=2, markersize=8)
plt.axhline(y=reconstruction_score, color='r', linestyle='--',
            label=f'Best F1={reconstruction_score:.4f} (t={best_threshold:.1f})')
plt.axhline(y=0.95, color='g', linestyle=':', label='Target Othello SAE (paper)=0.95')
plt.xlabel('Threshold t', fontsize=12)
plt.ylabel('F1 Score (micro-averaged global)', fontsize=12)
plt.title('Board Reconstruction: F1 por Threshold', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()


print(f"\n Threshold óptimo: t={best_threshold:.1f}")
print(f" Reconstruction Score: {reconstruction_score:.4f}")
```

```
=====
BARRIDO DE THRESHOLDS - F1 POR THRESHOLD
=====
```

```
t=0.0: F1=0.0426
t=0.1: F1=0.0426
t=0.1: F1=0.0445
t=0.2: F1=0.0487
t=0.2: F1=0.0536
t=0.2: F1=0.0579
t=0.3: F1=0.0607
t=0.3: F1=0.0634
t=0.4: F1=0.0655
t=0.5: F1=0.0671
t=0.5: F1=0.0679
t=0.6: F1=0.0681
t=0.6: F1=0.0680
t=0.6: F1=0.0682
t=0.7: F1=0.0680
t=0.8: F1=0.0683
t=0.8: F1=0.0686
t=0.9: F1=0.0702
t=0.9: F1=0.0732
t=0.9: F1=0.0752 ← MEJOR
```

t=1.0: F1=0.0000

=====



Threshold óptimo: t=0.9

Reconstruction Score: 0.0752

```
# Estadísticas sobre features de alta precisión por BSP
hp_per_bsp = hp_mask_TFB.sum(dim=1) # (T, n_bsps)
best_t_idx = torch.tensor(f1_per_threshold).argmax().item()
n_features_per_bsp = hp_per_bsp[best_t_idx].cpu().numpy()

print("="*60)
print("ESTADÍSTICAS - FEATURES POR BSP")
print("="*60)
print(f"Total BSPs: {len(n_features_per_bsp)}")
print(f"Features promedio por BSP: {np.mean(n_features_per_bsp):.1f}")
print(f"Features mediana por BSP: {np.median(n_features_per_bsp):.1f}")
print(f"BSPs con 0 features: {np.sum(n_features_per_bsp == 0)}")
print(f"BSPs con 1+ features: {np.sum(n_features_per_bsp > 0)}")
print(f"BSPs con 5+ features: {np.sum(n_features_per_bsp >= 5)}")
print(f"BSPs con 10+ features: {np.sum(n_features_per_bsp >= 10)}")
print(f"Max features en una BSP: {np.max(n_features_per_bsp)}")
print("="*60)
```

=====

ESTADÍSTICAS - FEATURES POR BSP

=====

Total BSPs: 128
Features promedio por BSP: 817.7
Features mediana por BSP: 547.5
BSPs con 0 features: 0
BSPs con 1+ features: 128

BSPs con 5+ features: 128
 BSPs con 10+ features: 128
 Max features en una BSP: 3823

0.9 9. Guardar Resultados

```
results = {
    'reconstruction_score': reconstruction_score,
    'best_threshold': best_threshold,
    'f1_per_threshold': f1_per_threshold,
    'thresholds': np.array(THRESHOLDS),
    'phase1_time_seconds': phase1_time,
    'phase2_time_seconds': phase2_time,
    'total_time_seconds': total_time,
    'bsp_names': bsp_pieces_names,
}

metrics_dir = get_metrics_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)
output_path = metrics_dir / f"reconstruction_results_{filter_suffix}.npz"
np.savez(output_path, **results)

print(f" Resultados guardados en: {output_path}")
print(f"\nContenido:")
print(f" reconstruction_score: {reconstruction_score:.4f}")
print(f" best_threshold: {best_threshold:.1f}")
print(f" f1_per_threshold: {f1_per_threshold}")

# Comparación con el paper
print()
print("=" * 60)
print("COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)")
print("=" * 60)
print(f"{'Modelo':<25} {'Reconstruction':>14}")
print("-" * 40)
print(f"{'SAE trained (paper)':<25} {'0.95':>14} ← Objetivo Othello")
print(f"{'Nuestro SAE':<25} {reconstruction_score:>14.4f} ← Resultado")
print("=" * 60)

if reconstruction_score >= 0.90:
    print("\n Excelente: Muy cercano al objetivo del paper")
elif reconstruction_score >= 0.50:
    print("\n~ Moderado: Por encima de random, revisar SAE")
else:
    print("\n Bajo: Revisar entrenamiento del SAE o calidad de las activaciones")
```

Resultados guardados en: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_topk\sae_batch-topk_topk\25000games\ex5\metrics\reconstruction_results_absolute.npz

Contenido:

```
reconstruction_score: 0.0752
best_threshold: 0.9
f1_per_threshold: [0.04264032 0.04264032 0.04451206 0.04868792 0.05358377 0.05785631 0.06073521 0.06338259 0.06541206 0.068541206 0.07167032 0.07479968 0.07792896 0.08105824 0.08418752 0.0873168 0.09044608 0.09357536 0.09670464 0.09983392 0.1029632 0.10609248 0.10922176 0.11235104 0.11548032 0.1186096 0.12173888 0.12486816 0.12799744 0.13112672 0.134256 0.13738528 0.14051456 0.14364384 0.14677312 0.1499024 0.15303168 0.15616096 0.15929024 0.16241952 0.1655488 0.16867808 0.17180736 0.17493664 0.17806592 0.1811952 0.18432448 0.18745376 0.19058304 0.19371232 0.1968416 0.19997088 0.20310016 0.20622944 0.20935872 0.212488 0.21561728 0.21874656 0.22187584 0.22500512 0.2281344 0.23126368 0.23439296 0.23752224 0.24065152 0.2437808 0.24691008 0.25003936 0.25316864 0.25629792 0.2594272 0.26255648 0.26568576 0.26881504 0.27194432 0.2750736 0.27820288 0.28133216 0.28446144 0.28759072 0.29072 0.29384928 0.29697856 0.30010784 0.30323712 0.3063664 0.30949568 0.31262496 0.31575424 0.31888352 0.3220128 0.32514208 0.32827136 0.33140064 0.33452992 0.3376592 0.34078848 0.34391776 0.34704704 0.35017632 0.3533056 0.35643488 0.35956416 0.36269344 0.36582272 0.368952 0.37208128 0.37521056 0.37833984 0.38146912 0.3845984 0.38772768 0.39085696 0.39398624 0.39711552 0.4002448 0.40337408 0.40650336 0.40963264 0.41276192 0.4158912 0.41902048 0.42214976 0.42527904 0.42840832 0.4315376 0.43466688 0.43779616 0.44092544 0.44405472 0.447184 0.45031328 0.45344256 0.45657184 0.45970112 0.4628304 0.46595968 0.46908896 0.47221824 0.47534752 0.4784768 0.48160608 0.48473536 0.48786464 0.49099392 0.4941232 0.49725248 0.50038176 0.50351104 0.50664032 0.5097696 0.51289888 0.51602816 0.51915744 0.52228672 0.525416 0.52854528 0.53167456 0.53480384 0.53793312 0.5410624 0.54419168 0.54732096 0.55045024 0.55357952 0.5567088 0.55983808 0.56296736 0.56609664 0.56922592 0.5723552 0.57548448 0.57861376 0.58174304 0.58487232 0.5880016 0.59113088 0.59426016 0.59738944 0.60051872 0.603648 0.60677728 0.60990656 0.61303584 0.61616512 0.6192944 0.62242368 0.62555296 0.62868224 0.63181152 0.6349408 0.63807008 0.64119936 0.64432864 0.64745792 0.6505872 0.65371648 0.65684576 0.65997504 0.66310432 0.6662336 0.66936288 0.67249216 0.67562144 0.67875072 0.68188 0.68500928 0.68813856 0.69126784 0.69439712 0.6975264 0.70065568 0.70378496 0.70691424 0.71004352 0.7131728 0.71630208 0.71943136 0.72256064 0.72568992 0.7288192 0.73194848 0.73507776 0.73820704 0.74133632 0.7444656 0.74759488 0.75072416 0.75385344 0.75698272 0.760112 0.76324128 0.76637056 0.7695 0.77262928 0.77575856 0.77888784 0.78201712 0.7851464 0.78827568 0.79140496 0.79453424 0.79766352 0.8007928 0.80392208 0.80705136 0.81018064 0.81330992 0.8164392 0.81956848 0.82269776 0.82582704 0.82895632 0.8320856 0.83521488 0.83834416 0.84147344 0.84460272 0.847732 0.85086128 0.85399056 0.85711984 0.86024912 0.8633784 0.86650768 0.86963696 0.87276624 0.87589552 0.8790248 0.88215408 0.88528336 0.88841264 0.89154192 0.8946712 0.89780048 0.90092976 0.90405904 0.90718832 0.9103176 0.91344688 0.91657616 0.91970544 0.92283472 0.925964 0.92909328 0.93222256 0.93535184 0.93848112 0.9416104 0.94473968 0.94786896 0.95099824 0.95412752 0.9572568 0.96038608 0.96351536 0.96664464 0.96977392 0.9729032 0.97603248 0.97916176 0.98229104 0.98542032 0.9885496 0.99167888 0.99480816 0.99793744 1.00106672 1.004196 1.00732528 1.01045456 1.01358384 1.01671312 1.0198424 1.02297168 1.02610096 1.02923024 1.03235952 1.0354888 1.03861808 1.04174736 1.04487664 1.04800592 1.0511352 1.05426448 1.05739376 1.06052304 1.06365232 1.0667816 1.06991088 1.07304016 1.07616944 1.07929872 1.082428 1.08555728 1.08868656 1.09181584 1.09494512 1.0980744 1.10120368 1.10433296 1.10746224 1.11059152 1.1137208 1.11685008 1.11997936 1.12310864 1.12623792 1.1293672 1.13249648 1.13562576 1.13875504 1.14188432 1.1450136 1.14814288 1.15127216 1.15440144 1.15753072 1.16066 1.16378928 1.16691856 1.17004784 1.17317712 1.1763064 1.17943568 1.18256496 1.18569424 1.18882352 1.1919528 1.19508208 1.19821136 1.20134064 1.20446992 1.2075992 1.21072848 1.21385776 1.21698704 1.22011632 1.2232456 1.22637488 1.22950416 1.23263344 1.23576272 1.238892 1.24202128 1.24515056 1.24827984 1.25140912 1.2545384 1.25766768 1.26079696 1.26392624 1.26705552 1.2701848 1.27331408 1.27644336 1.27957264 1.28270192 1.2858312 1.28896048 1.29208976 1.29521904 1.29834832 1.3014776 1.30460688 1.30773616 1.31086544 1.31399472 1.317124 1.32025328 1.32338256 1.32651184 1.32964112 1.3327704 1.33589968 1.33902896 1.34215824 1.34528752 1.3484168 1.35154608 1.35467536 1.35780464 1.36093392 1.3640632 1.36719248 1.37032176 1.37345104 1.37658032 1.3797096 1.38283888 1.38596816 1.38909744 1.39222672 1.395356 1.39848528 1.40161456 1.40474384 1.40787312 1.4110024 1.41413168 1.41726096 1.42039024 1.42351952 1.4266488 1.42977808 1.43290736 1.43603664 1.43916592 1.4422952 1.44542448 1.44855376 1.45168304 1.45481232 1.4579416 1.46107088 1.4642 1.46732928 1.47045856 1.47358784 1.47671712 1.4798464 1.48297568 1.48610496 1.48923424 1.49236352 1.4954928 1.49862208 1.50175136 1.50488064 1.50800992 1.5111392 1.51426848 1.51739776 1.52052704 1.52365632 1.5267856 1.52991488 1.53304416 1.53617344 1.53930272 1.542432 1.54556128 1.54869056 1.55181984 1.55494912 1.5580784 1.56120768 1.56433696 1.56746624 1.57059552 1.5737248 1.57685408 1.57998336 1.58311264 1.58624192 1.5893712 1.59250048 1.59562976 1.59875904 1.60188832 1.6050176 1.60814688 1.61127616 1.61440544 1.61753472 1.620664 1.62379328 1.62692256 1.63005184 1.63318112 1.6363104 1.63943968 1.64256896 1.64569824 1.64882752 1.6519568 1.65508608 1.65821536 1.66134464 1.66447392 1.6676032 1.67073248 1.67386176 1.67699104 1.68012032 1.6832496 1.68637888 1.68950816 1.69263744 1.69576672 1.698896 1.70202528 1.70515456 1.70828384 1.71141312 1.7145424 1.71767168 1.72080096 1.72393024 1.72705952 1.7301888 1.73331808 1.73644736 1.73957664 1.74270592 1.7458352 1.74896448 1.75209376 1.75522304 1.75835232 1.7614816 1.76461088 1.76774016 1.77086944 1.774000 1.77712928 1.78025856 1.78338784 1.78651712 1.7896464 1.79277568 1.79590496 1.79903424 1.80216352 1.8052928 1.80842208 1.81155136 1.81468064 1.81780992 1.8209392 1.82406848 1.82719776 1.83032704 1.83345632 1.8365856 1.83971488 1.84284416 1.84597344 1.84910272 1.852232 1.85536128 1.85849056 1.86161984 1.86474912 1.8678784 1.87100768 1.87413696 1.87726624 1.88039552 1.8835248 1.88665408 1.88978336 1.89291264 1.89604192 1.8991712 1.90230048 1.90542976 1.90855904 1.91168832 1.9148176 1.91794688 1.92107616 1.92420544 1.92733472 1.930464 1.93359328 1.93672256 1.93985184 1.94298112 1.9461104 1.94923968 1.95236896 1.95549824 1.95862752 1.9617568 1.96488608 1.96801536 1.97114464 1.97427392 1.9774032 1.98053248 1.98366176 1.98679104 1.98992032 1.9930496 1.99617888 1.99930816 2.00243744 2.00556672 2.008696 2.01182528 2.01495456 2.01808384 2.02121312 2.0243424 2.02747168 2.03060096 2.03373024 2.03685952 2.040000 2.04312928 2.04625856 2.04938784 2.05251712 2.0556464 2.05877568 2.06190496 2.06503424 2.06816352 2.0712928 2.07442208 2.07755136 2.08068064 2.08380992 2.0869392 2.09006848 2.09319776 2.09632704 2.09945632 2.1025856 2.10571488 2.10884416 2.11197344 2.11510272 2.118232 2.12136128 2.12449056 2.12761984 2.13074912 2.1338784 2.13700768 2.14013696 2.14326624 2.14639552 2.1495248 2.15265408 2.15578336 2.15891264 2.16204192 2.1651712 2.16830048 2.17142976 2.17455904 2.17768832 2.1808176 2.18394688 2.18707616 2.19020544 2.19333472 2.196464 2.19959328 2.20272256 2.20585184 2.20898112 2.2121104 2.21523968 2.21836896 2.22149824 2.22462752 2.2277568 2.23088608 2.23401536 2.23714464 2.24027392 2.2434032 2.24653248 2.24966176 2.25279104 2.25592032 2.2590496 2.26217888 2.26530816 2.26843744 2.27156672 2.274696 2.27782528 2.28095456 2.28408384 2.28721312 2.2903424 2.29347168 2.29660096 2.29973024 2.30285952 2.3059888 2.30911808 2.31224736 2.31537664 2.31850592 2.3216352 2.32476448 2.32789376 2.33102304 2.33415232 2.3372816 2.34041088 2.34354016 2.34666944 2.34979872 2.352928 2.35605728 2.35918656 2.36231584 2.36544512 2.3685744 2.37170368 2.37483296 2.37796224 2.38109152 2.3842208 2.38735008 2.39047936 2.39360864 2.39673792 2.3998672 2.40299648 2.40612576 2.40925504 2.41238432 2.4155136 2.41864288 2.42177216 2.42490144 2.42803072 2.43116 2.43428928 2.43741856 2.44054784 2.44367712 2.4468064 2.44993568 2.45306496 2.45619424 2.45932352 2.4624528 2.46558208 2.46871136 2.47184064 2.47496992 2.4780992 2.48122848 2.48435776 2.48748704 2.49061632 2.4937456 2.49687488 2.49990416 2.50303344 2.50616272 2.509292 2.51242128 2.51555056 2.51867984 2.52180912 2.5249384 2.52806768 2.53119696 2.53432624 2.53745552 2.5405848 2.54371408 2.54684336 2.54997264 2.55310192 2.5562312 2.55936048 2.56248976 2.56561904 2.56874832 2.5718776 2.57500688 2.57813616 2.58126544 2.58439472 2.587524 2.59065328 2.59378256 2.59691184 2.59904112 2.6021704 2.60530008 2.60842936 2.61155864 2.61468792 2.6178172 2.62094648 2.62407576 2.62720504 2.63033432 2.6334636 2.63659288 2.63972216 2.64285144 2.64598072 2.64911008 2.65223936 2.65536864 2.65849792 2.6616272 2.66475648 2.66788576 2.67101504 2.67414432 2.6772736 2.68040288 2.68353216 2.68666144 2.68979072 2.69292008 2.69604936 2.69917864 2.70230792 2.7054372 2.70856648 2.71169576 2.71482504 2.71795432 2.7210836 2.72421288 2.72734216 2.73047144 2.73360072 2.73673008 2.73985936 2.74298864 2.74611792 2.7492472 2.75237648 2.75550576 2.75863504 2.76176432 2.7648936 2.76802288 2.77115216 2.77428144 2.77741072 2.78054008 2.78366936 2.78679864 2.78992792 2.7930572 2.79618648 2.79931576 2.80244504 2.80557432 2.8087036 2.81183288 2.81496216 2.81809144 2.82122072 2.82435008 2.82747936 2.83060864 2.83373792 2.8368672 2.840000 2.84312928 2.84625856 2.84938784 2.85251712 2.8556464 2.85877568 2.86190496 2.86503424 2.86816352 2.8712928 2.87442208 2.87755136 2.88068064 2.88380992 2.8869392 2.89006848 2.89319776 2.89632704 2.89945632 2.9025856 2.90571488 2.90884416 2.91197344 2.91510272 2.918232 2.92136128 2.92449056 2.92761984 2.93074912 2.9338784 2.93700768 2.94013696 2.94326624 2.94639552 2.9495248 2.95265408 2.95578336 2.95891264 2.96204192 2.9651712 2.96830048 2.97142976 2.97455904 2.97768832 2.9808176 2.98394688 2.98707616 2.99020544 2.99333472 2.996464 2.99959328 3.00272256 3.00585184 3.00898112 3.0121104 3.01523968 3.01836896 3.02149824 3.02462752 3.0277568 3.03088608 3.03401536 3.03714464 3.04027392 3.0434032 3.04653248 3.04966176 3.05279104 3.05592032 3.0590496 3.06217888 3.06530816 3.06843744 3.07156672 3.074696 3.07782528 3.08095456 3.08408384 3.08721312 3.0903424 3.09347168 3.09660096 3.09973024 3.10285952 3.1059888 3.10911808 3.11224736 3.11537664 3.11850592 3.1216352 3.12476448 3.12789376 3.13102304 3.13415232 3.1372816 3.14041088 3.14354016 3.14666944 3.14979872 3.152928 3.15605728 3.15918656 3.16231584 3.16544512 3.1685744 3.17170368 3.17483296 3.17796224 3.18109152 3.1842208 3.18735008 3.19047936 3.19360864 3.19673792 3.1998672 3.20299648 3.20612576 3.20925504 3.21238432 3.2155136 3.21864288 3.22177216 3.22490144 3.22803072 3.23116 3.23428928 3.23741856 3.24054784 3.24367712 3.2468064 3.24993568 3.25306496 3.25619424 3.25932352 3.2624528 3.26558208 3.26871136 3.27184064 3.27496992 3.2780992 3.28122848 3.28435776 3.28748704 3.29061632 3.2937456 3.29687488 3.29990416 3.30303344 3.30616272 3.309292 3.31242128 3.31555056 3.31867984 3.32180912 3.3249384 3.32806768 3.33119696 3.33432624 3.33745552 3.3405848 3.34371408 3.34684336 3.34997264 3.35310192 3.3562312 3.35936048 3.36248976 3.36561904 3.36874832 3.3718776 3.37500688 3.37813616 3.38126544 3.38439472 3.387524 3.39065328 3
```

```
0.06707413 0.06785361 0.0680982 0.06800633 0.06822339 0.06800495 0.06825854 0.06864146 0.07017744
0.07323278 0.0752174 0.          ]
```

```
=====
COMPARACIÓN CON PAPER (Karvonen et al., NeurIPS 2024)
=====
Modelo                Reconstruction
-----
SAE trained (paper)           0.95  ← Objetivo Othello
Nuestro SAE                   0.0752 ← Resultado
=====
```

Bajo: Revisar entrenamiento del SAE o calidad de las activaciones

0.10 10. Generar PDF con Quarto

```
output_dir = get_metrics_dir(
    NAMING_META['experiment_base_path'],
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    extras_id=extras_id
)

pdf_name = get_report_name(
    NAMING_META['variante'],
    NAMING_META['tecnica'],
    NAMING_META['capa'],
    NAMING_META['juegos'],
    f'reconstruction_{filter_suffix}',
    extras_id=extras_id
)

notebook_name = 'reconstruction_sae_batchtopk.ipynb'
output_path = output_dir / pdf_name

print(f' Generando PDF con Quarto...')
print(f' Archivo de salida: {output_path}')
```

```
result = subprocess.run(
    f'quarto render {notebook_name} --to pdf --output-dir "{output_dir}" --output {pdf_name}',
    shell=True,
    capture_output=True,
    text=True
)

if result.returncode == 0:
    print(f' PDF generado exitosamente: {output_path}')
```

```
else:
    print(f' Error al generar PDF:')
    print(result.stderr)
```

```
Generando PDF con Quarto...
Archivo de salida: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_topk\sae_batch-
topk_topk\25000games\ex5\metrics\sae_topk_batch-topk_l6_25000g_ex5_reconstruction_absolute.pdf
PDF generado exitosamente: c:\Users\Esposa\Documents\Repos\othello_world\sae\experiments\layer_06\sae_topk\sae_batch-
topk_topk\25000games\ex5\metrics\sae_topk_batch-topk_l6_25000g_ex5_reconstruction_absolute.pdf
```