

Learning Multi-Level Features with Matryoshka Sparse Autoencoders

Bart Bussmann* Noa Nabeshima* Adam Karvonen Neel Nanda

Abstract

Sparse autoencoders (SAEs) have emerged as a powerful tool for interpreting neural networks by extracting the concepts represented in their activations. However, choosing the size of the SAE dictionary (i.e. number of learned concepts) creates a tension: as dictionary size increases to capture more relevant concepts, sparsity incentivizes features to be split or absorbed into more specific features, leaving high-level features missing or warped. We introduce Matryoshka SAEs, a novel variant that addresses these issues by simultaneously training multiple nested dictionaries of increasing size, forcing the smaller dictionaries to independently reconstruct the inputs without using the larger dictionaries. This organizes features hierarchically - the smaller dictionaries learn general concepts, while the larger dictionaries learn more specific concepts, without incentive to absorb the high-level features. We train Matryoshka SAEs on Gemma-2-2B and TinyStories and find superior performance on sparse probing and targeted concept erasure tasks, more disentangled concept representations, and reduced feature absorption. While there is a minor trade-off with reconstruction performance, we believe Matryoshka SAEs are a superior alternative for practical tasks, as they enable training arbitrarily large SAEs while retaining interpretable features at different levels of abstraction.

practice, however, sparsity is only an imperfect proxy for interpretability; optimizing to represent inputs using as few active latents as possible can lead to pathological solutions.

A core issue lies in the mismatch between the flat sparsity objective and the inherent hierarchical structure of real-world features. For example, a broad, general concept (e.g. “punctuation marks”) might be a parent feature to more specific ones (“punctuation mark and period”, “punctuation mark and question mark,” and “punctuation mark and comma”). When an SAE has the capacity to learn many latents, the sparsity penalty can push it to replace each general concept with a set of narrowly specialized features without retaining the original high-level category—a phenomenon called *feature splitting* (Bricken et al., 2023). Even more problematic is *feature absorption* (Chanin et al., 2024), where a parent feature only partially splits and a number of specific instances of a general feature are absorbed by more specialized latents, leaving “holes” in the representation of the more general feature (e.g. a latent that activates on all tokens starting with an E, except if the word is Elephant). Finally, *feature composition* occurs when the SAE merges distinct concepts (like “red” and “triangle”) into single composite features (“red triangle”) to minimize the number of active latents, instead of learning the underlying features (Anders et al., 2024; Wattenberg & Viégas, 2024; Leask et al., 2025).

These problems become more severe as SAEs scale to larger dictionary sizes (i.e., more learned concepts), since the model exploits the additional capacity to further minimize active latents through mechanisms like feature splitting (Karvonen et al., 2024a). This presents a stark contrast to typical deep learning, where scaling consistently improves both training loss and model capabilities (Kaplan et al., 2020; Hoffmann et al., 2022). While larger SAEs do achieve lower training loss (Gao et al., 2024; Templeton, 2024), their performance on downstream tasks often deteriorates (Karvonen et al., 2024a). This fundamental tension creates significant practical challenges for deploying SAEs in interpretability research - they become less reliable for probing model behavior, steering models, analyzing feature circuits, or understanding how models process information. The current recommended approach is to train a sweep of different sizes and evaluate each of them, but this is costly (Lieberum et al., 2024). Ideally, we would like to have a single SAE that can capture both high-level and fine-grained features.

1. Introduction

Sparse autoencoders (SAEs) have emerged as a powerful tool for interpreting the internal representations of neural networks (Bricken et al., 2023; Cunningham et al., 2023). By training autoencoders to map dense activations of neural networks into a *sparse* latent space, researchers aim to extract meaningful, human-interpretable features that provide insight into a model’s decision-making process. In

*Joint first authors. Author order determined by coinflip. Correspondence to: Bart Bussmann <bartbussmann@gmail.com>.




Figure 1. Architecture and Performance of Matryoshka SAEs. (Left) The model learns multiple nested reconstructions simultaneously, with each reconstruction using only a subset of the total latents. This creates pressure for early latents to capture general features while later latents can specialize in more specific concepts. (Right) Comparative metrics between SAEs with average sparsity of 40 active latents per token ($L_0=40$) showing that while Matryoshka SAEs sacrifice a small amount of reconstruction fidelity (higher variance unexplained), they achieve significantly lower feature absorption rates and less feature composition (lower decoder cosine similarity).

Therefore, we introduce **Matryoshka SAEs**, a novel hierarchical approach to SAE training, named after the Russian nesting dolls and inspired by Kusupati et al. (2024). Matryoshka SAEs learn a single feature space that retains both abstract and specific features by training multiple nested SAEs of increasing size simultaneously. Each nested subSAE is optimized to reconstruct the input using only a subset of the total latents (including the sub-SAEs nested within itself). This nested structure prevents later, more specialized latents from absorbing the roles of earlier, more general ones, effectively regularizing the SAE to learn features at multiple levels of abstraction. Despite the need for reconstructions at multiple scales, Matryoshka SAEs can be implemented efficiently with a modest increase in training time compared to traditional SAEs, depending on the number of nested sub-SAEs.

Through extensive experiments on both synthetic and real-world datasets, we demonstrate that Matryoshka SAEs mitigate feature absorption across a wide range of model sizes and sparsity levels. We find that Matryoshka SAEs have more disentangled latent representations (as measured by maximum cosine similarity of decoder vectors) and also improve performance on probing and targeted concept erasure tasks compared to standard SAE baselines. Furthermore, we find that with increasing dictionary size, Matryoshka SAEs often improve or maintain their performance on downstream tasks, where other architectures deteriorate in performance. While Matryoshka SAEs exhibit a small trade-off in reconstruction error due to the additional constraints, we argue that the improved quality of the learned latents, which also get better with scaling, make them an attractive alternative for practical applications.

We consider our results a positive sign for the use of SAEs in interpretability research. We find that concerning failure modes that were thought to be fundamental barriers to using SAEs, can be mitigated through simple changes in the

training procedure. This indicates that other problems with SAEs might be easier to solve than thought, and illustrates the promise of more research in this area.

2. Background

2.1. Interpretability with Sparse Autoencoders

Early mechanistic interpretability work sought to understand neural networks by examining individual architectural components, but faced challenges due to polysemanticity - where components seem to exhibit multiple unrelated behaviors (Olah et al., 2020; Elhage et al., 2022). In recent years, SAEs emerged as a promising solution for extracting monosemantic and interpretable features from model activations (Bricken et al., 2023; Cunningham et al., 2023; Templeton, 2024; Gao et al., 2024; Rajamanoharan et al., 2024a;b), which researchers can analyze through manual inspection (Bricken et al., 2023; Lin, 2023) or automated interpretability techniques (Gao et al., 2024; Juang et al., 2024). Their effectiveness has been demonstrated across various applications, such as analyzing attention mechanisms in GPT-2 (Kissane et al., 2024), circuit analysis (Marks et al., 2024; Makelov et al., 2024), and model-diffing (Lindsey et al., 2024).

Technically, an SAE consists of an encoder and a decoder:

$$\mathbf{f}(\mathbf{x}) = \sigma(\mathbf{W}^{\text{enc}}\mathbf{x} + \mathbf{b}^{\text{enc}}) \quad (1)$$

$$\hat{\mathbf{x}} = \mathbf{W}^{\text{dec}}\mathbf{f}(\mathbf{x}) + \mathbf{b}^{\text{dec}} \quad (2)$$

The encoder is trained to map high-dimensional activations or hidden states from inside a neural network into a sparse, overcomplete representation $\mathbf{f}(\mathbf{x})$, called the SAE *latents* or *features*. The activation function σ enforces non-negativity and sparsity in $\mathbf{f}(\mathbf{x})$, with a latent i typically considered active when $f_i(\mathbf{x}) > 0$.

The decoder maps the sparse features back to the original input space through a linear transformation, producing an approximate reconstruction $\hat{x}$. This reconstruction aims to minimize the mean squared error with respect to the original input x while maintaining the sparsity constraints on the intermediate representation.

Recent architectural improvements include JumpReLU SAEs (Rajamanoharan et al., 2024b), which learn individual activation thresholds, and TopK SAEs (Gao et al., 2024), which enforce exactly K active features per sample. For our Gemma-2-2B experiments, we use the BatchTopK (Bussmann et al., 2024) activation function, which we further describe in Section 3.2.

2.2. Challenges in Scaling SAEs

At the heart of SAE training lies a fundamental tension between reconstruction quality and sparsity. The sparsity objective, typically implemented via L1 regularization or explicit L0 constraints such as TopK, encourages SAEs to represent inputs using as few active latents as possible. While some degree of sparsity is necessary to extract interpretable features from superposition, optimizing for the reconstruction + sparsity loss often leads to pathological solutions. This manifests in several interrelated phenomena:

Feature splitting (Bricken et al., 2023) causes unified concepts to fragment into many specialized latents – for example, a feature responding to punctuation marks may split into separate features for question marks, periods, and commas. Even though each individual, specialized latent might be interpretable, the high-level concept of “punctuation mark” - which might be functionally used by the LLM - is now missing from the dictionary. Instead, the “punctuation mark” decoder direction is now implicitly represented in the decoder vector of each specific latent, as for example “comma” also implies “punctuation mark”. Leask et al. (2025) found that when doubling the dictionary size of an SAE, only one-third of the latents in the larger SAE represent novel concepts, while the other latents represent concepts similar to those of the smaller SAE in a sparser way.

Feature absorption (Chanin et al., 2024) occurs when a latent representing a general feature develops systematic blind spots in cases handled by more specialized latents. For example, consider a latent that activates on female names like “Mary”, “Jane”, “Sarah”, and “Lily”. If a specialized latent splits for the name “Lily”, the latent for female names might instead become an “all female names except Lily” latent. Unlike feature splitting where general concepts completely fragment, absorption creates specific “holes” in otherwise intact general features, making the general feature less reliable for downstream tasks and harder to interpret.

Feature composition represents another sparsity-driven dis-

tortion that occurs when features naturally co-occur. When presented with a direct product space of independent features (e.g. colors and shapes), the sparsity objective incentivizes learning single latents that capture specific combinations (like “red triangle”) rather than representing the underlying independent features (“red” and “triangle”) separately (Anders et al., 2024; Wattenberg & Viégas, 2024). Even when these features are conceptually and functionally independent in the model, combining them allows the SAE to achieve the same reconstruction with fewer active latents. Leask et al. (2025) have shown using “meta-SAEs” that the latents of sparse autoencoders are often composed of more fundamental “meta-latents”.

These issues become more severe as dictionary size increases. While larger dictionaries allow for representing more relevant features, they also enable more opportunities for sparsity-driven distortions. The reconstruction objective alone does not prevent these distortions since they achieve similar reconstruction with fewer active latents. Their persistence suggests they arise from fundamental limitations in standard SAE training objectives rather than optimization challenges. This creates a pressing need for new training approaches that can preserve features at multiple levels of abstraction while maintaining the benefits of sparsity. Ideally, we want SAEs that can capture both high-level concepts and their refinements without letting specialized features absorb or fragment their general counterparts.

2.3. Matryoshka Representation Learning

Matryoshka Representation Learning (MRL) (Kusupati et al., 2024) is an approach to representation learning that encodes information at varying levels of granularity within a single embedding vector. This allows the representation to be adaptable to the computational constraints of diverse downstream tasks. MRL is designed to modify existing representation learning pipelines with minimal overhead during training and no additional costs during inference and deployment. The key distinction of MRL lies in its loss function, which is modified to incorporate the performance when using only part of the embedding vector. The name “Matryoshka” is inspired by the nested Russian dolls, reflecting the nested nature of the representations.

3. Matryoshka Sparse Autoencoders

Matryoshka SAEs extend traditional sparse autoencoders by simultaneously training multiple nested autoencoders of increasing size, as illustrated in Figure 1. Given a maximum dictionary size m , we define a sequence of nested dictionary sizes $\mathcal{M} = m_1, m_2, \dots, m_n$ where $m_1 < m_2 < \dots < m_n = m$. Each size m_i corresponds to a sub-SAE that must reconstruct the input using only the first m_i latents.

We experiment with two types of Matryoshka SAE: *random prefix*, where $\mathcal{M}$ is randomly sampled from a distribution per batch, and *fixed prefix* where $\mathcal{M}$ is set as a hyperparameter before training. Formally, for an input $\mathbf{x} \in \mathbb{R}^n$, the encoder and decoder are defined as:

$$\mathbf{f}(\mathbf{x}) = \sigma(\mathbf{W}^{\text{enc}} \mathbf{x} + \mathbf{b}^{\text{enc}}) \quad (3)$$

$$\hat{\mathbf{x}}_i(\mathbf{f}) = \mathbf{W}_{0:m_i, :}^{\text{dec}} \mathbf{f}_{0:m_i} + \mathbf{b}^{\text{dec}} \quad \text{for } m_i \in \mathcal{M} \quad (4)$$

where $\mathbf{W}^{\text{enc}} \in \mathbb{R}^{m \times n}$ is the encoder matrix, $\mathbf{W}^{\text{dec}} \in \mathbb{R}^{n \times m}$ is the decoder matrix, $\mathbf{b}^{\text{enc}} \in \mathbb{R}^m$ is the encoder bias, $\mathbf{b}^{\text{dec}} \in \mathbb{R}^n$ is the decoder bias, and the subscript notation $0:m_i$ indicates taking the first m_i rows or elements. Each nested decoder $\mathbf{W}_{0:m_i, :}^{\text{dec}}$ must learn to reconstruct the input using only a subset of the latents, creating a hierarchy of representations at different scales.

3.1. Training Objective

The key innovation in Matryoshka SAEs is the training objective that enforces good reconstruction at multiple scales simultaneously:

$$\mathcal{L}(\mathbf{x}) = \sum_{m \in \mathcal{M}} \|\mathbf{x} - \underbrace{(\mathbf{f}(\mathbf{x})_{0:m} \mathbf{W}_{0:m, :}^{\text{dec}} + \mathbf{b}^{\text{dec}})}_{\text{reconstruction using first } m \text{ latents}}\|_2^2 + \alpha \mathcal{L}_{\text{aux}} \quad (5)$$

where $\mathcal{L}_{\text{aux}}$ is the standard auxiliary loss as used in Gao et al. (2024). The first m_1 latents must learn to reconstruct the input as well as possible on their own, the first m_2 latents must do the same, and so on. Early latents must be able to reconstruct the input when m_i is smaller, creating pressure for them to capture general, widely applicable features. Later latents only participate in reconstructing larger dictionaries, allowing them to specialize in more specific features.

3.2. BatchTopK Activation Function

For our Gemma-2-2B language model experiments, we use the BatchTopK activation function (Bussmann et al., 2024) to determine which latents are active. The BatchTopK activation function improves upon standard element-wise approaches by considering sparsity across batches rather than individual examples. For a batch $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_B]$, BatchTopK retains the $B \times K$ largest activations across the entire batch while setting others to zero:

$$\text{BatchTopK}(\mathbf{X}) = \mathbf{X} \odot \mathbf{1}[\mathbf{X} \geq \tau(\mathbf{X})] \quad (6)$$

where $\tau(\mathbf{X})$ is the $(B \times K)$ th largest value in $\mathbf{X}$, $\odot$ denotes element-wise multiplication, and $\mathbf{1}[\cdot]$ is the indicator function. This allows the number of active latents per example

to vary naturally while maintaining an average sparsity of K across the batch. During inference, BatchTopK is replaced with a global threshold to ensure consistent behavior independent of the batch:

$$\text{BatchTopK}(\mathbf{x}) = \mathbf{x} \odot \mathbf{1}[\mathbf{x} \geq \theta] \quad (7)$$

where θ is calibrated on the training data to maintain the desired average sparsity.

4. Experiments

To understand how Matryoshka SAEs prevent feature absorption and fragmentation, we conduct experiments across three settings of increasing complexity: a synthetic toy model designed to exhibit feature absorption, a qualitative investigation of the features learned on a small 4-layer TinyStories model, and finally benchmarks on the larger Gemma-2-2B architecture.

4.1. Toy Model Demonstration of Feature Absorption

We first demonstrate the ability of Matryoshka SAEs to avoid feature absorption in a controlled, synthetic setting, with a toy model following the model introduced in Chanin et al. (2024). Concretely, we construct a tree-structured set of L binary features, each mapped to a unique direction in $\mathbb{R}^d$ (we use $d = 20$). The root feature is always present (though it is excluded from L). Every other feature is sampled conditional on its parent being present, with an associated edge probability $p_{(\text{parent} \rightarrow \text{child})} \in (0, 1)$. If a feature is active, we add its direction vector to produce the final d -dimensional input. This induces a hierarchical dependency analogous to “comma” $\implies$ “punctuation mark” in real text: child features *always* appear with their parent features.

We train two types of autoencoders on data generated from this model: a standard sparse autoencoder (*Vanilla SAE*) with L latents, matching the true number of features and a *Matryoshka SAE* with the same total latents L , but with additional nested reconstruction objectives on sub-prefixes of the latents. For the Matryoshka SAE, we sample the prefix lengths from a truncated Pareto distribution at each training step, and use a ReLU activation function. Both models have the same encoder and decoder architecture (see Appendix C.1 for more training details).

Figure 3 shows the ground-truth features against the closest-matching Matryoshka and Vanilla activations on a single batch for a single parent and its children features (for a plot of the activations of all features see Figure 11). The Vanilla SAE exhibits feature absorption - when one of the child features is active, the corresponding parent feature does not activate. In contrast, the Matryoshka SAE recovers the true




Figure 2. Toy model decoder vector similarity. Graphical representation of the toy model (left). The heatmaps show the cosine similarity between learned latent vectors and ground-truth feature vectors for the Matryoshka SAE (middle) and Vanilla SAE (right). The Matryoshka SAE shows a clear diagonal structure, which demonstrates disentanglement of the hierarchical features and learning the ground truth. The Vanilla SAE, however, exhibits high similarity between parent and child latents, indicating feature absorption.




Figure 3. Toy model activations. Ground-truth feature activations alongside Matryoshka SAE activations and Vanilla activations for one of the parents and its children on the toy model. Notice that in the Vanilla SAE the parent latent (column bracketed in blue) does not fire when its children (bracketed in red) fire. For the activations of all features in the toy model, see Figure 11.

underlying feature structure, with parent features remaining active regardless of which child features are active.

To further see whether the autoencoders have found the ground truth features, we measure the cosine similarity between the learned decoder vectors and the ground-truth feature vectors (Figure 2). For the Vanilla SAE, we observe high similarity between parent and child latents, indicating that the model is learning redundant representations and failing to disentangle the hierarchical features. The Matryoshka SAE, on the other hand, shows a clear diagonal structure in its similarity matrix, demonstrating that it is able to recover the true feature vectors with minimal overlap, with only small nonzero off-diagonal terms.

4.2. TinyStories Investigation

We now examine feature absorption in a 4-layer Transformer¹ (hidden size 768) trained on TinyStories (Eldan & Li, 2023), a dataset of simple English children’s stories. To systematically study how features evolve as SAE dictionary size increases, we train a family of small “reference” SAEs with varying dictionary sizes (30, 100, 300, 1k, 3k, and 10k latents) on three model locations: attention block outputs, MLP outputs, and the residual stream.

These reference SAEs reveal a consistent pattern: smaller SAEs (like S/1 with 100 latents) often capture broad, general concepts, while larger ones tend to fragment these concepts into specialized cases. To visualize these relationships, we developed an interactive tool (sparselatents.com/tree-view) that traces how latents evolve across SAE sizes. Technical details about the visualization methodology can be found in Appendix F.

“Female tokens” Absorption Example. In Figure 4, we highlight latent 65 in our 300-latent reference SAE (S/2/65), which consistently fires on female-coded tokens: “she,” “her,” “girl,” and female names like “Lily” and “Sue.” When we examine the corresponding latents in the 1000-latent SAE (S/3), we find that the general female-words latent (S/3/66) has developed systematic “holes”—it stops firing on “Lily” because a specialized “Lily-specific” latent (S/3/359) has absorbed that case. The original concept becomes fragmented, obscuring that the model treats “Lily” as a female name.

In contrast, when we train a 25k-latent Matryoshka SAE with the same average sparsity, it maintains the broad coverage of female-coded tokens while still developing specialized child latents (e.g. a Sue latent, Lily latent, and a female names latent without holes, see [here](#)). This preserves the hierarchical relationship between general and specific features, making the model’s representations easier to interpret.

¹<https://github.com/noanabeshima/tinymodel>



Figure 4. Feature absorption in a TinyStories model. Example activations showing how a general “female words” latent (S/2/65) develops holes in a larger SAE (S/3/66) as specialized latents for “Lily” (S/3/359) and “Sue” (S/3/861) absorb specific cases. Feature absorption locations are circled in red.

4.3. Larger-Scale Validation with Gemma-2-2B

In our Gemma-2-2B language model experiment, we use fixed prefix Matryoshka over random as we found that they attained slightly better eval metrics (see Appendix G for details). We use a dictionary size of $D = 65536$ with five nested sub-SAEs where $\mathcal{M} = \{2048, 6144, 14336, 30720, 65536\}$. We train five Matryoshka SAEs with an average sparsity of respectively 20, 40, 80, 160, and 320 active latents per token.

The SAEs were trained on the residual stream activations from layer 12 of Gemma 2-2B using 500M tokens sampled from The Pile (Gao et al., 2020). We use the Adam optimizer with a learning rate of 3×10^{-4} and batch size of 2048. No additional regularization terms are used beyond the implicit regularization from the multiple reconstruction objectives.

We evaluate Matryoshka SAEs against six alternative SAE architectures (see Appendix A) to assess both basic performance and effectiveness in addressing feature absorption and composition challenges. We use the evaluations pro-

vided by SAE Bench (Karvonen et al., 2024a) to compare Matryoshka SAEs to the baseline SAEs across a wide variety of tasks. In these evaluations, we use all latents of the Matryoshka SAE. In Appendix D, we study the performance of the different nested sub-SAEs.

Reconstruction and Downstream Performance. We first examine basic performance metrics of the SAEs: reconstruction quality (measured by variance of the input explained by the reconstruction) and downstream performance (measured by cross-entropy loss when feeding reconstructed activations back into the language model). Figure 5 shows these metrics across different sparsity levels.




Figure 5. Reconstruction performance. The variance explained of Matryoshka SAEs is slightly worse than some competing architectures. However, looking at the downstream LLM CE loss, Matryoshka SAEs perform comparable, especially at larger L0s.

Matryoshka SAEs consistently show slightly worse reconstruction compared to standard BatchTopK SAEs. For example, at a L0 of 40 the reconstruction of Matryoshka SAEs explains 70% of the variance in model activations, whereas BatchTopK SAEs explain 72% of the variance. This performance gap is expected, as the nested reconstruction constraints prevent Matryoshka SAEs from splitting and absorbing broad latents to optimize sparsity. Therefore, it would be surprising if there was not mild degradation here.

However, downstream model loss tells a different story. Despite worse reconstruction, at larger L0s, Matryoshka SAEs achieve comparable cross-entropy loss when their reconstructed activations are fed back into the language model. This suggests that Matryoshka SAEs may be learning more meaningful features that better capture the language model’s internal representations, even though they reconstruct the raw activations less accurately.

Feature Absorption and Splitting. We evaluate the extent to which Matryoshka SAEs reduce two key pathologies that emerge when scaling sparse autoencoders: feature absorption and feature splitting. Following the methodology of Chanin et al. (2024), we use first-letter classification tasks as a probe to measure these phenomena.

To quantify feature absorption, we first train logistic regres-




Figure 6. Feature absorption and splitting. Matryoshka SAEs significantly reduce feature absorption and splitting compared to baselines.

sion probes on the raw activations to establish a ground truth for which directions encode first-letter information. We then find the corresponding latents and measure when these SAE latents fail to activate on tokens starting with that letter and get absorbed into specific token-aligned latents.

Secondly, we also analyze feature splitting by measuring when additional latents significantly improve classification performance for first-letter classification, indicating the first-letter concept is split over multiple latents.

As shown in Figure 6, Matryoshka SAEs reduces both feature absorption and splitting compared to the baseline architectures. For example, at an average sparsity (L0) of 40 active latents per token, Matryoshka SAEs exhibit an absorption rate of just 0.05, in contrast to 0.49 for BatchTopK SAEs. This indicates that Matryoshka SAEs are far more successful at preserving general first-letter features in the presence of specialized latents.

Similarly, Matryoshka SAEs require fewer latents to capture first-letter features, demonstrating reduced feature splitting. At an L0 of 40, Matryoshka SAEs need only one latent per first letter on average, while BatchTopK SAEs split this information across three latents.

These results demonstrate that the hierarchical structure imposed by Matryoshka SAEs is effective at combating the feature absorption and splitting that typically emerge when increasing SAE dictionary size. Matryoshka SAEs are able to maintain coherent high-level features even in the presence of specialized low-level features.

Sparse Probing and Targeted Concept Removal. We evaluate the quality of the features learned by the SAEs through three complementary analyses that measure different aspects of feature quality and disentanglement.

First, we perform sparse probing to assess whether the learned latents encode semantically meaningful concepts (Gao et al., 2024). We use a suite of 35 binary classification tasks spanning five domains: professions, sentiment, language, code, and news. For each task, we encode the inputs




Figure 7. Probing and concept isolation. Matryoshka SAEs score considerably better at SCR and TPP than baseline models, and better at sparse probing at lower sparsities.

through the SAE, apply mean pooling over non-padding tokens, select the most relevant latent, and train a logistic regression probe. Strong performance on these tasks indicates that the SAE has learned relevant latents and are not missing, entangled, or split across multiple latents.

Second, we test the SAE’s ability to separate dataset biases, such as the spurious correlation between a profession like “nurse” always being associated with “female”, using spurious correlation removal (SCR) (Karvonen et al., 2024b). We train classifiers on intentionally biased data, then ablate the latents that encode the spurious signal. A higher resulting accuracy suggests the SAE was more effective at isolating the spurious correlation (e.g. gender), enabling the classifier to focus on the intended task (e.g. profession classification).

Third, we evaluate how well the SAE isolates individual concepts using Targeted Probe Perturbation (TPP) (Karvonen et al., 2024b). We ablate latents specific to one class and measure the change in accuracy of a classifier trained for that class as well as classifiers for other classes. If the ablated latents are properly disentangled, it should have an isolated causal effect - reducing accuracy on the relevant class probe while leaving other class probes unaffected.

As shown in Figure 7, Matryoshka SAEs significantly outperform all benchmark architectures on the TPP and SCR metrics, demonstrating they learn more disentangled and isolated concept representations. On the sparse probing tasks, Matryoshka SAEs achieve the best performance at lower sparsity levels and consistently outperform standard BatchTopK SAEs across all sparsity levels. These results strongly suggest that the hierarchical structure imposed by Matryoshka SAEs leads to higher quality and more relevant learned features.




Figure 8. **Average maximum cosine similarity.** The maximum cosine similarity is considerably lower for Matryoshka SAEs, indicating less feature composition.

Feature Composition Analysis. To quantify feature composition and shared information between latents, we look at the average maximum cosine similarity between two latents in an SAE. If the SAE has many latents with high cosine similarity, this means that multiple latents are representing similar information, indicating composition of information.

In Figure 8, we see that the average maximum cosine similarity between decoder vectors is substantially lower for Matryoshka SAEs, indicating feature composition is less of a problem for Matryoshka SAEs. In Appendix D.1, we investigate this phenomenon further using meta-SAEs and find that the same holds.

Automated Interpretability. We evaluate feature interpretability using gpt4o-mini as an LLM judge. Following Bills et al. (2023) and Paulo et al. (2024), we generate feature explanations for 1000 latents from a range of activating dataset examples and score the explanation by asking an LLM judge to predict on which inputs a latent will be active. Figure 9 shows that the interpretability of latents of Matryoshka SAEs are comparable to standard BatchTopK SAEs and more interpretable than many alternative baseline architectures.




Figure 9. **Automatic interpretability score.** Matryoshka SAE latents are among the most interpretable.

Scaling SAEs To evaluate how Matryoshka SAEs performance changes when increasing dictionary size, we train




Figure 10. **Scaling evaluations.** Performance across three different dictionary sizes. Average performance across L0s 40, 80, 160 are reported.

Matryoshka SAEs with dictionary sizes of 4k, 16k, and 65k. Figure 10 shows that on almost all metrics Matryoshka SAEs improve or are stable with scale, whereas alternative architectures often degrade. These results indicate that as one increases the dictionary size of the SAE to capture more of features of the LLM, Matryoshka SAEs may be the superior choice for many downstream tasks.

5. Limitations

While Matryoshka SAEs demonstrate improved performance of downstream tasks, this comes at the cost of slightly reduced reconstruction performance compared to standard SAEs. Although downstream loss remains competitive, applications requiring precise activation reconstruction may find this trade-off problematic.

Furthermore, the multiple nested reconstruction objectives increases training time compared to standard SAEs. This increase is dependent on the length of the set of nested dictionary sizes ($\mathcal{M}$). In our experiments with 5 nested SAEs, this increased training time with circa 50%.

Finally, our evaluation relies heavily on quantitative measures of feature quality and automated interpretability metrics. While these provide valuable insights, they may not fully capture human-relevant aspects of feature interpretability or the practical utility of the learned representations for manual downstream analysis. Future work investigating the human interpretability of Matryoshka SAE latents and their applicability to practical interpretability tasks would be valuable.

6. Discussion and Conclusion

We have presented Matryoshka SAEs, a novel variant of sparse autoencoders that addresses key challenges with spar-

sity when scaling SAEs. By enforcing hierarchical feature learning through nested dictionaries, our approach achieves lower feature absorption rates and improved concept isolation compared to standard approaches. This demonstrates that the pathologies of scaling SAEs are not fundamental limitations, but rather artifacts of myopic training objectives that can be overcome through simple modifications to the training process.

Looking forward, the success of this approach suggests that many apparent problems with SAEs may have surprisingly simple solutions waiting to be discovered. As we work to understand increasingly large and capable neural networks, approaches to interpretability that scale as well will become essential to better understanding these systems.

Impact Statement

This work advances mechanistic interpretability techniques for neural networks. Better understanding of model internals could help identify failure modes, verify intended behaviors, and guide the development of more reliable and safe AI systems. However, advances in model interpretability may also accelerate progress in large language model development and deployment, potentially amplifying both positive and negative societal impacts of these technologies.

Acknowledgments

We are extremely grateful for feedback, advice, edits, helpful discussions, and support from Joel Becker, Gytis Daujotas, Julian D’Costa, Leo Gao, Collin Gray, Dan Hendrycks, Benjamin Hoffner-Brodsky, Patrick Leask, Mason Krug, Hunter Lightman, Mark Lippmann, Charlie Rogers-Smith, Logan R. Smith, Glen Taggart, Adly Templeton, and Matthew Wearden.

This research was made possible by funding from Light-speed Grants and AI Safety Support.

References

- Anders, E., Neo, C., Hoelscher-Obermaier, J., and Howard, J. N. Sparse autoencoders find composed features in small toy models. <https://www.lesswrong.com/posts/a5wwqza2cY3W7L9cj/sparse-autoencoders/find-composed-features-in-small-toy>, 2024.
- Anthropic Interpretability Team. Training sparse autoencoders. <https://transformer-circuits.pub/2024/april-update/index.html#training-saes>, 2024. [Accessed January 20, 2025].
- Biderman, S., Schoelkopf, H., Anthony, Q., Bradley, H., O’Brien, K., Hallahan, E., Khan, M. A., Purohit, S., Prashanth, U. S., Raff, E., Skowron, A., Sutawika, L., and van der Wal, O. Pythia: A suite for analyzing large language models across training and scaling, 2023. URL <https://arxiv.org/abs/2304.01373>.
- Bills, S., Cammarata, N., Mossing, D., Tillman, H., Gao, L., Goh, G., Sutskever, I., Leike, J., Wu, J., and Saunders, W. Language models can explain neurons in language models. URL <https://openaipublic.blob.core.windows.net/neuron-explainer/paper/index.html>. (Date accessed: 14.05. 2023), 2, 2023.
- Bricken, T., Templeton, A., Batson, J., Chen, B., Jermyn, A., Conerly, T., Turner, N., Anil, C., Denison, C., Askell, A., et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2, 2023.
- Busmann, B., Leask, P., and Nanda, N. Batchtopk sparse autoencoders. *arXiv preprint arXiv:2412.06410*, 2024.
- Chanin, D., Wilken-Smith, J., Dulka, T., Bhatnagar, H., and Bloom, J. A is for absorption: Studying feature splitting and absorption in sparse autoencoders. *arXiv preprint arXiv:2409.14507*, 2024.
- Cunningham, H., Ewart, A., Riggs, L., Huben, R., and Sharkey, L. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023.
- Eldan, R. and Li, Y. Tinstories: How small can language models be and still speak coherent english?, 2023. URL <https://arxiv.org/abs/2305.07759>.
- Elhage, N., Hume, T., Olsson, C., Schiefer, N., Henighan, T., Kravec, S., Hatfield-Dodds, Z., Lasenby, R., Drain, D., Chen, C., et al. Toy models of superposition. *arXiv preprint arXiv:2209.10652*, 2022.
- Gao, L., Biderman, S., Black, S., Golding, L., Hoppe, T., Foster, C., Phang, J., He, H., Thite, A., Nabeshima, N., et al. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027*, 2020.
- Gao, L., la Tour, T. D., Tillman, H., Goh, G., Troll, R., Radford, A., Sutskever, I., Leike, J., and Wu, J. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2406.04093*, 2024.
- Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., Casas, D. d. L., Hendricks, L. A., Welbl, J., Clark, A., et al. Training compute-optimal large language models. *arXiv preprint arXiv:2203.15556*, 2022.